



JenOS User Guide

JN-UG-3075
Revision 1.8
25 August 2016

Contents

Preface	9
Organisation	9
Conventions	10
Acronyms and Abbreviations	10
Related Documents	11
Support Resources	11
Trademarks	11
Chip Compatibility	11

Part I: Concept and Operational Information

1. Introduction	15
1.1 Modules and Architecture	15
1.1.1 JenOS Modules	15
1.1.2 Software Architecture	16
1.2 Resources	17
2. Real-time Operating System (RTOS)	19
2.1 RTOS Fundamentals	19
2.2 Introduction to the JenOS RTOS	20
2.3 RTOS Configuration	20
2.4 RTOS Concepts and Features	21
2.4.1 User Tasks	21
2.4.2 Interrupt Service Routines (ISRs)	22
2.4.3 Priorities and Scheduling	22
2.4.4 Task/ISR States	23
2.4.5 State Transitions	24
2.4.6 Activity Scheduling (using Software Timers)	25
2.4.7 Mutual Exclusion (Mutex)	27
2.4.8 Inter-task Communication (using Messages)	28
2.5 OS Error Callback Function	30
2.5.1 Strict Error Checks	30
2.5.2 Handling OS Errors	30

3. Persistent Data Manager (PDM) for Flash Memory	33
3.1 Overview	33
3.2 Initialising the PDM	34
3.3 Data Storage in NVM	35
3.4 Recovering Data from NVM	36
3.5 Saving Data to NVM	37
3.6 Deleting Data in NVM	37
3.7 Mutexes in PDM	38
3.8 Ensuring Consistency of PDM Records	38
4. Persistent Data Manager (PDM) for EEPROM	39
4.1 Overview	39
4.2 Initialising the PDM and Building a File System	40
4.3 Managing Data in EEPROM	41
4.3.1 Saving Data to EEPROM	42
4.3.2 Recovering Data from EEPROM	43
4.3.3 Deleting Data in EEPROM	43
4.4 Storing Counters in EEPROM	44
4.4.1 Creating a Counter	44
4.4.2 Incrementing a Counter	44
4.4.3 Reading a Counter	44
4.4.4 Deleting a Counter	45
4.5 PDM Features	45
4.5.1 Mutex in PDM	45
4.5.2 Event and Error Handler for EEPROM	45
4.5.3 EEPROM Capacity	46
4.5.4 EEPROM Wear Count	46
4.5.5 Ensuring Consistency of PDM Records	47
5. Power Manager (PWRM)	49
5.1 Low-Power Modes	49
5.1.1 Doze Mode	49
5.1.2 Sleep Mode with Memory Held	49
5.1.3 Sleep Mode without Memory Held	50
5.1.4 Deep Sleep Mode	50
5.2 Callback Functions for Power Manager	51
5.2.1 Essential Callback Function	51
5.2.2 Pre-sleep and Post-sleep Callback Functions	51
5.2.3 Wake Timer Callback Function	52
5.3 Initialising and Starting the Power Manager	52
5.4 Enabling Power-Saving	53

5.5 Non-interruptible Activities	53
5.6 Terminating Low-Power Mode	54
5.7 Scheduling Wake Events	55
5.8 Doze Mode	55
5.8.1 Circumstances that Lead to Doze Mode	56
5.8.2 Doze Mode Monitoring During Development	57
6. Protocol Data Unit Manager (PDUM)	59
6.1 Message Assembly and Disassembly	59
6.2 Preparing the PDU Manager	60
6.3 Inserting Data into Outgoing Message	61
6.4 Extracting Data from Incoming Message	62
7. Debug (DBG) Module	63
7.1 Overview	63
7.2 Enabling the Debug Module	64
7.3 Initialising and Configuring the Debug Module	64
7.3.1 Using JN516x UART Input/Output	64
7.3.2 Using Alternative Serial Output	65
7.4 Debug Configuration Flags	66
7.5 Example Diagnostic Code	67
 Part II: Reference Information	
8. RTOS API	71
8.1 RTOS Macros	71
OS_TASK	72
OS_ISR	73
OS_SWTIMER_CALLBACK	74
OS_HWCOUNTER_ENABLE_CALLBACK	75
OS_HWCOUNTER_DISABLE_CALLBACK	76
OS_HWCOUNTER_SET_CALLBACK	77
OS_HWCOUNTER_GET_CALLBACK	78
8.2 RTOS Functions	79
8.2.1 Initialisation Functions	79
OS_vStart	80
OS_vRestart	81
8.2.2 User Task Functions	82
OS_eActivateTask	83
OS_eGetCurrentTask	84

Contents

8.2.3 Interrupt Functions	85
OS_eDisableAllInterrupts	86
OS_eEnableAllInterrupts	87
OS_eSuspendOSInterrupts	88
OS_eResumeOSInterrupts	89
8.2.4 Mutex Functions	90
OS_eEnterCriticalSection	91
OS_eExitCriticalSection	92
8.2.5 Messaging Functions	93
OS_ePostMessage	94
OS_eCollectMessage	95
OS_eGetMessageStatus	96
8.2.6 Software Timer Functions	97
OS_eStartSWTimer	98
OS_eStopSWTimer	99
OS_eExpireSWTimers	100
OS_eContinueSWTimer	101
OS_eGetSWTimerStatus	102
9. PDM API for Flash Memory	103
PDM_vInit	104
PDM_vSPIFlashConfig	106
PDM_eLoadRecord	107
PDM_vSaveRecord	109
PDM_vSave	110
PDM_vDeleteRecord	111
PDM_vDelete	112
PDM_vWarmInitHw	113
PDM_vRegisterSystemCallback	114
10. PDM API for EEPROM	115
10.1 EEPROM PDM Functions	116
PDM_eInitialise	117
PDM_eSaveRecordData	118
PDM_eReadDataFromRecord	119
PDM_eDeleteData	120
PDM_eDeleteAllData	121
PDM_u8GetSegmentCapacity	122
PDM_u8GetSegmentOccupancy	123
PDM_bDoesDataExist	124
10.2 EEPROM PDM Bitmap Counter Functions	125
PDM_eCreateBitmap	126
PDM_eIncrementBitmap	127
PDM_eGetBitmap	128
PDM_eDeleteBitmap	129

10.3 EEPROM PDM Miscellaneous Functions	130
PDM_vRegisterSystemCallback	131
PDM_vSetWearCountTriggerLevel	132
PDM_eGetSegmentWearCount	133
11. PWRM API	135
11.1 Core Functions	135
PWRM_vInit	136
PWRM_eStartActivity	137
PWRM_eFinishActivity	138
PWRM_u16GetActivityCount	139
PWRM_eScheduleActivity	140
PWRM_vManagePower	141
11.2 Callback Set-up Functions	142
vAppMain	143
PWRM_vRegisterPreSleepCallback	144
PWRM_vRegisterWakeupCallback	145
vAppRegisterPWRMCallbacks	146
PWRM_vWakeInterruptCallback	147
11.3 Debugging Function	148
PWRM_vSetupDozeMonitor	149
12. PDUM API	151
PDUM_vInit	152
PDUM_hAPduAllocateAPduInstance	153
PDUM_eAPduFreeAPduInstance	154
PDUM_u16APduInstanceReadNBO	155
PDUM_u16APduInstanceWriteNBO	156
PDUM_u16APduInstanceWriteStrNBO	157
PDUM_u16SizeNBO	158
PDUM_u16APduGetSize	159
PDUM_pvAPduInstanceGetPayload	160
PDUM_u16APduInstanceGetPayloadSize	161
PDUM_eAPduInstanceSetPayloadSize	162
PDUM_vDBGPrintAPduInstance	163
13. DBG API	165
DBG_vInit	166
DBG_vUartInit	167
DBG_vPrintf	168
DBG_vAssert	169
DBG_vDumpStack	170
DBG_vFlush	171
DBG_iGetChar	172

14. JenOS Structures	173
14.1 PDM_tsHwFncTable	173
14.2 tSPIfIashFncTable	174
14.3 PWRM_teSleepMode	176
14.4 DBG_tsFunctionTbl	176
14.5 tsReg128	177
14.6 PDM_tpfvSystemEventCallback	177
14.7 PDM_eSystemEventCode	177
14.8 PDM_teStatus	180
14.9 OS_teStatus	182

Part III: Configuration Information

15. JenOS Configuration	187
15.1 CPU Stack and Heap Sizes	187
15.2 Configuration Principles	187
15.3 Configuring JenOS Resources	189

Part IV: Appendices

A. Hardware Counter Details	193
A.1 Hardware Counter Operation	193
A.2 Use of Tick Timer as Hardware Counter	194
B. Clearing Interrupts	195

Preface

This manual provides a single point of reference for information relating to the Jennic Operating System, referred to as JenOS. The manual provides both conceptual and practical information concerning JenOS, and provides guidance on use of the JenOS Application Programming Interfaces (APIs). The API resources (functions and structures) are fully detailed.

JenOS is designed to be used with the NXP ZigBee PRO stack on the NXP JN516x wireless microcontrollers. This manual should be used throughout ZigBee PRO wireless network application development, along with the *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)*.

Organisation

This manual is divided into four parts:

- **Part I: Concept and Operational Information** comprises seven chapters:
 - **Chapter 1** introduces JenOS, including its five modules and APIs
 - **Chapter 2** describes how to use the Real-time Operating System (RTOS)
 - **Chapter 3** describes how to use the Persistent Data Manager (PDM) for Flash memory
 - **Chapter 4** describes how to use the Persistent Data Manager (PDM) for EEPROM
 - **Chapter 5** describes how to use the Power Manager (PWRM)
 - **Chapter 6** describes how to use the Protocol Data Unit Manager (PDUM)
 - **Chapter 7** describes how to use the Debug (DBG) module
- **Part II: Reference Information** comprises seven chapters:
 - **Chapter 8** describes the functions of the RTOS API
 - **Chapter 9** describes the functions of the PDM API for Flash memory
 - **Chapter 10** describes the functions of the PDM API for EEPROM
 - **Chapter 11** describes the functions of the PWRM API
 - **Chapter 12** describes the functions of the PDUM API
 - **Chapter 13** describes the functions of the DBG API
 - **Chapter 14** details the structures used by the JenOS APIs
- **Part III: Configuration Information** comprises one chapter:
 - **Chapter 15** outlines the static configuration required to use JenOS and its resources
- **Part IV: Appendices** comprises two appendices that describe the use of hardware counters and clearing interrupts.

Conventions

Files, folders, functions and parameter types are represented in **bold** type.

Function parameters are represented in *italics* type.

Code fragments are represented in the `Courier New` typeface.



This is a **Tip**. It indicates useful or practical information.



This is a **Note**. It highlights important additional information.



*This is a **Caution**. It warns of situations that may result in equipment malfunction or damage.*

Acronyms and Abbreviations

APDU	Application Protocol Data Unit
API	Application Programming Interface
DBG	Debug
EEPROM	Electrically Erasable Programmable Read-Only Memory
ISR	Interrupt Service Routine
JenOS	Jennic Operating System
MAC	Media Access Control
PAN	Personal Area Network
NPDU	Network Protocol Data Unit
NVM	Non-Volatile Memory
OS	Operating System
PDU	Protocol Data Unit
PDUM	Protocol Data Unit Manager
PDM	Persistent Data Manager

PIC	Programmable Interrupt Controller
PWRM	Power Manager
RTOS	Real-Time Operating System
SDK	Software Developer's Kit
UART	Universal Asynchronous Receiver-Transmitter
ZPS	ZigBee PRO Stack

Related Documents

JN-UG-3048	ZigBee PRO Stack User Guide (for SE)
JN-UG-3101	ZigBee PRO Stack User Guide (for ZLL and HA)
JN-UG-3087	JN516x Integrated Peripherals API User Guide
JN-AN-1135	ZigBee Smart Energy HAN Solutions Application Note
JN-AN-1171	ZigBee Light Link Solution Application Note
JN-AN-1189	ZigBee Home Automation Demonstration Application Note
JN-DS-JN516x	JN516x Data Sheet (for JN5168, JN5164 and JN5161)
JN5169	JN5169 Data Sheet

Support Resources

To access online support resources such as SDKs, Application Notes and User Guides, visit the Wireless Connectivity area of the NXP web site:

www.nxp.com/products/interface-and-connectivity/wireless-connectivity

ZigBee resources can be accessed from the ZigBee page, which can be reached via the short-cut www.nxp.com/zigbee.

All NXP resources referred to in this manual can be found at the above addresses, unless otherwise stated.

Trademarks

All trademarks are the property of their respective owners.

Chip Compatibility

The software described in this manual can be used on the NXP JN516x family of wireless microcontrollers with the exception of JN5161 device. However, the supported devices will be referred to as JN516x.

Part I: Concept and Operational Information

1. Introduction

The Jennic Operating System (JenOS) is designed for use in wireless network applications for the NXP JN516x device, providing an interface which simplifies the programming of a range of operations that are not specific to wireless networking.

JenOS is primarily intended for use with the NXP ZigBee PRO stack in order to develop wireless network applications based on the ZigBee standard. Therefore, this manual should be studied in conjunction with the *ZigBee PRO Stack User Guide* (JN-UG-3048 or JN-UG-3101) - other reference resources are detailed in [Section 1.2](#).

1.1 Modules and Architecture



JenOS is organised into five modules, each with a dedicated Application Programming Interface (API) to facilitate easy interaction between the application and the JenOS module. Each API consists of a set of C functions and associated resources.

In addition, a configuration editor is provided which allows the graphical configuration of the JenOS resources used by the application. This tool is known as the JenOS Configuration Editor and is a plug-in for the Eclipse IDE (Integrated Development Environment) - the tool is described in [Chapter 15](#) and [Chapter 14](#).

1.1.1 JenOS Modules



The JenOS modules are briefly described below:

- **Real-time Operating System (RTOS):** This module provides a mechanism for reacting to real-time events in a way that optimises the efficiency and reliability of the system. The RTOS module is described in [Chapter 2](#).
- **Persistent Data Manager (PDM):** This module handles the storage of context and application data in Non-Volatile Memory (NVM), and the retrieval of this data. It provides a mechanism by which the JN516x device can resume operation without loss of continuity following a power loss. The PDM module is available in two editions - one for external SPI Flash memory, described in [Chapter 3](#), and another for JN516x internal EEPROM, described in [Chapter 4](#).
- **Power Manager (PWRM):** This module manages the transitions of the JN516x device into and out of low-power modes, such as sleep mode. The PWRM module is described in [Chapter 5](#).
- **Protocol Data Unit Manager (PDUM):** This module is concerned with managing memory, as well as inserting data into messages to be transmitted and extracting data from messages that have been received. The PDUM module is described in [Chapter 6](#).
- **Debug module (DBG):** This module allows diagnostic messages to be output when the application runs, as an aid to debugging the application code. The DBG module is described in [Chapter 7](#).

1.1.2 Software Architecture

On a JN516x-based node in a ZigBee PRO wireless network, JenOS interacts with the following software blocks:

- User application (through use of the JenOS APIs in the application code)
- ZigBee PRO stack
- JN516x integrated peripherals

JenOS can be envisaged as sitting alongside the ZigBee PRO stack and the JN516x Integrated Peripherals API, as depicted in the diagram below.

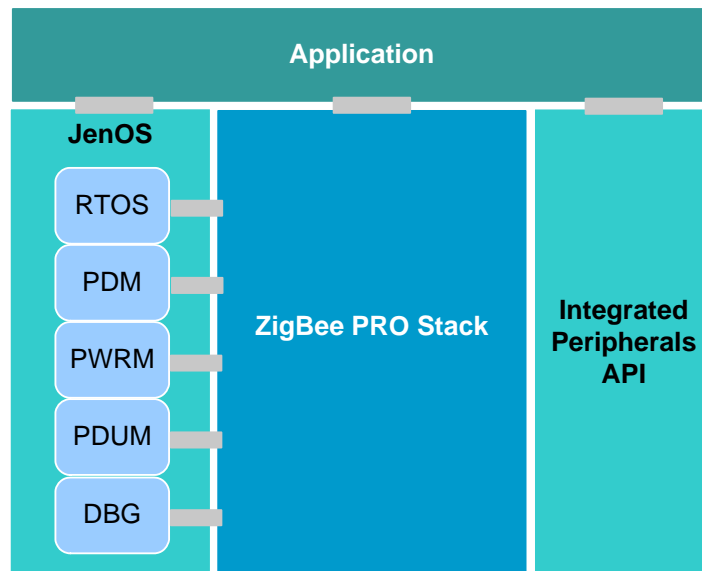


Figure 1: Basic Software Architecture

1.2 Resources

JenOS and related components are supplied in the NXP Software Developer's Kit (SDK) installers, as follows:

- JenOS is installed as part of the 'JN516x SDK Libraries' for the ZigBee application profiles (ZigBee Light Link, Home Automation, Smart Energy).
- The JenOS Configuration Editor (NXP plug-in for the Eclipse IDE) is also provided in these installers.

In addition to this manual, a number of other reference resources are available from the Wireless Connectivity area of the NXP web site (see ["Support Resources" on page 11](#)) to aid the development of application code which uses JenOS:

- *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)* provides some guidance on the use of JenOS API functions in ZigBee PRO application code.
- Application Notes which provide example code for different ZigBee PRO application profiles:
 - JN-AN-1135 for Smart Energy
 - JN-AN-1171 for ZigBee Light Link
 - JN-AN-1189 for Home Automation

2. Real-time Operating System (RTOS)

This chapter introduces the concept of a Real-time Operating System (RTOS) and introduces the main features of the RTOS which is a module of JenOS.

An RTOS allows the host system to react to multiple real-time events and schedule processing to meet the deadlines of these events. For example:

- An application that controls a set of traffic lights at a crossroads, where pedestrians can press buttons to request crossing any of the roads
- An application that handles data sampling, compression and storage in a digital sound-recording system

The above applications each have a number of tasks with different urgencies competing for CPU time.

2.1 RTOS Fundamentals

An operating system reacts to events, such as interrupts, and allocates CPU usage to the processing of tasks that arise from these events. The simplest operating system works on a 'first come, first served' basis. Thus, the tasks are handled in the order that they occur and the processing for one task is allowed to finish before processing for the next task starts.

Most operating systems allow multiple tasks to seemingly run concurrently. In fact, the CPU may be able to process only one task at a time, but shares its time between the tasks by switching between them, giving the illusion of concurrency - this is called multi-tasking. A conventional operating system may use a simple scheduling algorithm to allocate CPU time to multiple tasks. A common scheduler is the round-robin algorithm, which allocates a time-slice to each task in cyclic order. In such a 'time slicing' scheme, all tasks are allocated equal slices of time.

In a real-time application, some tasks may have strict deadlines and must not be delayed by other tasks. For example, in a digital sound-recording application, a sampling task must be completed within a certain time of it starting. This task will be more important than a simultaneous task to compress the previous sample. This suggests the need for precedence to ensure that the sampling task is allowed to meet its deadline.

An RTOS accommodates this idea of precedence, as follows:

- Uses priorities assigned to the different tasks to be performed
- Schedules CPU usage such that higher priority tasks are handled before lower priority tasks
- May allow the processing of a task to be temporarily suspended while a newer, higher priority task is handled (pre-emptive RTOS only)

Therefore, an RTOS multi-tasks without time-slicing, but instead allocates CPU time according to the priorities of the current tasks.

2.2 Introduction to the JenOS RTOS

The RTOS within JenOS offers both immediacy and flexibility in scheduling real-time tasks with different priorities. It is a pre-emptive RTOS, so switches CPU usage to the highest priority task, but also provides a 'mutex' feature that allows execution of a critical section of the currently running task to be completed before any switch occurs.

The RTOS identifies the following set of high-level priorities which apply to four general classes of event:

1. Uncontrolled interrupts corresponding to events external to the operating system (not handled by RTOS)
2. Operating system housekeeping tasks
3. Controlled interrupts corresponding to events internal to the operating system (handled by RTOS)
4. User tasks (handled by RTOS)

Therefore, use of the RTOS in the user application is only concerned with the last two categories - controlled interrupts and user tasks. There is a separate set of user-defined priorities within each of these categories, but a controlled interrupt will always take priority over a user task.

The JenOS RTOS provides an 'idle task', which is executed when there are no other tasks to be run.

The RTOS is started using the function **OS_vStart()**. Following a warm start with memory held, it can be re-started using the function **OS_vRestart()**.

2.3 RTOS Configuration

The JenOS RTOS is configured at build time. This ensures that provision can be made for all types of application.

RTOS configuration is performed during application development using the JenOS Configuration Editor. This tool allows OS resources to be easily assigned and configured through a graphical interface. The tool then generates the necessary OS configuration files to feed into the application build process. The JenOS Configuration Editor and the build process are outlined in [Chapter 15](#) before a more detailed account of the JenOS Configuration Editor is presented in [Chapter 14](#).

2.4 RTOS Concepts and Features



This section details the following concepts and features of the JenOS RTOS:

- **User tasks:** The user application must be divided into user tasks, described in [Section 2.4.1](#).
- **Interrupt Service Routines (ISRs):** The user application must define ISRs to deal with controlled interrupts (those handled by the RTOS) - see [Section 2.4.2](#).
- **Priorities and scheduling:** Priorities must be assigned to tasks and ISRs statically by the system developer, using the method described in [Section 2.4.3](#).
- **States and transitions:** The possible states of user tasks and ISRs, and the transitions between these states, are described in [Section 2.4.5](#).
- **Scheduled activities:** Individual activities within tasks/ISRs can be scheduled to start at certain times, as described in [Section 2.4.6](#).
- **Mutually exclusive access (mutex):** The mutex feature allows the priority system to be effectively over-ridden when tasks are competing for a shared resource and task switching should be avoided - see [Section 2.4.7](#).
- **Inter-task communication:** Communications between tasks can be managed using message queues, as described in [Section 2.4.8](#).

2.4.1 User Tasks



A user application can be conveniently subdivided into sections that are executed according to their real-time requirements. These sections can be implemented as tasks, where a task provides the framework for the execution of functions. The JenOS RTOS defines a task as consisting of a main C function and all its sub-functions. The main function of a task can only be invoked by the operating system (no user function is allowed to call it).

A task is defined using the RTOS macro **OS_TASK()**. The task must be given a reference handle, which is assigned in the RTOS configuration. The handle of the currently running task can be obtained using the function **OS_eGetCurrentTask()**.

The set of tasks in an application must be assigned priorities, used to determine the allocation of CPU time in a multi-tasking environment. This is described further in [Section 2.4.3](#).



Caution: To allow the RTOS to operate correctly, user tasks must be designed so that they do not block.

2.4.2 Interrupt Service Routines (ISRs)

The user application may define Interrupt Service Routines (ISRs) to handle controlled interrupts. These are interrupts that are managed by the RTOS via its control mechanisms (e.g. mutex) and API calls. Each interrupt is assigned an ISR, which is invoked by the operating system when the interrupt occurs.

An ISR is defined using the RTOS macro **OS_ISR()**. The ISR must be given a reference handle, which is assigned in the RTOS configuration.

The set of ISRs in an application must be assigned priorities, used to determine the allocation of CPU time when multiple interrupts are pending. This is described further in [Section 2.4.3](#).



Note 1: The RTOS API contains functions to enable/disable interrupts. You can enable/disable all interrupts or just the controlled interrupts that the RTOS handles. For details of these functions, refer to [Section 8.2.3](#).

Note 2: The RTOS cannot clear interrupts and it is the responsibility of the application to do this - refer to [Appendix B](#).



Caution: The RTOS uses the Programmable Interrupt Controller (PIC) of the JN516x device. When using the RTOS, you must therefore not use the PIC directly.

2.4.3 Priorities and Scheduling

The user tasks and ISRs can both be prioritised for CPU allocation, each being assigned its own set of priorities. These priorities are statically set in the RTOS configuration by the system developer at build time. The sections below describe the priority system and the related topic of CPU scheduling.

Assigning Priorities

Within each of the two task categories (user task and ISR), the following method should be used to assign priorities:

1. Rank the tasks in order of importance with respect to their deadlines (i.e. deadline-monotonic analysis).
2. Assign priorities to the tasks as integer values, starting with 1 for the lowest priority task and incrementing the assigned value until all tasks have a priority.

Therefore, the assigned priorities are 1, 2,... n, where the higher the value, the higher the priority. The priority assigned to an ISR is referred to as its Interrupt Priority Level (IPL).

CPU Scheduling

In terms of scheduling, among the user tasks/ISRs waiting for CPU time, the one with the highest priority will be handled next. However, an ISR will always take precedence over a user task (the lowest ISR priority takes precedence over the highest user task priority). If an event then occurs which has a corresponding user task/ISR with higher priority than the one currently being executed, processing will switch to the higher priority process, with the existing process being temporarily suspended. In this case, the higher priority process is said to pre-empt the lower priority process.

Co-operative Tasks

Tasks can be grouped as co-operative tasks. A task in a Co-operative Task Group will not pre-empt another task from the same group, irrespective of their relative priorities. A running co-operative task temporarily takes the highest priority level assigned to the members of its group. However, any task from within the group can pre-empt or be pre-empted by a task from outside the group.

2.4.4 Task/ISR States

At any one time, a user task or ISR can be in one of three states:

- **Running:** In the running state, the CPU is executing the functions/instructions that make up the task/ISR. Only one task/ISR can be in this state at any one time.
- **Pending:** In the pending state, the task/ISR has been activated (and may have been previously run), and is waiting to become the highest priority task/ISR. When it does, it will be switched into the running state.
- **Dormant:** In the dormant state, a task/ISR is awaiting activation (which will put it in the pending state) or has already been run.

Transitions between these states are further described in [Section 2.4.5](#).



A user task can be activated (moved from the dormant state to pending state) in the application code. The task is subsequently managed automatically by the RTOS. Activation of a user task is performed using the **OS_eActivateTask()** function. When this function is called, the activation counter for the task is incremented. It is possible to call this function even when the user task is already in the pending state - the activation counter keeps a record of the number of times the user task must be run before it can return to the dormant state.



Note: A user task must initially be activated from the main task but, once in the running state, can activate itself using **OS_eActivateTask()**, in which case its activation counter will be incremented.

An ISR is automatically activated when the corresponding interrupt is generated.

Once activated, a user task/ISR will be executed when it becomes the highest priority pending user task/ISR - see [Section 2.4.3](#).

2.4.5 State Transitions

As detailed in [Section 2.4.4](#), at any one time, a user task or ISR can be in the running, pending or dormant state. The possible transitions between these states are described below.

1. **Activate:** A user task must first be activated (using `eOS_ActivateTask()`) to move it from the dormant state to the pending state. Note that ISRs do not need to be explicitly activated.
2. **Start:** When a pending task/ISR becomes the highest priority task/ISR, it is automatically started by the RTOS, which moves it to the running state.
3. **Pre-empt:** A running task/ISR can be pre-empted by a higher priority task/ISR by suspending execution and moving it back to the pending state. Here it will stay until it becomes the highest priority task/ISR again. It will then be moved back in the running state and execution will be resumed from where it left off.
4. **Complete:** Once execution of a task/ISR has completed, it is automatically moved from the running state to the dormant state.

The possible state transitions are illustrated in the figure below.




Figure 2: State Transitions for User Task or ISR



Note: An ISR always takes precedence over a user task. Therefore, a user task can never pre-empt an ISR.

2.4.6 Activity Scheduling (using Software Timers)

Within a user task or ISR, it may be necessary to schedule an activity to occur in the future or on a regular basis (i.e. periodically). For example, during a 'button detect' task, we may need to sample the state of a button every 10 ms. For such scheduled activities, the RTOS provides software timers and a number of timer functions.

A software timer is derived from a source counter, which can be either of:

- A hardware timer, such as the on-chip tick timer or an external timer
- Another software timer

The software timers and source counter required by the application are pre-defined in the RTOS configuration using the JenOS Configuration Editor. Each software timer has a handle, which is also assigned in the RTOS configuration.

A software timer is started using the function **OS_eStartSWTimer()**. As part of this function call, you must specify the number of ticks (of the source counter) before the software timer expires. This value is entered into a 'compare register' for the source counter. The counter increments and when the count reaches the value in the compare register, an interrupt may be generated and an ISR invoked, where this ISR is pre-defined in the RTOS configuration for the application. Data for this ISR can be specified through the function **OS_eStartSWTimer()**. The ISR must call the function **OS_ExpireSWTimers()**.



Note: Several software timers may be based on the same source counter. Their different expiration times are handled in a relay, by passing ownership of the source counter from one timer to the next - when one timer expires, the source counter's compare register is updated with the number of ticks until the next timer expires, and so on.

The function **OS_ExpireSWTimers()** sets the software timer status to 'expired' and checks whether there are any other pending software timers for the same source counter:

- If there are no pending software timers, the function disables the source counter.
- If there is at least one pending software timer, the function updates the source counter's compare register with the required number of ticks (until the next software timer expires).

When a software timer expires, if the same timer is to be re-started immediately (possibly with a different timed period), the function **OS_eContinueSWTimer()** can be used to re-start the timer without loss of continuity - there will be no break between the last expiry point and the new timer run, provided that this function is called before the next counter period starts.



Note: If the tick timer is used as the source counter and the maximum count of the tick timer is T (before the tick timer wraps around), there must be no more than $T/2$ ticks between consecutive software timer expiry events (e.g. if the tick timer wraps around every 60 seconds, a software timer must expire every 30 seconds or less).

Once started, a software timer can be prematurely stopped at any time using the function **OS_eStopSWTimer()**.



Caution: To allow the JN516x device to enter sleep mode, no software timers should be active. Any running software timers must first be stopped and any expired timers must be de-activated. Both can be achieved using the function **OS_eStopSWTimer()**, which must be called individually for each running and expired timer.

Callback Functions and Macros

The software timer functions mentioned above use callback functions that are user-defined, e.g. to enable/disable the hardware counter. The callback functions must be defined using macros provided in the RTOS module. These software timer functions are listed below with their associated callback functions and macros:

- **OS_eStartSWTimer():**
 - Calls the user-defined 'hardware counter enable' callback function, defined using the macro **OS_HWCOUNTER_ENABLE_CALLBACK()**
 - Calls the user-defined 'hardware counter get' callback function, defined using the macro **OS_HWCOUNTER_GET_CALLBACK()**
 - Calls the user-defined 'hardware counter set' callback function, defined using the macro **OS_HWCOUNTER_SET_CALLBACK()**
- **OS_eStopSWTimer():**
 - Calls the user-defined 'hardware counter disable' callback function, defined using the macro **OS_HWCOUNTER_DISABLE_CALLBACK()**
- **OS_eContinueSWTimer():**
 - Calls the user-defined 'hardware counter set' callback function, defined using the macro **OS_HWCOUNTER_SET_CALLBACK()**
- **OS_eExpireSWTimers():**
 - Calls the user-defined 'software timer expired' callback function, defined using the macro **OS_SWTIMER_CALLBACK()**
 - Calls the user-defined 'hardware counter set' callback function, defined using the macro **OS_HWCOUNTER_SET_CALLBACK()**

For full details of the above macros and callback functions, refer to [Chapter 8](#).

2.4.7 Mutual Exclusion (Mutex)

The pre-emptive task scheduling of an RTOS can be problematic when two competing tasks need to access the same resource. In this case, switching the running task to a higher priority task, where both tasks can control a shared resource, may lead to undesirable results. For example, if both tasks need to access the same memory block and the running task is already writing data to memory, it is not desirable for this process to be suspended and for another task to start writing to the memory block. In such circumstances, it is important to allow the (lower priority) running task to complete its access before the higher priority task starts.

The JenOS RTOS provides a mutual exclusion (mutex) feature to prevent task switching during execution of a critical part of a user task or ISR. The critical section of code within the user task/ISR must be delimited with the functions **OS_eEnterCriticalSection()** and **OS_eExitCriticalSection()**.

For this feature, RTOS implements a Priority Inheritance Protocol, which works by temporarily raising the priority of the running task during the critical section of code (the task's priority is returned to its normal value outside the critical section). A running task and pending task can only be managed in this way if they belong to the same mutex group. Each mutex group is given a unique handle and the user tasks/ISRs in a given mutex group are pre-defined in the RTOS configuration for the application. During execution of the critical section, the priority of the running task is changed to the highest possible priority of the tasks within the same mutex group.

Use of the mutex feature in handling a critical section of code is illustrated in [Figure 3](#).

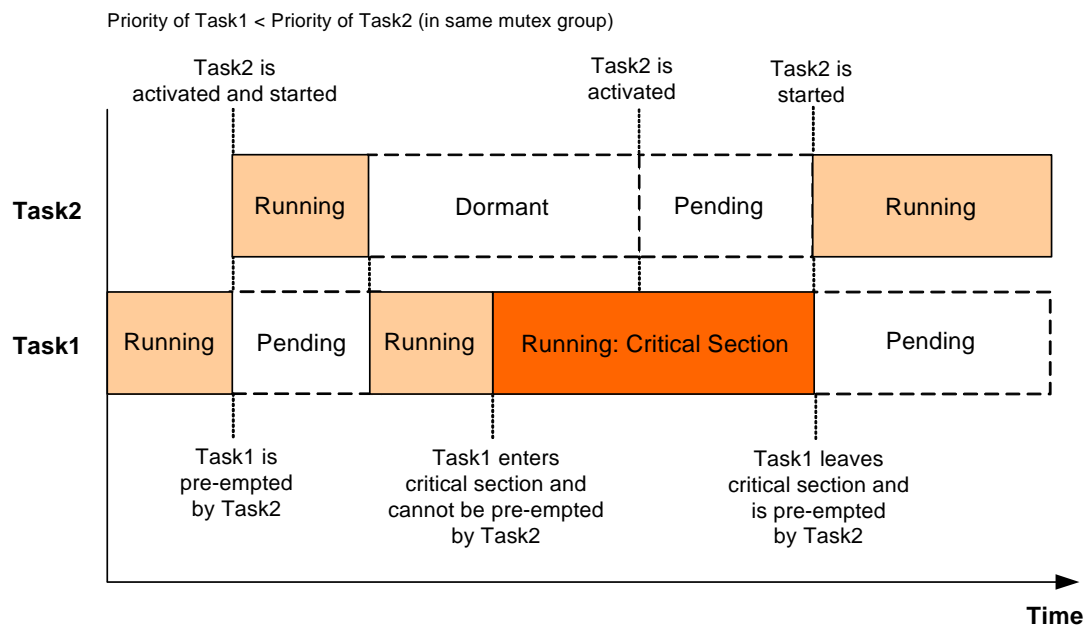


Figure 3: Mutex Handling of Critical Section

2.4.8 Inter-task Communication (using Messages)

A user task may often need to communicate with another user task, or an ISR may need to communicate with a user task. This communication can be implemented using messages. The message types required by the application are pre-defined in the header file **os_msg_types.h**. The identity of the destination task or ISR for a message type, and whether the message type is queued, are specified in the RTOS configuration for the application.

The function **OS_ePostMessage()** is used to send a message from a user task/ISR. While this user task/ISR is running, the destination task will be dormant or pending, and the message cannot be delivered immediately.

- If the message has associated user data, the message can be placed in a queue - this queue is specifically associated with the message type and the destination task, and is set up in the RTOS configuration. Alternatively, the message may be unqueued.
- If the message has no associated user data, the message is not queued.

If a message is not queued, it replaces any previous message of the same type that was waiting to be delivered to the destination task.

The destination task will run as soon as it becomes the highest priority task. Other messages may arrive in the message queue before the task runs. Once run, the task will collect the messages waiting in its queue. A message is collected using the function **OS_eCollectMessage()**, which can be called repeatedly until the queue is empty.

The function **OS_eGetMessageStatus()** is also provided which allows the destination task to determine if there are valid messages in its queue.

The above process of message sending and collection, for a queued message type, is illustrated in [Figure 4](#) below.

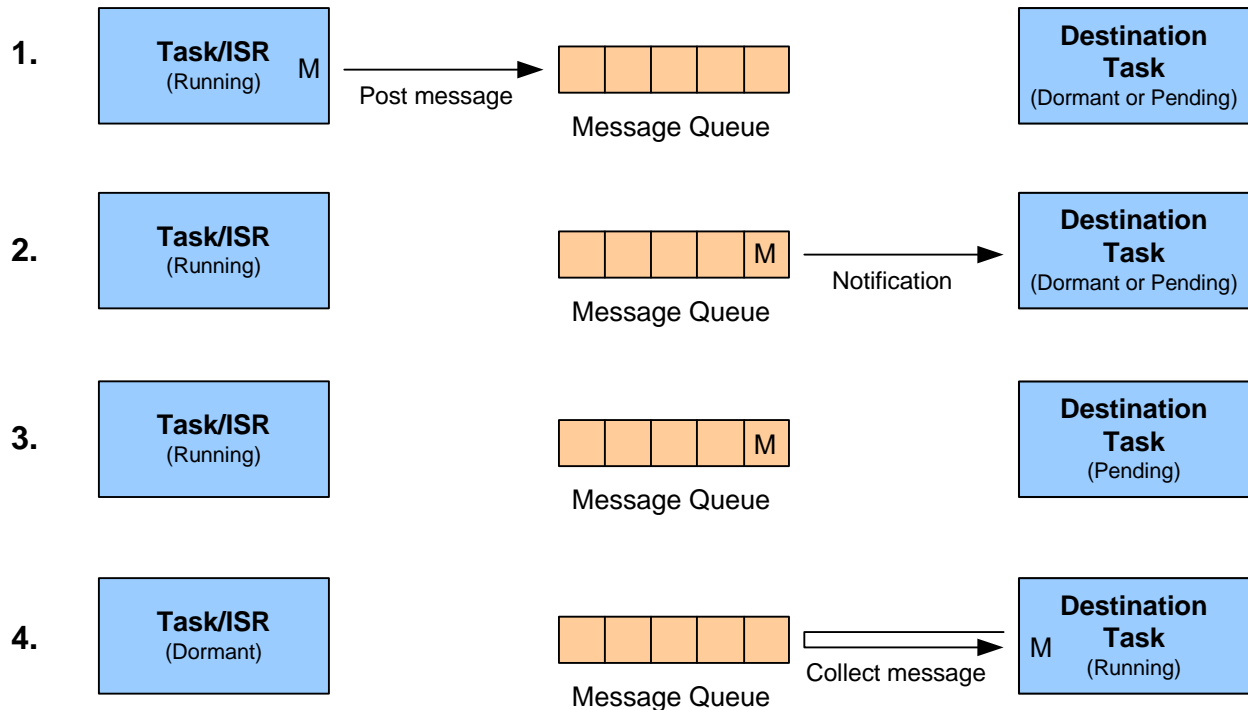


Figure 4: Message Queuing

When sending a message of a given type using **OS_ePostMessage()**, two alternative options are available:

- Invoke a callback function which performs user actions
- Activate the associated destination task through a notification - that is, to increment the task's activation counter and, if the task is not already pending, move it from the dormant to pending state.

The required send option is set in the RTOS configuration.

Data structures are generated at compile time detailing message senders, receivers and queue sizes, to ensure only the minimum resources necessary are allocated. Each task lists the message types that it can transmit and receive.



Note: The above messaging system is used by the NXP ZigBee PRO software to send and receive responses, notifications and events. These messages are automatically generated and sent by the stack software, but the application is responsible for collecting them using **OS_eCollectMessage()**.

2.5 OS Error Callback Function



Caution: Errors affecting the OS can cause system instability. It is recommended that all applications register an OS error handler and enable strict error checks.

Errors can be detected by testing the return codes from the calls to OS functions. However, this requires application code to test the return code from every OS call.

Registering an error callback function provides a robust alternative to checking the return codes. This function is invoked whenever an OS function returns an error. The function is registered as a parameter of **OS_vStart()**. The error callback option must also be enabled in the JenOS Configuration Editor.

Many OS errors leave the scheduler in an undefined state. For example, nesting a mutex by calling **OS_eEnterCriticalSection()** twice before calling **OS_eExitCriticalSection()** causes a **OS_E_BAD_NESTING** error. Once an error of this nature has occurred, the OS scheduler is in an undefined state.

2.5.1 Strict Error Checks

The OS scheduler will enter an undefined state if there are inconsistencies between the OS configuration diagram (in the JenOS Configuration Editor) and the application code. A strict error check option can be enabled in the JenOS Configuration Editor to check for inconsistencies between the OS configuration diagram and the software. The strict mode has a slight overhead in code space and execution time but it is good practice to enable strict checking where possible. For example, calling **OS_eEnterCriticalSection()** from a task which is not in the group for the mutex will generate **OS_E_CURRENT_TASK_NOT_A_MUTEX_MEMBER** with strict checking enabled. If strict checks were not enabled, the scheduler operation would be undefined and the system may become unstable.

2.5.2 Handling OS Errors

During testing, an application's error callback function should stop the application with a stack dump and the error should be fixed. The OS passes two parameters to the error callback - the status code of the error and a pointer to the handle which caused the error. These parameters should be printed out to help determine the cause of the error.

In production code, the device must be re-started from cold by calling **vAHI_SwReset()**. Data in the PDM module does not normally need to be erased, so the device can rejoin a ZigBee PRO network with existing security material.

The error callback function will be called on some non-fatal errors. Depending on the application design, the following errors can be ignored by the error callback function:

- OS_E_QUEUE_EMPTY
- OS_E_SWTIMER_STOPPED
- OS_E_SWTIMER_EXPIRED
- OS_E_SWTIMER_RUNNING

3. Persistent Data Manager (PDM) for Flash Memory

This chapter describes the Persistent Data Manager (PDM) module which handles the storage of stack context and application data in Flash memory which is external to the JN516x device. This edition of the PDM is supplied in the JN516x ZigBee Smart Energy SDK (JN-SW-4064).

For the later edition of the PDM (supplied in other SDKs) that supports JN516x internal EEPROM, refer to [Chapter 4](#).



Tip 1: In this chapter, the storage medium for persisted data is referred to as Non-Volatile Memory (NVM) but in practice it is SPI-connected external Flash memory.

Tip 2: In this chapter, a cold start refers to either a first-time start or a re-start without memory (RAM) held. A warm start refers to a re-start with memory held (for example following sleep with memory held).

3.1 Overview

Data needed for the operation of a network node is normally stored in on-chip RAM. This includes data that may evolve during node operation, e.g. context data for the network stack and application data. This data is only maintained in memory while the node is powered and will be lost during an interruption to the power supply (e.g. power failure or battery replacement).

In order for the node to recover from a power interruption with continuity of service, provision must be made for storing a back-up of the operational data in NVM, such as Flash memory. This data can then be recovered during a re-boot following power loss, allowing the node to resume its role in the network.

The storage and recovery of operational data can be handled using the Persistent Data Manager (PDM) module, as described in the rest of this chapter. The NVM device used by the PDM is expected to be external SPI-connected Flash memory.



Caution: When using the PDM, do not use the JN516x Integrated Peripherals API to interact with the Flash memory device connected to the JN516x chip.

An overview of the use of the PDM API functions in application code is presented in [Figure 5](#) below.

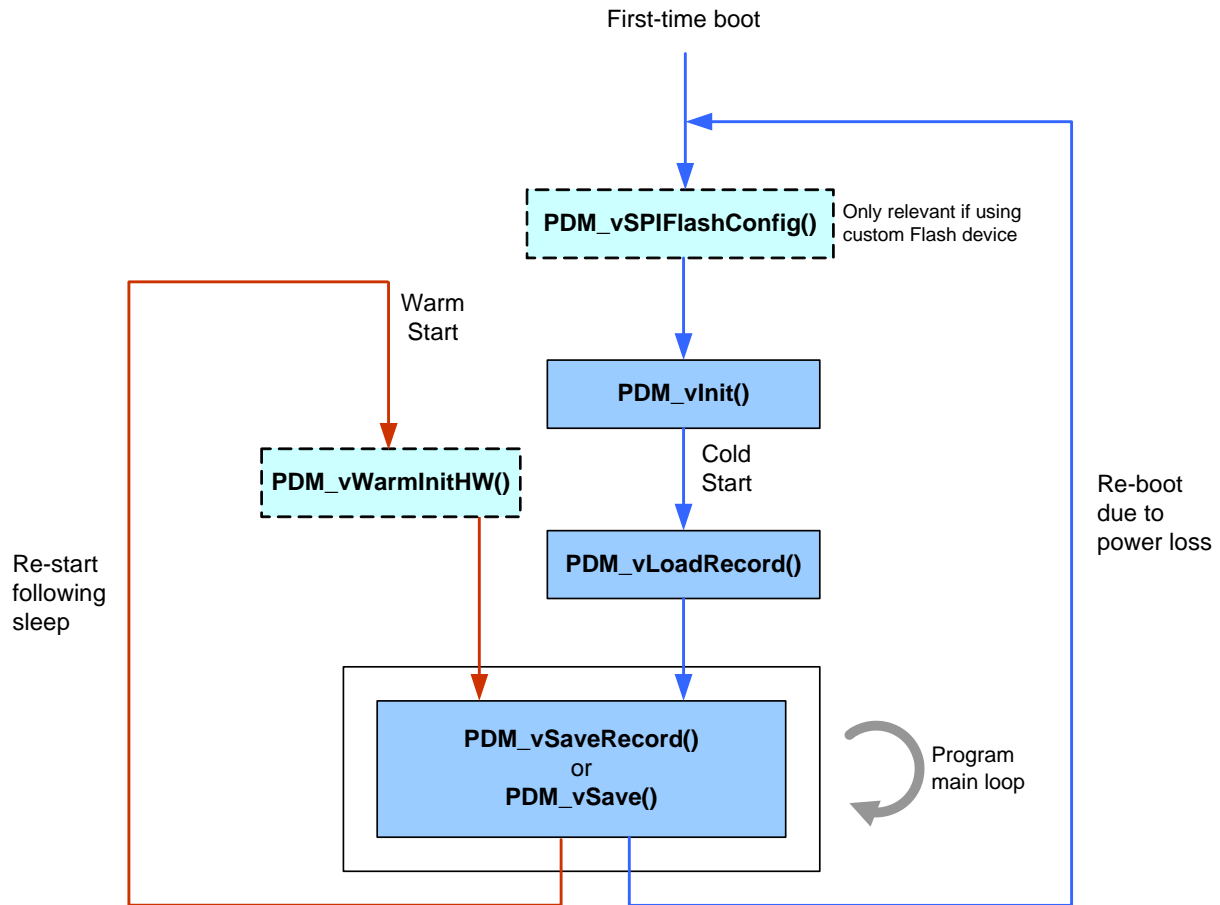


Figure 5: PDM Overview

3.2 Initialising the PDM

The initialisation of the PDM depends on the type of NVM device used. For a cold start (a first-time start or a re-start following power loss), the initialisation is as follows:

1. If NVM is a custom device, call the function **PDM_vSPIFlashConfig()**. This function requires you to specify a pointer to a set of custom functions that will be used by the PDM to interact with the NVM device (e.g. read and write functions). There is no need to call this function if using a supported NVM device.
2. Call the function **PDM_vInit()**. If using a supported NVM device, this function will auto-detect the device type. If required, optional mutexes can be specified through this function - for information on these mutexes, refer to [Section 3.7](#). A security key can also be specified which will be used by the PDM module to encrypt data saved to NVM and to decrypt the data when read from NVM (the key can be specified by the application or obtained from eFuse) - this security will be automatically applied to stack context data but must be explicitly enabled for application data (see [Section 3.4](#)).

Both of the above functions require you to identify and specify the size of the NVM sectors to be managed by the PDM.

For a warm start (following sleep), there is no need to call **PDM_vSPIFlashConfig()** or **PDM_vInit()** but the function **PDM_vWarmInitHW()** must be called.

3.3 Data Storage in NVM

Data is stored in NVM in terms of 'records', where a record is an area of NVM used to store data and associated housekeeping information. Each NVM record corresponds to a data buffer in RAM, where the NVM record is used to back up the RAM buffer. Any number of NVM records of different lengths can be created, provided that they do not exceed the NVM capacity.

NVM records are created automatically for stack context data and by the application (as indicated in [Section 3.4](#)) for application data.

The stack context data which is stored in NVM includes the following:

- Application layer data:
 - AIB members, such as the EPID and ZDO state
 - Group Address table
 - Binding table
 - Application key-pair descriptor
 - Trust Centre device table
- Network layer data:
 - NIB members, such as PAN ID and radio channel
 - Neighbour table
 - Network keys
 - Address Map table

The PDM manages a Cyclic Redundancy Code (CRC) for each record and writes the CRC only on a complete record write. Therefore, if the CRC is not present then the data will not be recoverable.

3.4 Recovering Data from NVM

Application data and stack context data are loaded from NVM to RAM as described below.

Application Data

During a cold start, after the PDM module has been initialised (see [Section 3.2](#)), the function **PDM_eLoadRecord()** must be called for each record of application data in NVM and its corresponding data buffer in RAM. This function requires a descriptor to be specified for the NVM record, where this record descriptor is held in RAM.



Note: The function **PDM_eLoadRecord()** must be called before calling any function which automatically saves application data to NVM - for example, before calling **ZPS_eAplAflnit()** to initialise the ZigBee PRO stack and **ZPS_vAplSecSetInitialSecurityState()** to initialise the ZigBee security state on the node.

Depending on the type of cold start, **PDM_eLoadRecord()** will do one of two things:

- During a first-time cold start, the function creates an NVM record that corresponds to the specified record descriptor (but stores no data in NVM). This NVM record will be used to back up application data from the specified buffer in RAM (see [Section 3.5](#)). Initial data values must be provided in this RAM buffer.
- During any subsequent cold start, the function loads data from the specified NVM record to the associated RAM buffer, thus recovering application data previously saved to NVM.



Caution: ***PDM_eLoadRecord()** must not be called during a warm start, otherwise data held in RAM will be overwritten by possibly older data from NVM.*

PDM_eLoadRecord() also allows security to be enabled for the application data record. If security is enabled, data saved to the NVM record will be encrypted using the key specified through **PDM_vlnit()** and will therefore be decrypted using the same key when read from the record.

Stack Context Data

The function **PDM_eLoadRecord()**, described above, is not used for records of stack context data. Loading this data from NVM to RAM is handled automatically by the stack (provided that the PDM has been initialised). When saved to NVM, the stack context data is automatically encrypted using the security key specified through **PDM_vlnit()** and is therefore decrypted using the same key when read from NVM.

3.5 Saving Data to NVM

Application data and stack context data are saved from RAM to NVM as described below.



Note: If, during a data save, NVM needs to be defragmented and purged, this will be performed automatically resulting in all records being re-saved.

Application Data

Once you have application data in RAM to be backed up, you can save the RAM data associated with an individual record in NVM using the function **PDM_vSaveRecord()**. Alternatively, you can save the RAM data corresponding to all records in NVM using the function **PDM_vSave()**. In both cases, the saved application data will be encrypted if security was enabled for the corresponding NVM record through the function **PDM_eLoadRecord()** (see [Section 3.4](#)). You should save data to NVM when important changes have been made to the data in RAM.

Stack Context Data

The function **PDM_vSave()** can be used to save all records of stack context data as well as all records of application data. The saved stack context data is encrypted using the security key specified when the PDM module was initialised using **PDM_vInit()** (see [Section 3.2](#)). Note that the NXP ZigBee PRO stack automatically saves its own context data from RAM to NVM (encrypted as detailed above) when certain data items change.

3.6 Deleting Data in NVM

An individual record of application data in NVM can be deleted using the function **PDM_vDeleteRecord()**. Alternatively, all records (application data and stack context data) in NVM can be deleted using the function **PDM_vDelete()**. Note that when a record is deleted in NVM, the corresponding data buffer in RAM is not deleted.



Caution: You are not recommended to delete records of ZigBee PRO stack context data by calling **PDM_vDelete()** before a rejoin of the same secured network. If these records are deleted, data sent by the node after the rejoin will be rejected by the destination node since the frame counter has been reset on the source node. For more information and advice, refer to the “Application Design Notes” appendix in the ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101).

3.7 Mutexes in PDM

The PDM module can use optional mutexes, as follows:

- PDM functions are not re-entrant and a mutex can be optionally used to prevent concurrent PDM function calls - if enabled, the mutex is applied automatically during a PDM function call.
- External NVM is accessed via the SPI bus of the JN516x device, but this bus may also be used to access other resources. The PDM module can optionally use a mutex to prevent concurrent accesses to the SPI bus - if enabled, the mutex is applied automatically during a PDM access to NVM.

If required, the above mutexes can be specified when the PDM module is initialised using the function **PDM_vInit()** - see [Section 3.2](#). The user task must also be linked to the relevant mutex in the JenOS Configuration Editor - see [Section 14.7](#).

The principles of a mutex are described in [Section 2.4.7](#). All tasks that use a mutex must be connected to the mutex in the OS configuration diagram. The ZigBee PRO stack makes PDM calls. Therefore, any task that makes ZigBee PRO function calls must be connected to the PDM mutexes.

3.8 Ensuring Consistency of PDM Records

The data in the PDM may differ in structure from that expected by the application. The structures stored by the ZigBee PRO libraries can change due to altering table sizes in the ZPS Configuration Editor, as well as between releases of the ZigBee PRO stack libraries, and inconsistency can occur.

When an Over-The-Air (OTA) software update is performed, the PDM data is not erased. This is normally a benefit because it allows the application to rejoin the network. However, if any of the PDM structures change, a factory reset must be performed by calling **PDM_vDelete()**.

Applications normally contain a way to perform a factory reset of the PDM module - for example, by calling **PDM_vDelete()** if a button is held down during reset.

The application can automatically check for PDM consistency by storing an application-specific 'magic number' in a record. A new magic number should be used if the application software or ZigBee PRO libraries PDM usage is inconsistent with the previous version of the software. Immediately after calling **PDM_vInit()**, the application should call **PDM_eLoadRecord()**. If the magic number does not match, the application should call **PDM_vDelete()** to erase all records before attempting to start the ZigBee PRO stack. If the call to **PDM_eLoadRecord()** indicates that the record has not been found, the application should also call **PDM_vDelete()** because another application may have been running that does not use the same record ID but has written inconsistent ZigBee PRO records to the PDM module.

4. Persistent Data Manager (PDM) for EEPROM

This chapter describes the Persistent Data Manager (PDM) module which handles the storage of stack context data and application data in JN516x internal EEPROM. This edition of the PDM is supplied in the following SDKs:

- JN-SW-4168: JN516x ZigBee Light Link and Home Automation SDK
- JN-SW-4165: JN516x JenNet-IP SDK
- JN-SW-4163: JN516x IEEE802.15.4 SDK

For the earlier edition of the PDM (supplied in other SDKs) that supports external Flash memory devices, refer to [Chapter 3](#).



Note : The PDM functions referenced in this chapter are detailed in [Chapter 10](#).



Tip: In this chapter, a cold start refers to either a first-time start or a re-start without memory (RAM) held. A warm start refers to a re-start with memory held (for example following sleep with memory held).

4.1 Overview

If the data needed for the operation of a network node is stored only in on-chip RAM, this data is maintained in memory only while the node is powered and will be lost during an interruption to the power supply (e.g. power failure or battery replacement). This data includes context data for the network stack and application data.

In order for the node to recover from a power interruption with continuity of service, provision must be made for storing essential operational data in Non-Volatile Memory (NVM), such as EEPROM. This data can then be recovered during a re-boot following power loss, allowing the node to resume its role in the network.

The storage and recovery of operational data in JN516x EEPROM can be handled using the Persistent Data Manager (PDM) module, as described in the rest of this chapter, which covers the following topics:

- Initialising the PDM module - see [Section 4.2](#)
- Managing data in EEPROM - see [Section 4.3](#)
- Storing counters in EEPROM - see [Section 4.4](#)
- PDM features including mutexes, EEPROM wear counts and event handling - see [Section 4.5](#)

[Section 4.2](#) also provides information on using the PDM without the JenOS RTOS.



Note: Unlike the earlier version of the PDM (described in [Chapter 3](#)), this version does not base storage records on descriptors. The removal of descriptors simplifies the operation of the PDM and reduces the EEPROM space needed to hold a given amount of data.

4.2 Initialising the PDM and Building a File System

The PDM module must be initialised by the application following a cold or warm start, irrespective of the PDM functionality used (e.g. context data storage or counter implementation). PDM initialisation is performed using the function **PDM_eInitialise()**.

This function requires the following information to be specified:

- The number of EEPROM segments to be used by PDM (a zero value means use all segments)
- An optional mutex in order to serialise PDM function calls and prevent concurrent calls - for information on this mutex, refer to [Section 4.5](#)

Once the **PDM_eInitialise()** function has been called, the PDM module builds a file system in RAM containing information about the segments that it manages in EEPROM. The PDM reads the header data from each EEPROM segment and builds the file system.

The file system allows the PDM to perform efficient searches when operating on data, track the occupation of all the segments in the EEPROM and keep track of the number of segments available for data allocation at any time. It also helps to even out the wear across EEPROM segments - for more information on EEPROM segment wear, refer to [Section 4.5.4](#).

Using PDM without the JenOS RTOS (IEEE802.15.4 and JenNet-IP only)

The JenOS RTOS is not used in IEEE802.15.4 and JenNet-IP applications, but the PDM module can still be used.

IEEE802.15.4

To use the PDM in applications developed using the IEEE802.15.4 SDK (JN-SW-4163), the flag **PDM_NO_RTOS** must be defined in the makefile, as follows:

```
CFLAGS += -DPDM_NO_RTOS
```

The serialisation mutex cannot be used in this case and the relevant parameter is removed from the **PDM_eInitialise()** function.

JenNet-IP

When the PDM is used in applications developed using the JenNet-IP SDK (JN-SW-4165), no makefile modifications need to be made. However, the serialisation mutex is always implemented and a non-zero value must be passed to the relevant parameter in the **PDM_eInitialise()** function.

4.3 Managing Data in EEPROM

This section describes use of the PDM module to persist data in EEPROM in order to provide continuity of service when the JN516x device resumes operation after a cold start or a warm start without memory held.

Data is stored in EEPROM in terms of 'records'. A record occupies at least one EEPROM segment but may be larger than a segment and occupy multiple segments. Any number of records of different lengths can be created, provided that they do not exceed the EEPROM capacity. The records are created automatically for stack context data and by the application (as indicated in [Section 4.3.1](#)) for application data. Each record is identified by a unique 16-bit value which is assigned when the record is created - for application data, this identifier is user-defined.

The stack context data which is stored in EEPROM includes the following:

- Application layer data:
 - AIB members, such as the EPID and ZDO state
 - Group Address table
 - Binding table
 - Application key-pair descriptor
 - Trust Centre device table
- Network layer data:
 - NIB members, such as PAN ID and radio channel
 - Neighbour table
 - Network keys
 - Address Map table

On performing a JN516x cold start or warm start without RAM held, the PDM must be initialised in the application as described in [Section 4.2](#).

- If this is the first ever cold start, there will be no stack context data or application data preserved in the EEPROM.
- If it is a cold or warm start following previous use (such as after a reset), there should be stack context data and application data preserved in the EEPROM.

On start-up, the PDM builds a file system in RAM and scans the EEPROM for valid data. If any data is found, it is incorporated in the file system.

The PDM saves a Cyclic Redundancy Code (CRC) for each segment of a record. Any failure will result in the data being unrecoverable and the record becoming invalid.

Saving, recovering and deleting application data in EEPROM are described in the sub-sections below.

4.3.1 Saving Data to EEPROM

Application data and stack context data are saved from RAM to EEPROM as described below.



Note: During a data save, if the EEPROM needs to be defragmented and purged, this will be performed automatically resulting in all records being re-saved.

Application Data

You should save application data to EEPROM when important changes have been made to the data in RAM. Application data in RAM can be saved to an individual record in EEPROM using the function **PDM_eSaveRecordData()**. A buffer of data in RAM is saved to a single record in EEPROM (a record may span multiple EEPROM segments).

The first time that a record is saved using **PDM_eSaveRecordData()**, the record is created and the data is written in its entirety, provided there are enough free EEPROM segments to hold the data (you can first find out how many segments are available using the function **PDM_u8GetSegmentCapacity()**). When a record is first created, a unique 16-bit identifier must be assigned to the record by the application - this identifier is subsequently used to reference the record. The value used must not clash with those used by the NXP libraries - the ZigBee PRO stack libraries use values above 0x8000 and the JenNet-IP libraries use values between 0x3000 and 0x3007.

Subsequently, in performing a re-save to the same record (specified by its 16-bit identifier), the original EEPROM segments associated with the record will be overwritten but only the segment(s) containing data changes will be altered (if no data has changed, no write will be performed). This method of only making incremental saves improves the occupancy level of the size-restricted EEPROM.

If a save fails, the function **PDM_eSaveRecordData()** will return the code **PDM_E_STATUS_NOT_SAVED**. Alternatively, the callback event **E_PDM_SYSTEM_EVENT_DESCRIPTOR_SAVE_FAILED** can be used to notify the application of a save failure - this requires a PDM callback function to have been registered using the function **PDM_vRegisterSystemCallback()**, as described in [Section 4.5.2](#).

Stack Context Data

The NXP ZigBee PRO stack automatically saves its own context data from RAM to EEPROM when certain data items change. This data will not be encrypted.

4.3.2 Recovering Data from EEPROM

Application data and stack context data are loaded from the EEPROM to RAM as described below.

Application Data

Application data records in EEPROM can be read by the application using the function **PDM_eReadDataFromRecord()**. The record to be read is specified using its 16-bit identifier and a data buffer in RAM must also be specified in which the read data will be stored.

Before calling **PDM_eReadDataFromRecord()**, it may be useful to call the function **PDM_bDoesDataExist()** in order to determine whether a record with the specified identifier exists in the EEPROM and, if it does, to obtain its size and therefore the length of the required RAM buffer.

During a cold start or a warm start without memory held, once the PDM module has been initialised (see [Section 4.2](#)), **PDM_eReadDataFromRecord()** must be called for each record of application data in EEPROM that needs to be copied to RAM.

Stack Context Data

The function **PDM_eReadDataFromRecord()**, described above, is not used for records of stack context data. Loading this data from the EEPROM to RAM is handled automatically by the stack (provided that the PDM has been initialised).

4.3.3 Deleting Data in EEPROM

An individual record of application data in the EEPROM can be deleted using the function **PDM_eDeleteData()** - the record to be deleted is specified using its 16-bit identifier. Alternatively, all records (application data and stack context data) in the EEPROM can be deleted using the function **PDM_eDeleteAllData()**.



Caution: You are not recommended to delete records of ZigBee PRO stack context data by calling **PDM_eDeleteAllData()** before a rejoin of the same secured network. If these records are deleted, data sent by the node after the rejoin will be rejected by the destination node since the frame counter has been reset on the source node. For more information and advice, refer to the “Application Design Notes” appendix in the ZigBee PRO Stack User Guide (JN-UG-3101).

4.4 Storing Counters in EEPROM

The PDM provides a means of using the JN516x EEPROM to store counters, such as frame counters and acknowledgement sequences, as found in many communications protocols. A counter is implemented within a single EEPROM segment as follows:

- **Start value** which is held in pure binary form inside the counter's header
- **Incremental value** (over the start value) which is represented as a bitmap

Each bit of the bitmap represents an increment (by one) of the counter and is set when the corresponding increment occurs. There is a maximum incremental value that can be represented in one segment. When this value is reached, the counter is silently moved to a new segment in which the start value (in the header) is increased appropriately and the bitmap is reset to zero. To avoid increasing the segment wear count, the old bitmap segment is not formally deleted when a new segment is started. This process continues while there are segments free in the EEPROM.

The sub-sections below describe how to manage a counter in EEPROM using the PDM functions.

4.4.1 Creating a Counter

The function **PDM_eCreateBitmap()** can be used to create a counter in the EEPROM. In this function call, the new counter must be given a user-defined 16-bit identifier and a start value (these values will be stored in the counter's header).

4.4.2 Incrementing a Counter

The application can increment the counter by calling the function **PDM_eIncrementBitmap()**. When an increment results in the counter filling the current bitmap/segment, the function will indicate this by returning **PDM_E_STATUS_SATURATED_OK**. The next time the function is called, the counter will automatically be moved to a new bitmap/segment (as described above). However, if there is no free segment available, the function will be unable to perform the increment and will return **PDM_E_STATUS_USER_PDM_FULL**.

4.4.3 Reading a Counter

A counter in EEPROM can be read using the function **PDM_eGetBitmap()**. This function obtains the start value (stored in the counter's header) and the incremental value from the bitmap. The current value of the counter is then the sum of these two results.

The above function should be called when the JN516x device comes up from a cold start, to check whether a counter is present in EEPROM.

4.4.4 Deleting a Counter

Once a counter is no longer required, it can be removed from EEPROM using the function **PDM_eDeleteBitmap()**. This clears the current segment and all the older (expired) segments for the counter. This deletion increments the wear counts for these segments (see [Section 4.5.4](#)) and should be done only if absolutely necessary, as expired bitmap segments can be re-used directly via the PDM without formal deletion.

4.5 PDM Features

4.5.1 Mutex in PDM

PDM functions are not re-entrant and a mutex can be optionally used to prevent concurrent PDM function calls - if enabled, the mutex is applied automatically during a PDM function call.

If required, a mutex can be specified when the PDM module is initialised using the function **PDM_eInitialise()** - see [Section 4.2](#). If the mutex is used with the JenOS RTOS in a ZigBee application, the user task must be linked to the relevant mutex in the JenOS Configuration Editor - see [Section 14.7](#).

The principles of a mutex are described in [Section 2.4.7](#). All tasks that use a mutex must be connected to the mutex in the OS configuration diagram. The ZigBee PRO stack makes PDM calls. Therefore, any task that makes ZigBee PRO function calls must be connected to the PDM mutexes.



Note: The mutex does not remain optional when the PDM is used without the JenOS RTOS in IEEE802.15.4 and JenNet-IP applications. In applications developed using the IEEE802.15.4 SDK (JN-SW-4163), the mutex is not available and the relevant parameter is removed from **PDM_eInitialise()**. In applications developed using the JenNet-IP SDK (JN-SW-4165), the mutex is always implemented and a non-zero value must be passed to the relevant parameter of **PDM_eInitialise()**. For more information, refer to [Section 4.2](#).

4.5.2 Event and Error Handler for EEPROM

The internal PDM library allows a handler to be called to alert the application of events and error conditions in the JN516x internal EEPROM. This callback function is registered by calling the function **PDM_vRegisterSystemCallback()**. The PDM events/error conditions are listed and described in [Section 14.7](#).

An application must trap `E_PDM_SYSTEM_EVENT_PDM_NOT_ENOUGH_SPACE` and `E_PDM_SYSTEM_EVENT_DESCRIPTOR_SAVE_FAILED` callback errors

during testing. The ZigBee PRO stack uses multiple records. Once an 'out of space' error has occurred, the records will be in an inconsistent state. The software must be altered to use smaller record sizes or an external SPI Flash device. The PDM record sizes for the ZigBee PRO stack are dependent on table sizes set in the ZPS Configuration Editor.

The registered callback function may also be designed to handle a Wear Count event `E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED` which indicates that the Wear Count for an EEPROM segment has reached the configured trigger level (see [Section 4.5.4](#)).

4.5.3 EEPROM Capacity

The JN516x internal EEPROM consists of multiple small segments. On the JN5168 and JN5169 devices, there are 63 segments of 64 bytes each. The internal PDM library can store no more than one data record in each segment, although a large record may be stored across multiple segments. The PDM library needs to store some system information in each segment, so in practice each segment can hold only up to 56 bytes of record data. This means that a PDM record that has a single byte of information will need the same space as a 56-byte record and that a 57-byte record will need two segments (the same as a 112-byte record).

The function **PDM_u8GetSegmentCapacity()** returns the number of segments that are free for PDM. The function **PDM_u8GetSegmentOccupancy()** returns the number of segments that are in use. One of these functions may be called after all the records have been created and saved (including records in the ZigBee PRO stack). When updating a record, the PDM saves the new data before deleting the old data (to ensure that data is retained over any unexpected power cycles). Therefore, there must be sufficient capacity in the EEPROM to store another copy of a record before the old copy is deleted. To allow for the worst-case scenario, the value returned by **PDM_u8GetSegmentCapacity()** must be greater than the number of segments required to store the largest record.

4.5.4 EEPROM Wear Count

An EEPROM device supports a limited number of data writes to each byte before the storage medium begins to fail. For the JN516x EEPROM, at least 100000 writes are guaranteed and a million writes should be typically possible. For each EEPROM segment, a record is kept of the number of writes made to the segment so far. This is the 'Wear Count', which is stored and maintained in the segment header. The PDM manages the use of EEPROM segments in a way that minimises wear and attempts to spread the wear evenly across the segments.

The function **PDM_eGetSegmentWearCount()** allows the current value of the Wear Count of a particular segment to be obtained. It is also possible to set up the generation of an event when the Wear Count of any segment reaches a certain trigger level. This trigger level can be configured (for all segments) using the function **PDM_vSetWearCountTriggerLevel()**. The Wear Count event is `E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED` and

the user-defined PDM callback function (see [Section 4.5.2](#)) should be designed to process this Wear Count event.

4.5.5 Ensuring Consistency of PDM Records

The data in the PDM may differ in structure from that expected by the application. The structures stored by the ZigBee PRO libraries can change due to altering table sizes in the ZPS Configuration Editor, as well as between releases of the ZigBee PRO stack libraries. Inconsistency can occur under the following circumstances:

- The internal EEPROM on a JN516x device is not erased when programming an application with the JN51xx Flash Programmer. If multiple applications are run on the same hardware, it is unlikely that the structures will be consistent between the applications.
- When a ZigBee Over-The-Air (OTA) software update is performed, the PDM data is not erased. This is normally a benefit because it allows the application to rejoin the network. However, if any of the PDM structures change, a factory reset must be performed by calling **PDM_eDeleteAllData()**

Applications normally contain a way to perform a factory reset of the PDM module - for example, by calling **PDM_eDeleteAllData()** if a button is held down during reset.

The application can automatically check for PDM consistency by storing an application-specific 'magic number' in a record. A new magic number should be used if the application software or ZigBee PRO libraries PDM usage is inconsistent with the previous version of the software. Immediately after calling **PDM_eInitialise()**, the application should call **PDM_eReadDataFromRecord()**. If the magic number does not match, the application should call **PDM_eDeleteAllData()** to erase all records before attempting to start the ZigBee PRO stack. If the call to **PDM_eReadDataFromRecord()** indicates that the record has not been found, the application should also call **PDM_eDeleteAllData()** because another application may have been running that does not use the same record ID but has written inconsistent ZigBee PRO records to the PDM module.

Chapter 4
Persistent Data Manager (PDM) for EEPROM

5. Power Manager (PWRM)

This chapter describes the Power Manager (PWRM) module, which manages the transitions of the JN516x device into and out of low-power modes.

Low-power modes are typically used to prolong the battery life of a node by reducing the power consumption of the device during periods when the node does not need to receive, transmit or perform any other activities. Thus, low-power modes only apply to End Devices, as the Co-ordinator and Routers always need to remain fully alert for routing purposes.

5.1 Low-Power Modes

A number of low-power modes are available on the JN516x device. In descending order of power consumption, the modes are:

- Doze mode
- Sleep modes:
 - Sleep with memory held
 - Sleep without memory held
- Deep Sleep mode

When the node is inactive, the Power Manager will put the device into the lowest power mode possible.

The above low-power modes are described in the sub-sections below. For further information on the low-power modes of the JN516x device, refer to the *JN516x Data Sheet (JN-DS-JN516x)*.

5.1.1 Doze Mode

In Doze mode, the CPU of the chip pauses (the CPU clock is stopped) but all other parts of the JN516x device continue to run. Any interrupt will cause Doze mode to terminate and the application program will continue running from the next instruction. To prevent the Watchdog firing when in Doze mode, the application should ensure that a timer is running at a higher frequency than the Watchdog expiry period.

5.1.2 Sleep Mode with Memory Held

During Sleep with memory held, the contents of on-chip RAM are maintained, including stack context data and application data. Thus, on waking, the device can recover from sleep very quickly to continue normal operation from the next instruction.

In this mode, all power domains are powered down except those for the on-chip RAM and VDD supply. In addition, the 32-kHz on-chip oscillator can optionally be left running, which allows the device to be woken from sleep using wake timers.

Otherwise, the device can only be woken by changes on the DIO pins or the comparator input, or by a pulse counter expiry.

Although the contents of memory are held, on waking it is still necessary to re-configure the IEEE 802.15.4 stack layers and to re-initialise most of the on-chip peripherals. Wake callback functions can be registered for this purpose:

- You DO NOT have to re-initialise the DIOs, wake timers and comparator.
- You DO have to re-initialise everything else, including all other on-chip peripherals, the IEEE 802.15.4 MAC layer and, if using callbacks, the Programmable Interrupt Controller (PIC) - the callback functions must be re-registered. On the JN516x device, the SPI hardware must also be re-initialised.

5.1.3 Sleep Mode without Memory Held

During Sleep without memory held, on-chip RAM is powered down, and therefore stack context data and application data are not preserved on-chip. Normally, this data must be saved to external NVM (Non-Volatile Memory) before the chip enters sleep mode, and then recovered from NVM on waking (see [Chapter 3](#)).

In this mode, all power domains are powered down except the VDD power supply domain. Again, the 32-kHz on-chip oscillator can optionally be left running, which allows the device to be woken from sleep using wake timers. Otherwise, the device can only be woken by changes on the DIO pins or the comparator input, or by a pulse counter expiry.

On waking, the application program must be re-loaded from external NVM before the node can resume operation. All variables and peripherals must be re-initialised, except those used as wake sources and the DIO lines.

5.1.4 Deep Sleep Mode

In Deep Sleep mode, all switchable power domains are powered down and the 32-kHz oscillator is stopped. The device can be woken from deep sleep either via a hardware reset (by taking the RESETN pin low or by power cycling the device) or a change on the DIO pins.

On waking, the application program must be re-loaded from external NVM before the node can resume operation. All variables and peripherals must be re-initialised, including the DIO lines.

5.2 Callback Functions for Power Manager

If you intend to use the Power Manager, a number of callback functions must be available for the Power Manager to call in order to:

- start the application (see [Section 5.2.1](#))
- perform housekeeping tasks when entering and leaving low-power mode (see [Section 5.2.2](#))
- handle interrupts from Wake Timer 1 (see [Section 5.2.3](#))

5.2.1 Essential Callback Function

When your application uses the Power Manager, you must define and use the callback function **vAppMain()** in your code. The main task of your application must be included in this function (which must never return).

5.2.2 Pre-sleep and Post-sleep Callback Functions

In order to implement low-power modes, you must provide the Power Manager with user-defined callback functions to perform housekeeping tasks when the node enters and leaves low-power mode. Registration functions are provided for these callback functions, where the registration functions must be called in the user-defined callback function **vAppRegisterPWRMCallbacks()**.

- The pre-sleep callback function is called by the Power Manager just before putting the device into low-power mode. This callback function is registered in your code through the API function **PWRM_vRegisterPreSleepCallback()**.
- The post-sleep callback function is called by the Power Manager just after the device leaves low-power mode (irrespective of how the device was woken from sleep). This callback function is registered in your code through the API function **PWRM_vRegisterWakeupCallback()**.

vAppRegisterPWRMCallbacks() is called by the stack as part of a cold start.

The pre- and post-sleep callback function themselves must each be declared in the code using the macro

PWRM_CALLBACK(*fn_name*)

where *fn_name* is the name of the callback function.

Each of these callback functions must also have a descriptor. This is a structure that is used in the above registering functions to specify the callback function to register.

The callback descriptor must be declared using the macro

PWRM_DECLARE_CALLBACK_DESCRIPTOR(*desc_name*, *fn_name*)

where *desc_name* is the descriptor name and *fn_name* is the callback function name.

For example:

```
PWRM_CALLBACK(vPreSleepCB1);  
PWRM_DECLARE_CALLBACK_DESCRIPTOR(pscbl_desc, vPreSleepCB1);
```

5.2.3 Wake Timer Callback Function

If you intend to use wake events, derived from Wake Timer 1, your interrupt handler must call the pre-defined callback function **PWRM_WakeInterruptCallback()**. This function maintains the list of scheduled wake events - if required, it will re-start the wake timer for the next wake point. It also calls the user-defined callback function specified through **PWRM_vScheduleActivity()**.

For further information on waking the device using scheduled wake events, refer to [Section 5.6](#).

5.3 Initialising and Starting the Power Manager

The Power Manager is initialised and started using the function **PWRM_vInit()**. This function requires one of five possible low-power configurations to be specified:

- Sleep with 32-kHz oscillator running and memory held
- Sleep with 32-kHz oscillator running and memory not held
- Sleep with 32-kHz oscillator not running and memory held
- Sleep with 32-kHz oscillator not running and memory not held
- Deep sleep (oscillator not running and memory not held)

The specified configuration is the low-power mode in which the Power Manager will attempt to put the device during inactive periods. Note that Doze mode cannot be explicitly specified, but the Power Manager may put the device into Doze mode at times when the specified mode cannot be entered (see [Section 5.8.1](#)).

The criteria for selecting a sleep mode are as follows:

- **Oscillator setting:**
 - If the 32-kHz oscillator is left running during sleep, a wake point can be scheduled using **PWRM_vScheduleActivity()** - see [Section 5.6](#).
 - Otherwise, the device must be woken by an external event (a change on a DIO line or comparator input, a pulse counter expiry or a reset).
- **Memory setting:**
 - If memory is held during sleep, stack context data and application data will be preserved in memory, allowing the device to quickly resume operation through a warm re-start following sleep.
 - Sleep without memory held provides a greater power saving. However, stack context data and application data must be saved to external NVM (Non-Volatile Memory) before entering sleep mode and restored into on-chip memory during a cold re-start on exiting sleep (see [Chapter 3](#)).

5.4 Enabling Power-Saving

To enable the Power Manager to put the JN516x device into low-power mode at appropriate times, you must call the function **PWRM_vManagePower()**, normally called from the OS idle task. Once this function has been called, the Power Manager will, whenever possible, put the JN516x device into the sleep mode specified through **PWRM_vInit()** (or, alternatively, into Doze mode - see [Section 5.8.1](#)).



Note: Sleep mode cannot be entered while there are software timers active (in running or expired states). You must therefore de-activate any such timers to allow the Power Manager to put the JN516x device into sleep mode - refer to the Caution in [Section 2.4.6](#).

5.5 Non-interruptible Activities

In order to enter sleep mode, no activities must be running that must not be interrupted by sleep. This condition for entering sleep mode is monitored using an activity counter - sleep mode can only be entered when this counter is zero. The application is responsible for maintaining the activity counter, as follows:

- Whenever an activity is started that must not be interrupted by sleep, the application must notify the Power Manager using the function **PWRM_eStartActivity()**, which increments the activity counter.
- Whenever such an activity is completed, the application must notify the Power Manager using the function **PWRM_eFinishActivity()**, which decrements the activity counter.



Caution: ***PWRM_eFinishActivity()** must only be called by an application following a matching call to **PWRM_eStartActivity()**. The OS and ZigBee PRO stack both use the activity counter, so calling **PWRM_eFinishActivity()** inappropriately can leave the OS or ZigBee PRO stack in an inconsistent state.*

You can obtain the current value of the activity counter using the function **PWRM_u16GetActivityCount()**.

5.6 Terminating Low-Power Mode

Low-power modes can be terminated in a number of ways:

- **Any Interrupt:** When in Doze mode, the device can be woken by any interrupt.
- **Wake Timer:** When in Sleep mode in which the 32-kHz oscillator runs, the device can be woken by a scheduled wake event configured using the function **PWRM_vScheduleActivity()**. For more information on scheduled wake events, refer to [Section 5.6](#).
- **External Wake Event:** The following external wake events are available:
 - **DIO:** When in Sleep and Deep Sleep modes, the device can be woken by a change of state of a DIO line.
 - **Comparator input:** When in Sleep mode, the device can be woken by a change of state of the comparator input.
 - **Pulse counter:** When in Sleep mode, the device can be woken on expiry of the pulse counter, which counts pulses on an external input.

The above external wake events can be controlled by functions of the JN516x Integrated Peripherals API, described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)*. In addition, the applicable external wake events must be configured in the RTOS configuration for the application.

- **Hardware Reset:** When in Deep Sleep mode, the device can be woken by a hardware reset.

The valid wake sources for the different low-power modes are summarised in [Table 1](#) below.

On leaving low-power mode, the Power Manager will call the user-defined callback function that has been registered using **PWRM_vRegisterWakeupCallback()**.

Low-Power Mode	Wake Source					
	Any Interrupt	Wake Timer	DIO	Comparator	Pulse Counter	Hardware Reset
Doze mode	✓	✗	✗	✗	✗	✗
Sleep mode with oscillator running and memory held	✗	✓	✓	✓	✓	✗
Other Sleep modes	✗	✗	✓	✓	✓	✗
Deep Sleep mode	✗	✗	✓	✗	✗	✓

Table 1: Valid Wake Sources for Low-Power Modes

5.7 Scheduling Wake Events



Note: This section is only applicable to the sleep mode in which the 32-kHz oscillator is left running and memory is held.

In **PWRM_vInit()**, if you have selected the Sleep mode with the 32-kHz oscillator running and memory held, you can schedule wake events which ensure that the device will be awake at certain times - that is, if the device is sleeping, it will be woken at the scheduled time. This scheduling uses Wake Timer 1 of the JN516x device, which operates at 32 kHz.

A wake event can be scheduled using the function **PWRM_eScheduleActivity()**. This function requires you to specify the number of ticks of the wake timer until the wake event. You must also specify the user-defined callback function that must be called when the wake event occurs.

When the wake timer expires for a scheduled wake event, an interrupt is generated. The application's interrupt handler then calls the pre-defined callback function **PWRM_WakeInterruptCallback()**. This function maintains the list of scheduled wake events and, if necessary, re-starts the wake timer for the next scheduled wake event. The function also calls the user-defined callback function specified through **PWRM_eScheduleActivity()**.



Note: In addition, when the device wakes from sleep, the user-defined callback function registered through **PWRM_vRegisterWakeupCallback()** will also be called. However, this is a general-purpose wake-up function which is called irrespective of how the device was woken (it is not unique to scheduled wake events, but also called for external wake events).

5.8 Doze Mode

Doze mode is a lighter power-saving mode than the sleep modes, as all elements of the JN516x device remain powered but the CPU is paused (CPU clock is stopped).

This low-power mode cannot be explicitly selected in **PWRM_vInit()**. The Power Manager will put the JN516x device into Doze mode only in certain circumstances, described in [Section 5.8.1](#) below. However, to enter Doze mode, the Power Manager must have been initialised using **PWRM_vInit()** and the power-saving modes must have been enabled using **PWRM_vManagePower()**.

5.8.1 Circumstances that Lead to Doze Mode

Although Sleep and Deep Sleep modes cannot be entered while there are activities running that must not be interrupted by sleep (see [Section 5.5](#)), the Power Manager can put the device into Doze mode while the activity counter is non-zero.

Even when the activity counter is zero, if a sleep mode has been configured with the 32-kHz oscillator running (see [Section 5.3](#)) but no wake event has been scheduled (see [Section 5.6](#)), the Power Manager will put the device into Doze mode instead of Sleep mode.

The decision to put a device into a Sleep mode or Doze mode is illustrated in the flowchart in [Figure 6](#) below.

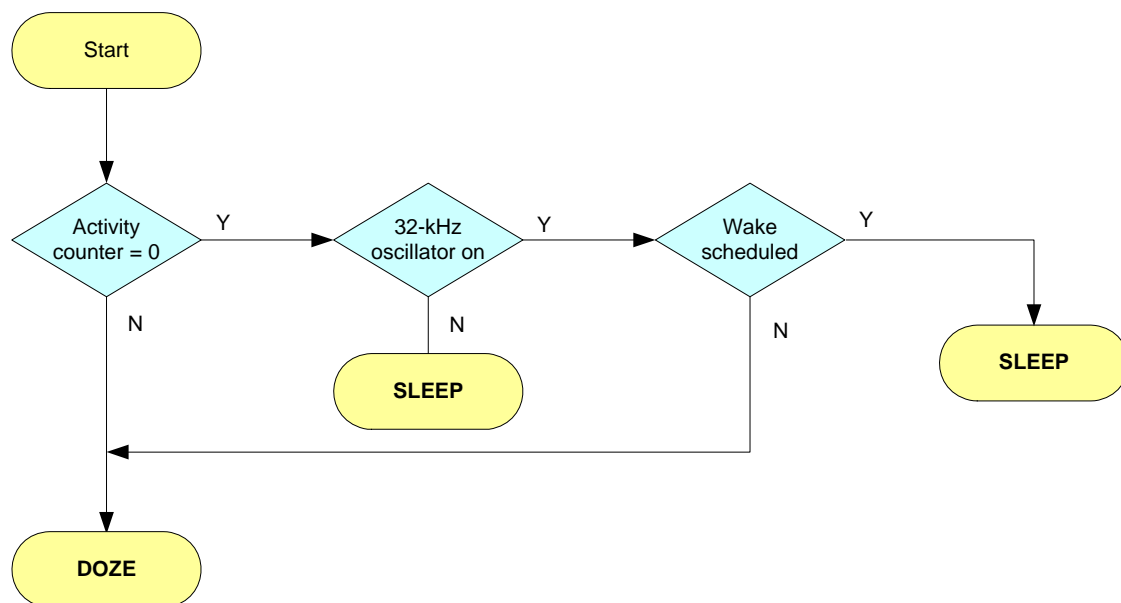


Figure 6: Flowchart of Decision to Enter Doze Mode

5.8.2 Doze Mode Monitoring During Development

Depending on the circumstances described in [Section 5.8.1](#), the JN516x device may spend a significant proportion of its time in Doze mode. The Power Manager API provides a function that allows you to investigate the fraction of time that the JN516x device typically spends in Doze mode for a given application. The function provides a doze monitoring output on the DIO1 pin of the JN516x device. This functionality can be used when the application is running in debug mode.

The function **PWRM_vSetupDozeMonitor()** must be called to start a monitoring session. The state of the DIO1 pin will then reflect the doze state of the device, allowing you to make doze state measurements using external equipment. The fraction of time that the JN516x device spends in Doze mode can then be estimated as (see [Figure 7](#)): *Total time in Doze mode during session / Elapsed time of session*




Figure 7: Doze Monitoring

To obtain sensible results, doze monitoring should be allowed to run for a significant period of time.

6. Protocol Data Unit Manager (PDUM)

Communication between nodes in a wireless network is implemented using messages which contain application data. The part of a message which contains this data is called the Application Protocol Data Unit (APDU). The Protocol Data Unit Manager (PDUM) is concerned with APDU memory management, and assembling and disassembling APDUs - that is, inserting data into APDUs to be transmitted and extracting data from received APDUs.

6.1 Message Assembly and Disassembly

A message travels over a wireless network as a packet (usually an 802.15.4 packet) containing application data surrounded by header and footer information relating to the different layers of the protocol stack.

A message to be sent is prepared at the application level, at the top of the protocol stack, by creating an Application Protocol Data Unit (APDU) containing the application data to be included in the message. This APDU is then passed down the layers of the stack, with each layer adding its own protocol information to the header and footer. On reaching the 'physical' layer at the bottom of the stack, the message is complete and ready to be transmitted.

For transmission, the message is converted to an NPDU (Network Protocol Data Unit). If the length of the message is greater than the packet size used in network communication (e.g. 802.15.4 packet size), the message is divided up and transmitted in multiple NPDUs (Network Protocol Data Units). You will need to be aware of this if using a sniffer to detect transmitted packets.



Note: Data is stored in memory in the JN516x device in big-endian byte order but is transmitted over the network in little-endian byte order.

A received message is passed up the protocol stack, with each stack layer stripping out the corresponding protocol information from the header and footer. On reaching the application level, only the APDU remains. The application data can then be extracted from this APDU.

The assembly and disassembly of a message, described above, are illustrated in the figure below, in which the lower stack layers (MAC and Physical) are provided by the IEEE 802.15.4 protocol.

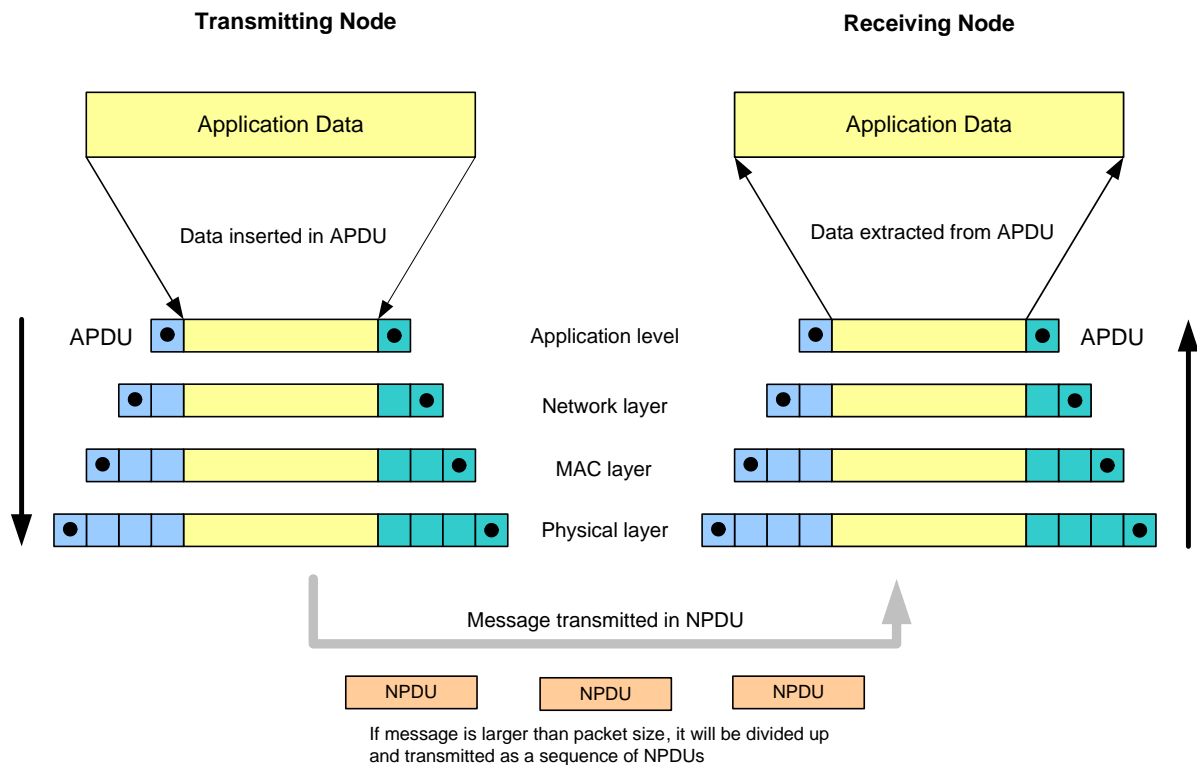


Figure 8: Message Assembly and Disassembly

6.2 Preparing the PDU Manager

In order to use the PDU Manager:

- You must statically define the required APDUs using the ZPS Configuration Editor (described in the *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)*). Each APDU is given a unique handle. While the data payload of an APDU can be of arbitrary length, a maximum length is set for an APDU.
- Before calling any other PDUM functions in your code, you must call the function **PDUM_vlnit()** to initialise the PDU Manager.

6.3 Inserting Data into Outgoing Message

When sending a message to another node, you must first create an APDU containing the application data to be sent. To do this, first allocate an APDU instance by calling the function **PDUM_hAPduAllocateAPdulInstance()** and then populate the APDU instance with data using **PDUM_u16APdulInstanceWriteNBO()**, in which you must specify:

- the handle of the APDU instance in which data is to be inserted (this is the handle returned by **PDUM_hAPduAllocateAPdulInstance()**)
- the starting position of the data in the APDU - that is, the position of the least significant data byte
- the format of the data payload - the data can be made up of a sequence of data values of different types
- the data values to be inserted in the data payload

Alternatively, the function **PDUM_u16APdulInstanceWriteStrNBO()** can be used to populate the APDU instance - this function allows a data structure to be inserted into the APDU.

You must then use the relevant ZigBee PRO API function to send the message - refer to the *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)*. Once the message has been sent, the ZigBee PRO stack automatically de-allocates the memory-space used for the APDU instance.

Note that **PDUM_u16APdulInstanceWriteNBO()** performs the necessary data conversion from big-endian byte order to little-endian byte order for transmission.

Alternatively, you can produce your own code to insert data into the payload of an APDU. To help you, two functions are provided:

- **PDUM_pvAPdulInstanceGetPayload()**: This function returns a pointer to the start of the payload section of the APDU instance.
- **PDUM_eAPdulInstanceSetPayloadSize()**: This function sets the size, in bytes, of the data payload.



Caution: Data must be stored in memory in big-endian order but is transmitted over the network in little-endian byte order. Therefore, if you use your own code to insert data into an APDU, you must reverse the byte order of the data before inserting it. Failure to change the endianness of the data will result in an alignment exception.

6.4 Extracting Data from Incoming Message

The function **PDUM_u16APduInstanceReadNBO()** provides an easy way of extracting the data payload from an incoming message, which is received using the RTOS function **OS_eCollectMessage()** - refer to [Section 2.4.8](#). The **PDUM_u16APduInstanceReadNBO()** function requires the following to be specified:

- the handle of the APDU instance containing the data to be extracted (this is the handle contained in the ZPS_EVENT_AF_DATA_INDICATION stack event which notified the application of the arrival of the data message)
- the starting position of the data in the APDU - that is, the position of the least significant data byte
- the format of the data payload - the data can be made up of a sequence of data values of different types
- a pointer to a structure in which the extracted data will be stored

Once the data has been extracted, you should de-allocate the memory space used for the APDU instance by calling the function **PDUM_eAPduFreeAPduInstance()**.

Note that **PDUM_u16APduInstanceReadNBO()** performs the necessary data conversion from little-endian byte order to big-endian byte order for storage.

Alternatively, you can produce your own code to extract the payload data from an APDU. To help you, two functions are provided:

- **PDUM_pvAPduInstanceGetPayload()**: This function returns a pointer to the start of the payload data in the APDU instance.
- **PDUM_u16APduInstanceGetPayloadSize()**: This function returns the size, in bytes, of the data payload.



Caution: Data is received from the network in little-endian byte order, but must be stored in memory in big-endian order. Therefore, if you use your own code to extract data from an APDU, you must reverse the byte order of the data before storing it. Failure to change the endianness of the data will result in an alignment exception.

7. Debug (DBG) Module

This chapter describes the Debug (DBG) module which allows application code to be debugged by means of diagnostic messages that are output to a display device.

7.1 Overview

The Debug module comprises an API containing diagnostic functions that can be embedded in your application code. Application debugging using the Debug module requires the JN516x device to be connected to a display device (such as a PC) via an IO interface, such as one of the on-chip UARTs. The display device must provide a dumb terminal through which output from the JN516x device can be viewed. A typical implementation is illustrated in the figure below.

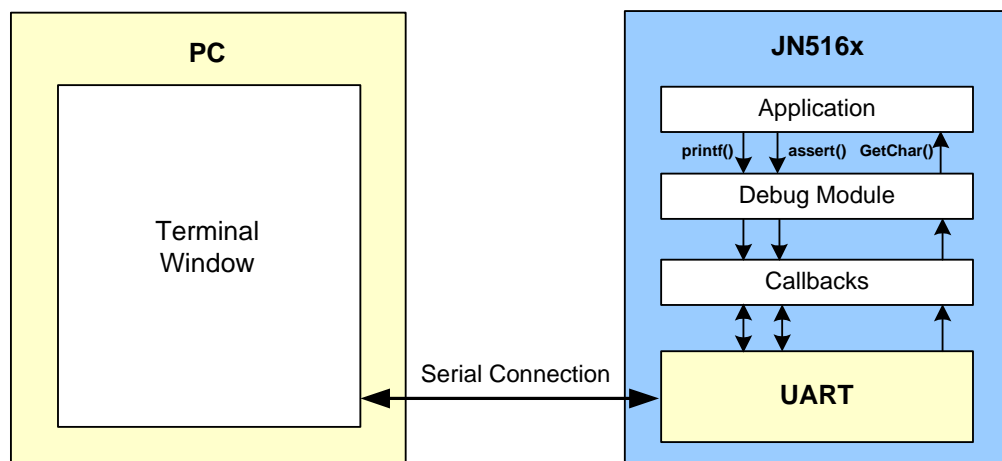


Figure 9: Typical Hardware Set-Up for Debugging

The API provides the essential printf- and assert-style debug functions, which can be strategically placed in your code:

- **DBG_vPrintf()** is used to output formatted strings and data values at an appropriate point during program execution, in order to indicate progress.
- **DBG_vAssert()** is used to test a logical condition, and to stop program execution if the test fails (condition is FALSE).

User-specified callback functions are used by the Debug module to control the IO interface (see [Section 7.3](#)).

The terminal on the PC can also supply input to the JN516x UART. The function **DBG_iGetChar()** can be used by the application to obtain a character from this input source. This input can then be handled by the JN516x application.

7.2 Enabling the Debug Module

The Debug module API is defined in the header file **DBG.h**, which must be included in your code.

In order to use the Debug module, it must be explicitly enabled at build time by defining **DBG_ENABLE** in the build - for example, by adding **-DDBG_ENABLE** to the compiler. If the module is not enabled in this way, all the Debug functions embedded in your code will be ignored.

In addition, the functions **DBG_vPrintf()** and **DBG_vAssert()** each include a Boolean parameter which can be used to enable/disable individual instances of these functions. Two or more instances of these functions can be grouped to form a 'stream' for which this Boolean parameter is a common constant used to enable/disable the whole function group. This constant can be set at build time (see [Section 7.5](#)).



Tip: By default, the Debug module will display each 'printf line' as passed. However, if **DBG_VERBOSE** is defined at build time then each line displayed will be prefixed with the file name and line number of the debug statement.

7.3 Initialising and Configuring the Debug Module

The way that the Debug module is configured and initialised depends on the serial IO interface which is to be used to output debug information from the JN516x device to the attached PC:

- If a JN516x UART is to be used for output, the required initialisation/configuration is as described in [Section 7.3.1](#). This option will be taken by most users.
- If any other serial IO interface is to be used for output, the required initialisation/configuration is as described in [Section 7.3.2](#).

Flags are provided in the global variable *DBG_u32Flags* for configuring certain aspects of the Debug module - for details, refer to [Section 7.4](#).

7.3.1 Using JN516x UART Input/Output

When a JN516x UART is to be used for the input/output of debug information, the configuration and initialisation of the Debug module is accomplished with a single call to the function **DBG_vUartInit()**, which allows selection of the UART (0 or 1) and the baud-rate to be used. This function is used both during a cold start of the JN516x device and during a warm start (where the latter is a device re-start with memory contents retained).

7.3.2 Using Alternative Serial Output

When an alternative to an on-chip UART is to be used for the output of debug information, the required IO interface must first be configured and enabled (using the relevant functions from the JN516x Integrated Peripherals API).

The Debug module must then be initialised using the function **DBG_vlnit()**. This function is used both during a cold start of the JN516x device and during a warm start (where the latter is a device re-start with memory contents retained). The function takes as input a structure which contains pointers to four callback functions needed for debugging:

```
typedef struct
{
    void (*prInitHardwareCb)(void);
    void (*prPutchCb)      (char c);
    void (*prFlushCb)      (void);
    void (*prFailedAssertCb)(void);
} tsDBG_FunctionTbl;
```

The callback functions are user-defined and are described in the table below.

Pointer	Callback Function
<i>*prInitHardwareCb</i>	Function which re-initialises the IO interface after a warm start, e.g. when JN516x device wakes from sleep.
<i>*prPutchCb</i>	Function used by DBG_vPrintf() to output a single character to the IO interface.
<i>*prFlushCb</i>	Function used by DBG_vPrintf() to flush the IO interface buffer to allow buffered output characters to be displayed. If the output is unbuffered, this function should do nothing or wait for the last character output using the putch() function to be made available. Note that the function should not append a newline character, as this should be handled by the formatting string passed to DBG_vPrintf() .
<i>*prFailedAssertCb</i>	Function which is called when DBG_vAssert() fails. The function should stop execution and may reset the device.

Table 2: Callback Functions Specified in DBG_vlnit()

7.4 Debug Configuration Flags

The Debug module has a global variable *DBG_u32Flags* which is a bitmap containing configuration flags. The bits/flags are enumerated, and are listed and described in the table below.

Flag/Enumeration	Description (if flag is set)
DBG_FLAG_NONE	None of the flags are set
DBG_FLAG_OUTGOING_NL_CRNL	Every \n character in the outgoing string is sent as \r\n. This is for compatibility with certain terminal programs.
DBG_FLAG_AUTO_FLUSH	DBG_vFlush() is called at the end of each DBG_vPrintf() invocation. The application may instead choose to flush the outgoing characters in idle time rather than at the end of each outputted string.
DBG_FLAG_FLUSH_WHEN_FULL	If the DBG back-end buffers outgoing characters then it will automatically flush the buffer when full. Otherwise, characters that do not fit in the buffer may be lost.

Table 3: Debug Configuration Flags



Note: The flags `DBG_FLAG_OUTGOING_NL_CRNL` and `DBG_FLAG_AUTO_FLUSH` are set by default.

7.5 Example Diagnostic Code

The following code fragment illustrates use of the Debug module API. The JN516x UART 0 is used. Two debug 'streams' (1 and 2) are used to separately enable/disable two groups of debug lines.

```
#include <jendefs.h>
#include "DBG.h"
#include "DBG_Uart.h"

#ifndef DBG_STREAM_1
#define DBG_STREAM_1 FALSE
#endif

#ifndef DBG_STREAM_2
#define DBG_STREAM_2 FALSE
#endif

void appColdStart(void)
{
    int i = 0;

    /* Initialise the standard UART hardware */
    DBG_vUartInit(DBG_E_UART_0, DBG_E_UART_BAUD_RATE_115200);

    /* Now we can use DBG_vPrintf() and DBG_vAssert() to output characters
       to the UART device */
    DBG_vPrintf(DBG_STREAM_1, "Printing to stream 1\n");
    DBG_vPrintf(DBG_STREAM_2, "Printing an integer %i to stream 2\n", 10);
    DBG_vAssert(DBG_STREAM_1, i == 1);
}
```

When building this application, you have following options:

- Debug disabled (the default)
- Debug enabled only for stream 1 - build with:
-DDBG_ENABLE -DDBG_STREAM_1=TRUE
- Debug enabled only for stream 2 - build with:
-DDBG_ENABLE -DDBG_STREAM_2=TRUE
- DBG enabled for both streams - build with:
-DDBG_ENABLE -DDBG_STREAM_1=TRUE -DDBG_STREAM_2=TRUE

Part II: Reference Information

8. RTOS API

The chapter contains descriptions of the macros and functions of the JenOS RTOS (Real-Time Operating System) API. The API is defined in the header files **os.h** and **os_lib.h**.

- The RTOS macros are described in [Section 8.1](#)
- The RTOS functions are described in [Section 8.2](#)



Caution: To control interrupts, you should only use the functions supplied as part of the RTOS.

8.1 RTOS Macros

The section contains descriptions of the macros of the RTOS API. These macros are used to define user tasks, ISRs (Interrupt Service Routines) and callback functions (related to the hardware counter and associated software timers), and are listed below.

Macro	Page
OS_TASK	72
OS_ISR	73
OS_SWTIMER_CALLBACK	74
OS_HWCOUNTER_ENABLE_CALLBACK	75
OS_HWCOUNTER_DISABLE_CALLBACK	76
OS_HWCOUNTER_SET_CALLBACK	77
OS_HWCOUNTER_GET_CALLBACK	78

OS_TASK

OS_TASK(*taskname*)

Description

This macro is used in the application code to define a user task. The name of the user task must be specified, where this task name has been declared in advance using the JenOS Configuration Editor (see [Section 14.4.1](#)).

The macro is followed by the task body, as illustrated in the code fragment below:

```
OS_TASK(TaskA) {  
    task body ...  
}
```

```
OS_TASK(TaskB) {  
    task body ...  
}
```

Parameters

<i>taskname</i>	Name of user task
-----------------	-------------------

OS_ISR

OS_ISR(*ISRname*)

Description

This macro is used in the application code to define an ISR (Interrupt Service Routine) which will be invoked when the corresponding interrupt occurs. The name of the ISR must be specified, where this ISR name has been declared in advance using the JenOS Configuration Editor (see [Section 14.4.2](#)). The association between this ISR and an interrupt source is also configured in this file.

The macro is followed by the ISR body, as illustrated in the code fragment below:

```
OS_ISR(TickTimerISR) {  
    ISR body ...  
}
```

Parameters

<i>ISRname</i>	Name of ISR
----------------	-------------

OS_SWTIMER_CALLBACK

OS_SWTIMER_CALLBACK(*CBname*, *pData*)

Description

This macro is used to define a 'software timer expiry' callback function. This user-defined function will be automatically invoked when a software timer expires which has been started using the function **OS_eStartSWTimer()** - that is, when the timer reaches the number of ticks specified in the start function. The callback function is called by the function **OS_eExpireSWTimers()** (as an alternative to activating a user task on expiry of the software timer).

The name of the callback function must be specified in the macro, where this name has been declared in advance using the JenOS Configuration Editor (see [Section 14.5.1](#)). A pointer to a buffer can optionally be specified, where this buffer will contain any data to be processed by the callback function. This data is passed to the callback function through a parameter of **OS_eStartSWTimer()**.

The macro is followed by the callback function definition, as illustrated in the code fragment below:

```
OS_SWTIMER_CALLBACK(MyCallback, pvMyData)
{
    Callback definition...
}
```

Parameters

<i>CBname</i>	Name of callback function
<i>pData</i>	Pointer to buffer containing data to be used by callback function (if no data is needed, specify a valid dummy pointer)

OS_HWCOUNTER_ENABLE_CALLBACK

OS_HWCOUNTER_ENABLE_CALLBACK(*CBname*)

Description

This macro is used to define a callback function to enable and start the hardware counter which may be used to drive one or more software timers. This user-defined function will be automatically invoked when a software timer is started using the function **OS_eStartSWTimer()**, provided that there are no other software timers already active that use the hardware counter.

The name of the callback function must be specified in the macro, where this name has been declared in advance using the JenOS Configuration Editor (see [Section 14.5.1](#)).

The macro is followed by the callback function definition, as illustrated in the code fragment below:

```
OS_HWCOUNTER_ENABLE_CALLBACK(EnableTickTimer)
{
    Callback definition...
}
```

Parameters

<i>CBname</i>	Name of callback function
---------------	---------------------------

OS_HWCOUNTER_DISABLE_CALLBACK

OS_HWCOUNTER_DISABLE_CALLBACK(*CBname*)

Description

This macro is used to define a callback function to stop and disable the hardware counter which may be used to drive one or more software timers. This user-defined function will be automatically invoked:

- when a software timer is stopped using the function **OS_eStopSWTimer()**
- by the function **OS_eExpireSWTimers()** when a software timer expires provided that there are no other software timers still active that use the hardware counter.

The name of the callback function must be specified in the macro, where this name has been declared in advance using the JenOS Configuration Editor (see [Section 14.5.1](#)).

The macro is followed by the callback function definition, as illustrated in the code fragment below:

```
OS_HWCOUNTER_DISABLE_CALLBACK(DisableTickTimer)
{
    Callback definition...
}
```

Parameters

<i>CBname</i>	Name of callback function
---------------	---------------------------

OS_HWCOUNTER_SET_CALLBACK

OS_HWCOUNTER_SET_CALLBACK(*CBname*, *CompValue*)

Description

This macro is used to define a callback function to set the value in the compare register for the hardware counter which may be used to drive one or more software timers. The compare register contains a value with which the incrementing counter value can be compared - action may be taken when the two values match.

This user-defined callback function is normally used to set a compare register value for the next software timer expiry point. Since the hardware counter increments continuously (and loops around), the compare register must be set to a value equal to the current counter value plus the number of ticks until the next expiry point.

The callback function will be automatically invoked when a software timer is started using the function **OS_eStartSWTimer()** or re-started using the function **OS_eContinueSWTimer()**. It is also called by **OS_eExpireSWTimers()** when a software timer expires and the compare register must be updated for the next software timer to expire (if any).

The name of the callback function must be specified in the macro, where this name has been declared in advance using the JenOS Configuration Editor (see [Section 14.5.1](#)). The required compare register value must also be specified.

The callback definition should follow the format below:

```
OS_HWCOUNTER_SET_CALLBACK(SetTickTimerCompare, u32CompVal)
{
    if (u32CompVal is in the future) {
        /* set the tick timer compare register */
        return TRUE;
    } else { /* compare value is in the past */
        /* expire timer immediately */
        return FALSE;
    }
}
```

Parameters

<i>CBname</i>	Name of callback function
<i>CompValue</i>	Value to set in the compare register for the hardware counter

OS_HWCOUNTER_GET_CALLBACK

OS_HWCOUNTER_GET_CALLBACK(*CBname*)

Description

This macro is used to define a callback function to obtain the current value of the hardware counter which may be used to drive one or more software timers. The function must return a **uint32** containing the hardware counter value.

The name of the callback function must be specified in the macro, where this name has been declared in advance using the JenOS Configuration Editor (see [Section 14.5.1](#)).

The macro is followed by the callback function definition, as illustrated in the code fragment below:

```
OS_HWCOUNTER_GET_CALLBACK(GetCurrentCount)
{
    Callback definition...
}
```

Parameters

<i>CBname</i>	Name of callback function
---------------	---------------------------

8.2 RTOS Functions

This section contains descriptions of the functions of the RTOS API. The API functions are sub-divided into the following areas, which are covered in separate sub-sections:

Functional Area	Section Reference
Initialisation	Section 8.2.1
User Tasks	Section 8.2.2
Interrupts	Section 8.2.3
Mutex	Section 8.2.4
Messaging	Section 8.2.5
Software Timers	Section 8.2.6

8.2.1 Initialisation Functions

This section describes the RTOS initialisation functions.

The functions are listed below, along with their page references:

Function	Page
OS_vStart	80
OS_vRestart	81

OS_vStart

```
void OS_vStart(  
    void (*prvInitFunction)(void),  
    void (*prvUnclaimedIRQ)(void),  
    void (*prvErrorHook)(OS_tStatus, void *osHandle));
```

Description

This function is used to start the RTOS.

Three optional user-defined functions can be specified:

- An initialisation function for the hardware and user application, e.g. initialises CPU peripherals and enables interrupt generation flags. This function is called with all controlled interrupts disabled. On completion of this function, all controlled interrupts are enabled.
- A function to capture any unclaimed interrupts that occur but do not have an ISR configured for them.
- A function to handle OS errors, as described in [Section 2.5](#).

On returning from **OS_vStart()**, the idle task is entered.



Caution: The user-defined initialisation function must not perform any operations that require the use of interrupts, since interrupts are not enabled until the RTOS has started, which occurs just before the function **OS_vStart()** returns.

Parameters

<i>*prvInitFunction</i>	Optional pointer to a user-defined initialisation function. If not used, this parameter must be set to NULL.
<i>*prvUnclaimedIRQ</i>	Optional pointer to user-defined function to capture unclaimed interrupts. If not used, this parameter must be set to NULL.
<i>*prvErrorHook</i>	Optional pointer to user-defined callback function to handle OS errors. Only required if error hook is set to true in the JenOS configuration editor.

Returns

None

OS_vRestart

```
void OS_vRestart(void);
```

Description

This function is used to restart the RTOS following a warm start with memory held. The function re-initiates the interrupt hardware.

Parameters

None

Returns

None

8.2.2 User Task Functions

This section provides descriptions of the RTOS API functions concerned with user tasks.

The functions are listed below, along with their page references:

Function	Page
OS_eActivateTask	83
OS_eGetCurrentTask	84

OS_eActivateTask

```
OS_teStatus OS_eActivateTask(OS_thTask hTask);
```

Description

This function is used to activate the specified user task - that is, to move the task from the dormant state to the pending state (i.e. ready to run). When it becomes the highest priority pending task, it will then execute.

If the task is already pending then its activation count is incremented. The activation count keeps track of the number of times the task must be executed. Once the task has been executed, the activation count is decremented. If this count reaches zero, the task is moved back to the dormant state, otherwise the task returns to the pending state.

A handle for the task is pre-defined using the JenOS Configuration Editor.

Parameters

<i>hTask</i>	Handle of the user task to be activated
--------------	---

Returns

- OS_E_OK (successful)
- OS_E_BADTASK (invalid task handle used)
- OS_E_OVERACTIVATION (maximum number of activations exceeded: 65535)

OS_eGetCurrentTask

```
OS_teStatus OS_eGetCurrentTask(OS_thTask *phTask);
```

Description

This function is used to obtain the handle of the user task that is currently in the running state.

Parameters

<i>*phTask</i>	Pointer to place to store task handle obtained
----------------	--

Returns

OS_E_OK (successful)
OS_BADVALUE (invalid pointer specified)

8.2.3 Interrupt Functions

This section provides descriptions of the RTOS API functions concerned with interrupts.

The functions are listed below, along with their page references:

Function	Page
OS_eDisableAllInterrupts	86
OS_eEnableAllInterrupts	87
OS_eSuspendOSInterrupts	88
OS_eResumeOSInterrupts	89



Caution: *These interrupt functions must **not** be called from within an Interrupt Service Routine (ISR).*



Note: The RTOS cannot clear interrupts and it is the responsibility of the application to do this - refer to [Appendix B](#).

OS_eDisableAllInterrupts

```
OS_tStatus OS_eDisableAllInterrupts(void);
```

Description

This function is used to disable all CPU interrupts - that is, both controlled and uncontrolled interrupts.

Note that nested calls to this function are not permitted and that the function cannot be called in an ISR.

These interrupts can be re-enabled using the function **OS_eEnableAllInterrupts()**.

Parameters

None

Returns

OS_E_OK (successful)

OS_eEnableAllInterrupts

```
OS_teStatus OS_eEnableAllInterrupts(void);
```

Description

This function is used to enable all CPU interrupts - that is, both controlled and uncontrolled interrupts.

Note that nested calls to this function are not permitted and that the function cannot be called in an ISR.

These interrupts can be disabled using the function **OS_eDisableAllInterrupts()**.

Parameters

None

Returns

OS_E_OK (successful)

OS_eSuspendOSInterrupts

```
OS_teStatus OS_eSuspendOSInterrupts(void);
```

Description

This function is used to disable interrupts managed by the RTOS - that is, controlled interrupts.

Note that nested calls to this function are permitted.

These interrupts can be subsequently re-enabled using the function **OS_eResumeOSInterrupts()**.

Parameters

None

Returns

OS_E_OK (successful)

OS_E_OSINTOVERFLOW (maximum nesting level exceeded: 0xFFFFFFFF)

OS_eResumeOSInterrupts

OS_teStatus OS_eResumeOSInterrupts(void)

Description

This function is used to re-enable interrupts managed by the RTOS - that is, controlled interrupts.

Note that nested calls to this function are permitted.

These interrupts can be disabled using the function **eSuspendOSInterrupts()**.

Parameters

None

Returns

OS_E_OK (successful)

OS_E_OSINTUNDERFLOW (too many resumes - unbalanced suspend/resumes)

8.2.4 Mutex Functions

This section provides descriptions of the RTOS API functions concerned with the mutex (Mutually Exclusive Activities) feature.

The functions are listed below, along with their page references:

Function	Page
OS_eEnterCriticalSection	91
OS_eExitCriticalSection	92

OS_eEnterCriticalSection

OS_teStatus OS_eEnterCriticalSection(OS_thMutex *hMutex*);

Description

This function is used at the start of a critical section of code, during which the currently running user task or ISR must not be pre-empted by another user task/ISR within the same “mutex” group (see below).

The function is paired with the function **OS_eExitCriticalSection()**, which is placed at the end of the critical section. Execution of the critical section will be allowed to complete before processing is switched to a higher priority task/ISR that may be pending within the same “mutex” group.

This mechanism, known as a “mutex”, applies to a group of user tasks or ISRs. Effectively, the mutex feature allows pre-set user task/ISR priorities to be over-ridden by temporarily assigning the highest priority in the mutex group to the task/ISR containing the critical section. The mutex feature is useful in allowing exclusive access to shared resources (e.g. hardware peripherals and shared memory). It allows the currently running user task/ISR to finish accessing a shared resource before processing is switched to a higher priority user task/ISR (from the same mutex group) that also needs to access the resource. In this way, access to a shared resource can be serialised.

Function calls to enter/exit the same critical section should not be nested. Nested calls to enter/exit different critical sections must be strictly nested. The critical section must be exited before termination of the task that uses it.

A handle for the mutex group is pre-defined using the JenOS Configuration Editor.

Parameters

<i>hMutex</i>	Handle of mutex group for critical section of code
---------------	--

Returns

- OS_E_OK (successful)
- OS_E_BADMUTEX (invalid mutex handle used)
- OS_E_BAD_NESTING (bad nesting)

OS_eExitCriticalSection

```
OS_teStatus OS_eExitCriticalSection(OS_thMutex hMutex);
```

Description

This function must be used at the end of a critical section of code during which the current user task or ISR cannot be pre-empted by another user task/ISR from the same mutex group.

The function is paired with the function **OS_eEnterCriticalSection()**, which is placed at the start of the critical section. Execution of the critical section will be allowed to complete before processing is switched to a higher priority user task/ISR that may be pending within the same mutex group.

The handle of the mutex group to which the critical section belongs must be specified in this function and must be the same as the handle used in the matching call to **OS_eEnterCriticalSection()**.

Function calls to enter/exit the same critical section should not be nested. Nested calls to enter/exit different critical sections must be strictly nested. The critical section must be exited before termination of the task that uses it.

Parameters

<i>hMutex</i>	Handle of mutex group for critical section of code
---------------	--

Returns

OS_E_OK (successful)
OS_E_BADMUTEX (invalid mutex handle used)
OS_E_BAD_NESTING (bad nesting)

8.2.5 Messaging Functions

This section provides descriptions of the RTOS API functions concerned with sending and collecting inter-task messages.

The functions are listed below, along with their page references:

Function	Page
OS_ePostMessage	94
OS_eCollectMessage	95
OS_eGetMessageStatus	96

OS_ePostMessage

**OS_teStatus OS_ePostMessage(OS_thMessage *hMessage*,
void **pvData*);**

Description

This function sends a message of the specified type, possibly containing user data, to another user task or ISR.

The message type must be defined in the header file **os_msg_types.h**. The identity of the destination task or ISR, and whether the message type is queued, are pre-configured using the JenOS Configuration Editor (see [Section 14.6.2](#)). The data content for a message type can be of arbitrary size, including zero length (no data).

- If the data content has non-zero length, the sent message may possibly be queued with other messages of the same type (depending on how the message type has been configured).
- If the data content has zero length, the sent message is not queued (regardless of how the message type has been configured).

If a sent message is not queued, the message overwrites any previous, uncollected message of the same type.

This function supports one-to-one and many-to-one communications, but does not support many-to-many communications.

Messages posted with this function are collected by the destination task using the function **OS_eCollectMessage()**.



Note: A message sent with this function can alternatively activate the destination task or invoke a callback function to perform user actions. These options are configured using the JenOS Configuration Editor.

Parameters

<i>hMessage</i>	Handle of message type
<i>*pvData</i>	Pointer to data to be sent in message. Set to NULL for a message with no data

Returns

OS_E_OK (successful)
OS_E_BADMESSAGE (invalid message type handle specified)
OS_E_QUEUE_FULL (message queue is full)

OS_eCollectMessage

```
OS_teStatus OS_eCollectMessage(  
    OS_thMessage hMessage,  
    void *pvData);
```

Description

This function is used to collect a posted message of the specified message type. Any extracted message data is placed in the specified location.

The function is used to collect a message sent using the function **OS_ePostMessage()**.

If the message is unqueued, it will always be successfully collected.

Parameters

<i>hMessage</i>	Handle of message type of message to collect
<i>*pvData</i>	Pointer to place where extracted message data will be stored. Set to NULL for a message with no data

Returns

- OS_E_OK (successful)
- OS_E_BADMESSAGE (invalid message handle specified)
- OS_E_BADVALUE (invalid data pointer specified)
- OS_E_QUEUE_EMPTY (message queue contains no messages)

OS_eGetMessageStatus

```
OS_teStatus OS_eGetMessageStatus(  
    OS_thMessage hMessage);
```

Description

This function is used to obtain the status of the incoming data message queue for the specified message type. The returned value indicates whether the queue contains messages or is empty.

Parameters

<i>hMessage</i>	Handle of message type for message queue to check
-----------------	---

Returns

- OS_E_BADMESSAGE (invalid message handle specified)
- OS_E_QUEUE_EMPTY (message queue contains no messages)
- OS_E_QUEUE_FULL (message queue contains uncollected messages)
- OS_E_UNQUEUED (message type is not queued)

8.2.6 Software Timer Functions

This section provides descriptions of the RTOS API functions concerned with the software timers that can be used to schedule the start of an activity within a user task or ISR.

The functions are listed below, along with their page references:

Function	Page
OS_eStartSWTimer	98
OS_eStopSWTimer	99
OS_eExpireSWTimers	100
OS_eContinueSWTimer	101
OS_eGetSWTimerStatus	102



Note 1: Some of the above software timer functions call user-defined callback functions that must be defined using macros described in [Section 8.1](#). The required callback functions and associated macros are mentioned in the function descriptions in this section.

Note 2: If the tick timer is used as the source counter and the maximum count of the tick timer is T (before the tick timer wraps around), there must be no more than $T/2$ ticks between consecutive software timer expiry events (e.g. if the tick timer wraps around every 60 seconds, a software timer must expire every 30 seconds or less).



Caution: *To allow the JN516x device to enter sleep mode, no software timers should be active. Any running software timers must first be stopped and any expired timers must be de-activated. Both can be achieved using the function **OS_eStopSWTimer()**, which must be called individually for each running and expired timer.*

OS_eStartSWTimer

```
OS_teStatus OS_eStartSWTimer(OS_thSWTimer hSWTimer,  
                              uint32 u32Ticks,  
                              void *pvData);
```

Description

This function is used to start the specified software timer and configure it to expire after the specified number of ticks. The software timer is defined and given a handle using the JenOS Configuration Editor (see [Section 14.5.2](#)). It is derived from a source counter - this could be a hardware counter, such as the on-chip tick timer, or another software timer. The source counter used determines the tick period and therefore the timed period.

In this function, you must specify the number of ticks of the source counter until the software timer expires. If the on-chip tick timer is used as the source counter, the specified value must not be greater than half the maximum count of the tick timer (when the tick timer wraps around) - since the tick timer is a 32-bit counter, you must not specify a value greater than 2147483647 ticks (0x7FFFFFFF ticks).

Using the JenOS Configuration Editor, it is possible to configure a task or callback function to be executed on expiry of the software timer. A parameter is provided which allows a pointer to data to be specified, which will be used by a callback function on expiry of the timer.

This function calls the user-defined 'hardware counter enable' callback function, defined using the macro **OS_HWCOUNTER_ENABLE_CALLBACK()**, provided that there are no other software timers already active that use the hardware counter. The function also calls the following user-defined callback functions:

- 'hardware counter get' function, defined using the macro **OS_HWCOUNTER_GET_CALLBACK()**, which obtains the current value of the hardware counter
- 'hardware counter set' function, defined using the macro **OS_HWCOUNTER_SET_CALLBACK()**, which sets the hardware counter's compare register to the appropriate value for the timed duration (equal to the specified number of ticks added to the current value of the hardware counter previously obtained)

Parameters

<i>hSWTimer</i>	Handle of software timer to start
<i>u32Ticks</i>	The number of ticks before software timer expires (if using on-chip tick timer, refer to above description)
<i>*pvData</i>	Pointer to data to be passed to optional callback function on expiry of timer (NULL if not required). Ignored for tasks

Returns

OS_E_OK (successful)
OS_E_BADSWTIMER (invalid software timer handle)
OS_E_SWTIMER_RUNNING (software timer already running)

OS_eStopSWTimer

```
OS_teStatus OS_eStopSWTimer(OS_thSWTimer hSWTimer);
```

Description

This function is used to de-activate the specified software timer - this can be a running timer or an expired timer.

This function calls the user-defined 'hardware counter disable' callback function, defined using the macro **OS_HWCOUNTER_DISABLE_CALLBACK()**, provided that there are no other software timers still active that use the hardware counter.

Parameters

<i>hSWTimer</i>	Handle of software timer to de-activate
-----------------	---

Returns

- OS_E_OK (successful)
- OS_E_BADSWTIMER (invalid software timer handle)
- OS_E_SWTIMER_STOPPED (software timer already de-activated)

OS_eExpireSWTimers

```
OS_teStatus OS_eExpireSWTimers(  
    OS_thHWCCounter hHWCCounter);
```

Description

This function is used to expire scheduled software timers associated with the specified source counter, where these timers have previously reached their target counts (source counter has matched the value in the counter's compare register).

The function should be called in the source counter's ISR that is invoked when the above match occurs. The function:

- sets the status of the software timer to 'expired' by calling the user-defined callback function that is defined using the macro **OS_SWTIMER_CALLBACK()**
- checks whether there are any other pending software timers and:
 - if there is at least one pending software timer, the function updates the source counter's compare register with the required number of ticks (until the next software timer expires) by calling the user-defined callback function that is defined using the macro **OS_HWCOUNTER_SET_CALLBACK()**
 - if there are no pending software timers, the function does nothing to the source counter (leaves it running)

Before this function returns, it deals with any software timers that expire in quick succession (without the need to re-call the function for the individual software timers).

For more information on the software timers, refer to [Section 2.4.6](#).

Parameters

hHWCCounter Handle of counter

Returns

OS_E_OK (successful)
OS_E_BADHWCOUNTER (invalid counter handle)
OS_E_NOTHINGTOEXPIRE (no software timers to expire)
OS_E_OVERACTIVATION (associated task activated too many times)
OS_E_BADTASK (invalid task handle used)
OS_E_HWCOUNTERIDLE (no active timers)

OS_eContinueSWTimer

```
OS_teStatus OS_eContinueSWTimer(  
    OS_thSWTimer hSWTimer,  
    uint32 u32Ticks,  
    void *pvData);
```

Description

This function is used to re-start the specified software timer, if it has previously expired, and configure it to expire again after the specified number of ticks. The timer will be re-started without any discontinuity between the previous expiry point and the new run, provided that the function is executed before the next timer cycle is complete.

The function is designed to be used immediately following the expiry of the timer. It should not be called a long time after the previous expiry (0x7FFFFFFF ticks).

OS_eStartSWTimer() should be used in these circumstances.

In this function, you must specify the number of ticks of the source counter until the software timer expires. If the on-chip tick timer is used as the source counter, the specified value must not be greater than half the maximum count of the tick timer (when the tick timer wraps around) - since the tick timer is a 32-bit counter, you must not specify a value greater than 2147483647 ticks (0x7FFFFFFF ticks).

It is possible to specify a pointer to data which will be used if the timer is configured to invoke a callback function on expiry (the callback function is specified using the JenOS Configuration Editor).

If the specified software timer will be the next to expire, this function calls the user-defined 'hardware counter set' callback function, defined using the macro

OS_HWCOUNTER_SET_CALLBACK(), which sets the hardware counter's compare register to the appropriate value for the remainder of the timed duration.

Parameters

<i>hSWTimer</i>	Handle of software timer to re-start
<i>u32Ticks</i>	The number of ticks before software timer expires
<i>*pvData</i>	Pointer to data to be passed to optional callback function on expiry of timer (NULL if not required)

Returns

OS_E_OK (successful)
OS_E_BADSWTIMER (invalid software timer handle)
OS_E_SWTIMER_RUNNING (software timer already running)

OS_eGetSWTimerStatus

```
OS_teStatus OS_eGetSWTimerStatus(  
    OS_thSWTimer hSWTimer);
```

Description

This function is used to obtain the status of the specified software timer. The timer is reported as running, stopped or expired.

Parameters

<i>hSWTimer</i>	Handle of software timer
-----------------	--------------------------

Returns

OS_E_BADSWTIMER (invalid software timer handle)
OS_E_SWTIMER_RUNNING (software timer is running)
OS_E_SWTIMER_STOPPED (software timer has been stopped)
OS_E_SWTIMER_EXPIRED (software timer has expired)

9. PDM API for Flash Memory

This chapter describes the functions of the JenOS Persistent Data Manager (PDM) API that supports context data and application data saving in Flash memory which is external to the JN516x device. This edition of the PDM is supplied in the JN516x ZigBee Smart Energy SDK (JN-SW-4064).

For the later edition of the PDM (supplied in other SDKs) that supports JN516x internal EEPROM, refer to [Chapter 10](#).



Tip: In this chapter, the storage medium for persisted data is referred to as Non-Volatile Memory (NVM) but in practice it is SPI-connected external Flash memory.

The API is defined in the header file **pdm.h**.

The PDM API functions are listed below, along with their page references:

Function	Page
PDM_vInit	104
PDM_vSPIFlashConfig	106
PDM_eLoadRecord	107
PDM_vSaveRecord	109
PDM_vSave	110
PDM_vDeleteRecord	111
PDM_vDelete	112
PDM_vWarmInitHw	113
PDM_vRegisterSystemCallback	114



Caution: When using the PDM, do not use the JN516x Integrated Peripherals API to interact with the Flash memory device connected to the JN516x chip.

PDM_vlnit

```
void PDM_vlnit(uint8 u8StartSector,  
               uint8 u8NumSectors,  
               uint32 u32SectorSize,  
               OS_thMutex hPdmMutex,  
               OS_thMutex hPdmMediaMutex,  
               PDM_tsHwFuncTable *psHwFuncTable,  
               const tsReg128 *psKey);
```

Description

This function initialises the PDM module, and must be called during a cold start.

The function takes as input details of the sectors of the NVM (Non-Volatile Memory) device to be managed by the module, as well as a set of functions that the PDM will use to interact with the NVM device. These functions are specified in the structure detailed in [Section 14.1](#).

Optional mutexes can be specified in order to:

- Serialise PDM function calls - if specified, this mutex is automatically applied during a PDM function call to prevent concurrent PDM function calls
- Serialise SPI bus access - if specified, this mutex is automatically applied during an access to NVM via the SPI bus to prevent concurrent accesses to the SPI bus (useful if other resources are also accessible via the SPI bus)

If a mutex is used, the user task must be linked to the relevant mutex in the JenOS Configuration Editor - refer to [Section 14.7](#).

The function also allows you to specify a key to be used by the PDM module to encrypt and decrypt saved data. An option is available to use a key stored in eFuse. Security based on this encryption key is applied to the context data that is automatically stored/restored by the stack, but security must be enabled for individual application data records when **PDM_eLoadRecord()** is called.

PDM_vlnit() will auto-detect the NVM device type (manufacturer/model), unless you have already called **PDM_vSPIFlashConfig()** which allows you to specify a specific or custom SPI Flash device type. Note that if you have specified a set of NVM device functions in a call to **PDM_vSPIFlashConfig()**, you can set a NULL pointer to these functions in **PDM_vlnit()**.

Parameters

<i>u8StartSector</i>	Number of the first sector of NVM to be managed by PDM module
<i>u8NumSectors</i>	Number of contiguous sectors of NVM to be managed by PDM module
<i>u32SectorSize</i>	Size of each sector, in bytes
<i>hPdmMutex</i>	Optional handle of the mutex to be used to serialise PDM function calls
<i>hPdmMediaMutex</i>	Optional handle of the mutex to be used to serialise access to NVM via the SPI bus.

<i>*psHwFuncTable</i>	Pointer to set of custom functions to be used with NVM device (see Section 14.1). Set to NULL if not needed
<i>*psKey</i>	Pointer to structure containing the 128-bit encryption key to be used by the PDM module (see Section 14.5). A NULL pointer indicates that a key stored in eFuse is to be used (if no key is stored, a zero value will be used).

Returns

None

PDM_vSPIFlashConfig

```
void PDM_vSPIFlashConfig(  
    teFlashChipType eFlashType,  
    tSPIflashFuncTable *psFlashFuncTable);
```

Description

This function should be used if the NVM device is an SPI Flash device and you do not wish to auto-detect the type of Flash device (manufacturer/model) - for example, if you are using an unsupported or custom Flash device.

If you wish to auto-detect the device type, you do not need to call this function, as auto-detection is implemented by default.

If required, this function must be called before **PDM_vInit()**.

The function requires you to specify the Flash device type. If a custom Flash device is selected (E_FL_CHIP_CUSTOM), you also need to specify a table of custom functions for the device. The structure through which these functions are specified are described in [Section 14.2](#).

Parameters

<i>eFlashType</i>	Flash device type, one of: E_FL_CHIP_ST_M25P10_A E_FL_CHIP_SST_25VF010 E_FL_CHIP_ATMEL_AT25F512 E_FL_CHIP_CUSTOM E_FL_CHIP_AUTO
<i>*psFlashFuncTable</i>	Pointer to custom function table (see Section 14.2), for case when a custom Flash device has been selected (E_FL_CHIP_CUSTOM). If a supported device has been selected, set this pointer to NULL.

Returns

None

PDM_eLoadRecord

```
PDM_teStatus PDM_eLoadRecord(  
    PDM_tsRecordDescriptor *psDesc,  
    uint16 u16IdValue,  
    void *pvData,  
    uint32 u32DataSize,  
    bool_t bSecure);
```

Description

This function is used during a cold start to load an individual record of application data from NVM into RAM, or to create an application data record in NVM:

- During a first-time cold start, the function defines a record to be created in NVM.
- During any subsequent cold start, the function restores (into RAM) application data previously stored in the NVM record.

The function must be called before calling any function which automatically saves application data to NVM - for example, before calling **ZPS_eAplAflnit()** to initialise the ZigBee PRO stack and before calling **ZPS_vAplSecSetInitialSecurityState()** to initialise the ZigBee security state on the node.

A pointer to a record descriptor and a unique ID value for the record must be specified. In addition, a pointer must be provided to the start of the corresponding data buffer in RAM and the data size must also be specified.



Caution: The application software must not use record identifier values that would clash with those used by the NXP libraries used with the application. The ZigBee PRO stack libraries use values above 0x8000.

When called, the function first checks whether the specified record already exists in NVM.

- If the record does exist, the data from the NVM record is loaded into RAM.
- If the record does not exist, the record will be created in NVM the next time **PDM_vSaveRecord()** or **PDM_vSave()** is called. In this case, the specified RAM buffer contents remain unchanged and are saved to the NVM record.

The function also allows the record to be secured when saving to external SPI Flash memory. If this option is enabled, data saved to the NVM record will be encrypted using the key specified in **PDM_vInit()** and the data will be decrypted using the same key when it is read from the record.

Note that this function must not be called after **PDM_vSaveRecord()** or **PDM_vSave()**, and must not be called during a warm start (for example, following sleep with memory held). Otherwise, the latest data in RAM will be overwritten.

Parameters

<i>*psDesc</i>	Pointer to record descriptor (you do not need to be concerned with the contents of this descriptor)
<i>u16IdValue</i>	Unique 16-bit record identifier (see Caution above)
<i>*pvData</i>	Pointer to start of RAM buffer in which to store data (in the case of record creation, initial data must be provided in the buffer)
<i>u32DataSize</i>	Size of record, in bytes
<i>bSecure</i>	Enable/disable data encryption for record: TRUE: Enable encryption FALSE: Disable encryption

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)

PDM_vSaveRecord

```
void PDM_vSaveRecord(PDM_tsRecordDescriptor *psDesc);
```

Description

This function saves the specified application data record from RAM to NVM.

Following a cold start, the function should only be called after all records have been created or loaded using **PDM_eLoadRecord()**.

The application data will be saved encrypted if security was enabled for the NVM record when the function **PDM_eLoadRecord()** was called.

Alternatively, you can save all records in RAM to NVM using the function **PDM_vSave()**.

Parameters

<i>*psDesc</i>	Pointer to descriptor of record to be saved to NVM
----------------	--

Returns

None

PDM_vSave

```
void PDM_vSave(void);
```

Description

This function saves all records, including both application data and stack context data, from RAM to NVM.

Following a cold start, the function should only be called after all application data records have been created or loaded using **PDM_eLoadRecord()**.

The stack context data is saved encrypted using the security key specified when **PDM_vInit()** was called. The application data will be saved encrypted using this key only if security was enabled for the corresponding NVM record when **PDM_eLoadRecord()** was called.

Alternatively, an individual application data record can be saved to NVM using the function **PDM_vSaveRecord()**.

Parameters

None

Returns

None

PDM_vDeleteRecord

```
void PDM_vDeleteRecord(  
    PDM_tsRecordDescriptor *psDesc);
```

Description

This function deletes the specified record of application data in NVM.

Alternatively, all records in NVM can be deleted using the function **PDM_vDelete()**.

Parameters

<i>*psDesc</i>	Pointer to descriptor of record to be deleted
----------------	---

Returns

None

PDM_vDelete

```
void PDM_vDelete(void);
```

Description

This function deletes all records in NVM, including both application data and stack context data.



Caution: You are not recommended to delete records of stack context data before a rejoin of the same secured network. If these records are deleted, data sent by the node after the rejoin will be rejected by the destination node since the frame counter has been reset on the source node. For more details, refer to “Application Design Notes” appendix in the ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101).

Alternatively, an individual record of application data can be deleted using the function **PDM_vDeleteRecord()**.

Parameters

None

Returns

None

PDM_vWarmInitHw

```
void PDM_vWarmHwInit(void);
```

Description

This function initialises the SPI Flash device following a warm start. The function must be called immediately following a warm start, before calling **OS_vRestart()**. Otherwise the ZigBee PRO stack may attempt to save records before the SPI hardware is ready.

Parameters

None

Returns

None

PDM_vRegisterSystemCallback

```
void PDM_vRegisterSystemCallback(  
    PDM_tpfvSystemEventCallback  
    fpvPDM_SystemEventCallback);
```

Description

This function registers a user-defined callback function to handle PDM errors and events.

Parameters

<i>fpvPDM_SystemEventCallback</i>	Pointer to the application callback function. The function type PDM_tpfvSystemEventCallback is documented in Section 14.6 . The events generated by the PDM library are documented in Section 14.7
-----------------------------------	---

Returns

None

10. PDM API for EEPROM

This chapter details the functions of the JenOS Persistent Data Manager (PDM) API that supports context data and application data saving in JN516x EEPROM, and is supplied in the following SDKs:

- JN-SW-4168: JN516x ZigBee Light Link and Home Automation SDK
- JN-SW-4165: JN516x JenNet-IP SDK
- JN-SW-4163: JN516x IEEE802.15.4 SDK

For the earlier edition of the PDM (supplied in other SDKs) that only supports external Flash memory devices, refer to [Chapter 12](#).

The API is defined in the header file **pdm.h** and is divided into the following categories:

- EEPROM PDM functions - see [Section 10.1](#)
- EEPROM PDM Bitmap Counter functions - see [Section 10.2](#)
- EEPROM PDM miscellaneous functions - see [Section 10.3](#)



Note: For more information on how to use the functions described in this chapter, refer to [Chapter 4](#).

10.1 EEPROM PDM Functions

The EEPROM PDM functions are listed below, along with their page references:

Function	Page
PDM_eInitialise	117
PDM_eSaveRecordData	118
PDM_eReadDataFromRecord	119
PDM_eDeleteData	120
PDM_eDeleteAllData	121
PDM_u8GetSegmentCapacity	122
PDM_u8GetSegmentOccupancy	123
PDM_bDoesDataExist	124



Note 1: For a description of how to use these functions, refer to [Section 4.3](#).

Note 2: Unlike the earlier PDM module detailed in [Chapter 9](#), the PDM module detailed in this section does not use descriptor-based records for data storage. This results in more efficient use of storage space and quicker operation.

PDM_eInitialise

```
PDM_teStatus PDM_eInitialise(  
    uint8 u8NumberOfEEPROMsegments  
#ifndef PDM_NO_RTOS  
    ,  
    OS_thMutex hPdmMutex  
#endif);
```

Description

This function initialises the PDM module and registers the required PDM functions. It must be called during both a warm start and a cold start.

The function initialises the PDM environment and builds the underlying EEPROM file system. A RAM-based file system is created to allow the PDM to map data to/from the EEPROM. The EEPROM sectors are scanned for evidence of any valid user data, which is mapped into the RAM file system. This routine handles any write errors that may have occurred if the EEPROM was powered down whilst data was being written to the PDM system. Once the file system has been constructed, you can then write data to and read data from the EEPROM via PDM.

The PDM can operate within any number of EEPROM segments, as specified through the parameter *u8NumberOfEEPROMsegments*. However, if a zero value is specified for this parameter, the function will auto-configure the PDM by interrogating the JN516x chip to obtain the variant and scaling the PDM accordingly, giving the application access to the full EEPROM.

An optional mutex can be specified in order to serialise PDM function calls. If specified, this mutex is automatically applied during a PDM function call to prevent concurrent calls. If the mutex is used with the JenOS RTOS in a ZigBee application, the user task must be linked to the relevant mutex in the JenOS Configuration Editor - refer to [Section 14.7](#). Note that when using the PDM without the JenOS RTOS:

- The mutex is not available when using the PDM in applications developed with the IEEE802.15.4 SDK (JN-SW-4163), in which case the flag `PDM_NO_RTOS` must be defined in the makefile - the function parameter *hPdmMutex* is then disabled.
- The mutex is always implemented when using the PDM in applications developed with the JenNet-IP SDK (JN-SW-4165), in which case the function parameter *hPdmMutex* must be set to a non-zero value.

For more information on using the PDM without the RTOS, refer to [Section 4.2](#).

Parameters

<i>u8NumberOfEEPROMsegments</i>	Number of contiguous EEPROM sectors to be managed. A zero value indicates that the full EEPROM should be used.
<i>hPdmMutex</i>	Optional handle of the mutex to be used to serialise PDM calls

Returns

PDM_E_STATUS_OK
PDM_E_STATUS_INTERNAL_ERROR

PDM_eSaveRecordData

```
PDM_teStatus PDM_eSaveRecordData(  
    uint16 u16IdValue,  
    uint8 *pu8DataBuffer,  
    uint16 u16DataLength);
```

Description

This function saves the specified application data from RAM to the specified record in EEPROM. The record is identified by means of a 16-bit user-defined value.



Caution: The application software must not use record identifier values that would clash with those used by the NXP libraries used with the application. The ZigBee PRO stack libraries use values above 0x8000. The JenNet-IP libraries use values between 0x3000 and 0x3007.

When a data record is saved to the EEPROM for the first time, the data is written provided there are enough EEPROM segments available to hold the data. Upon subsequent save requests, if there has been a change between the RAM-based and EEPROM-based data buffers then the PDM will attempt to re-save only the segments that have changed (if no data has changed, no save will be performed). This is advantageous due to the restricted size of the EEPROM and the constraint that old data must be preserved while saving changed data to the EEPROM.

Provided that you have registered a callback function with the PDM (see [Section 10.3](#)), the callback mechanism will signal when a save has failed. Upon failure, the callback function will be invoked and pass the event `E_PDM_SYSTEM_EVENT_DESCRIPTOR_SAVE_FAILED` to the application.

Parameters

<i>u16IdValue</i>	User-defined ID of the record to be saved (see Caution above)
<i>*pu8DataBuffer</i>	Pointer to data buffer to be saved in the record in EEPROM
<i>u16DataLength</i>	Length of data to be saved, in bytes

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (specified record ID is invalid)
PDM_E_STATUS_NOT_SAVED (save to EEPROM failed)

PDM_eReadDataFromRecord

```
PDM_teStatus PDM_eReadDataFromRecord(  
    uint16 u16IdValue,  
    void *pvDataBuffer,  
    uint16 u16DataBufferLength,  
    uint16 *pu16DataBytesRead);
```

Description

This function reads the specified record of application data from the EEPROM and stores the read data in the supplied data buffer in RAM. The record is specified using its unique 16-bit identifier.

Before calling this function, it may be useful to call **PDM_bDoesDataExist()** in order to determine whether a record with the specified identifier exists in the EEPROM and, if it does, to obtain its size.

Parameters

<i>u16IdValue</i>	User-defined ID of the record to be read
<i>*pvDataBuffer</i>	Pointer to the data buffer in RAM where the read data is to be stored
<i>u16DataBufferLength</i>	Length of the data buffer, in bytes
<i>*pu16DataBytesRead</i>	Pointer to a location to receive the number of data bytes read

Returns

PDM_E_STATUS_OK (success)

PDM_E_STATUS_INVLD_PARAM (specified record ID is invalid)

PDM_eDeleteData

```
PDM_teStatus PDM_eDeleteData(uint16 u16IdValue);
```

Description

This function deletes the specified record of application data in EEPROM.
Alternatively, all records in EEPROM can be deleted using the function **PDM_eDeleteAllData()**.

Parameters

<i>u16IdValue</i>	User-defined ID of the record to be deleted
-------------------	---

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (specified record ID is invalid)

PDM_eDeleteAllData

```
PDM_tStatus PDM_eDeleteAllData(void);
```

Description

This function deletes all records in EEPROM, including both application data and stack context data, resulting in an empty PDM file system. The EEPROM segment Wear Count values are preserved (and incremented) throughout this function call.



Caution: You are not recommended to delete records of stack context data before a rejoin of the same secured network. If these records are deleted, data sent by the node after the rejoin will be rejected by the destination node since the frame counter has been reset on the source node. For more details, refer to “Application Design Notes” appendix in the ZigBee PRO Stack User Guide (JN-UG-3101).

Alternatively, an individual record of application data can be deleted using the function **PDM_eDeleteData()**.

Parameters

None

Returns

None

PDM_u8GetSegmentCapacity

```
uint8 PDM_u8GetSegmentCapacity(void);
```

Description

This function returns the number of unused segments that remain in the EEPROM.

Parameters

None

Returns

Number of EEPROM segments free

PDM_u8GetSegmentOccupancy

`uint8 PDM_u8GetSegmentOccupancy(void);`

Description

This function returns the number of used segments in the EEPROM.

Parameters

None

Returns

Number of EEPROM segments used

PDM_bDoesDataExist

```
bool_t PDM_bDoesDataExist(uint16 u16IdValue,  
                           uint16 *pu16DataLength);
```

Description

This function checks whether data associated with the specified record ID exists in the EEPROM. If the data record exists, the function returns the data length, in bytes, in a location to which a pointer must be provided.

Parameters

<i>u16IdValue</i>	User-defined ID of the record to be found
<i>*pu16DataLength</i>	Pointer to location to receive length, in bytes, of data record (if any) associated with specified record ID

Returns

TRUE if data record found, FALSE otherwise

10.2 EEPROM PDM Bitmap Counter Functions

The EEPROM PDM Bitmap Counter functions are listed below, along with their page references:

Function	Page
PDM_eCreateBitmap	126
PDM_eIncrementBitmap	127
PDM_eGetBitmap	128
PDM_eDeleteBitmap	129



Note: For a description of how to use these functions, refer to [Section 4.4](#).

PDM_eCreateBitmap

```
PDM_teStatus PDM_eCreateBitmap(uint16 u16IdValue,  
                                uint32 u32InitialValue);
```

Description

The function creates a bitmap structure for a counter in a segment of the EEPROM. A user-defined ID and a start value for the bitmap counter must be specified.

The start value is stored in the counter's header. A bitmap is created to store the incremental value of the counter (over the start value). This bitmap can subsequently be incremented (by one) by calling the function **PDM_eIncrementBitmap()**. The incremental value stored in the bitmap and the start value stored in the header can be read at any time using the function **PDM_eGetBitmap()**.

If the specified ID value has already been used or the specified start value is NULL, the function returns PDM_E_STATUS_INVLD_PARAM. If the EEPROM has no free segments, the function returns PDM_E_STATUS_USER_PDM_FULL.

Parameters

<i>u16IdValue</i>	User-defined ID for bitmap counter to be created
<i>u32InitialValue</i>	Initial 32-bit value of bitmap counter

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)
PDM_E_STATUS_PDM_FULL (there is no space to store this bitmap)

PDM_eIncrementBitmap

```
PDM_teStatus PDM_eIncrementBitmap(uint16 u16IdValue);
```

Description

The function increments the bitmap value of the specified counter in the EEPROM. The counter must be identified using the user-defined ID value assigned when the counter was created using the function **PDM_eCreateBitmap()**.

The bitmap can be incremented within an EEPROM segment until its value saturates (contains all 1s). At this point, the function returns the code **PDM_E_STATUS_SATURATED_OK**. The next time that this function is called, the counter is automatically moved to a new segment (provided that one is available), the start value in its header is increased appropriately and the bitmap is reset to zero. To avoid increasing the segment Wear Count, the old segment is not formally deleted before a new segment is started. If the EEPROM has no free segments when the above overflow occurs, the function returns the code **PDM_E_STATUS_USER_PDM_FULL**.

If the specified ID value has already been used, the function returns **PDM_E_STATUS_INVLD_PARAM**.

Parameters

<i>u16IdValue</i>	User-defined ID of counter to be incremented
-------------------	--

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)
PDM_E_STATUS_PDM_FULL (no further EEPROM segments for the bitmap)
PDM_E_STATUS_BITMAP_SATURATED_OK (increment made but segment now saturated)

PDM_eGetBitmap

```
PDM_teStatus PDM_eGetBitmap(uint16 u16IdValue,  
                             uint32 *pu32InitialValue,  
                             uint32 *pu32BitmapValue);
```

Description

The function reads the specified counter value from the EEPROM. The counter must be identified using the user-defined ID value assigned when the counter was created using the function **PDM_eCreateBitmap()**. The function returns the counter's start value (from the counter's header) and incremental value (from the counter's bitmap).

The counter value is calculated as:

$$\text{Start Value} + \text{Incremental Value}$$

or in terms of the function parameters:

$$*pu32InitialValue + *pu32BitmapValue$$

Note that the start value may be different from the one specified when the counter was created, as the start value is updated each time the counter outgrows a segment and the bitmap is reset to zero.

This function should be called when the device comes up from a cold start, to check whether a bitmap counter is present in EEPROM.

If the specified ID value has already been used or a NULL pointer is provided for the received values, the function returns PDM_E_STATUS_INVLD_PARAM.

Parameters

<i>u16IdValue</i>	User-defined ID for bitmap counter to be accessed
<i>*pu32InitialValue</i>	Pointer to location to receive the start value of the counter
<i>*pu32BitmapValue</i>	Pointer to location to receive the incremental value of the counter

Returns

PDM_E_STATUS_OK (success)

PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)

PDM_eDeleteBitmap

```
PDM_teStatus PDM_eDeleteBitmap(uint16 u16IdValue);
```

Description

This function deletes the specified counter in the EEPROM. The counter must be identified using the user-defined ID value assigned when the bitmap was created using the function **PDM_eCreateBitmap()**.

The function can be used to formally delete a counter. It clears the current segment occupied by the counter and also all the older (expired) segments used for the counter. This deletion increments the Wear Counts for these segments and should be done only if absolutely necessary, as the expired segments can be re-used directly via the PDM without formal deletion.

If the specified ID value does not exist in the PDM, the function returns PDM_E_STATUS_INVLD_PARAM.

Parameters

<i>u16IdValue</i>	User-defined ID for bitmap counter to be deleted
-------------------	--

Returns

PDM_E_STATUS_OK (success)

PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)

10.3 EEPROM PDM Miscellaneous Functions

The EEPROM PDM miscellaneous functions include a function for registering a user-defined PDM system callback function and functions related to the Wear Counts of EEPROM segments. The functions are listed below, along with their page references:

Function	Page
PDM_vRegisterSystemCallback	131
PDM_vSetWearCountTriggerLevel	132
PDM_eGetSegmentWearCount	133



Note: For a description of how to use these functions, refer to [Section 4.5.2](#) and [Section 4.5.4](#).

PDM_vRegisterSystemCallback

```
void PDM_vRegisterSystemCallback(  
    PDM_tpfvSystemEventCallback  
        fpvPDM_SystemEventCallback);
```

Description

This function registers a user-defined callback function to handle PDM events and errors.

Parameters

<i>fpvPDM_SystemEventCallback</i>	Pointer to the application callback function. The function type PDM_tpfvSystemEventCallback is documented in Section 14.6 . The events generated by the PDM library are documented in Section 14.7
-----------------------------------	---

Returns

None

PDM_vSetWearCountTriggerLevel

```
void PDM_vSetWearCountTriggerLevel(  
    uint32 u32WearCountTriggerLevel);
```

Description

This function sets the Wear Count value of an EEPROM segment at which a Wear Count event will be triggered and the PDM callback function will be activated. The invoked callback function is user-defined and is registered using the function **PDM_vRegisterSystemCallback()**.

The callback function will only be invoked once for a particular segment, when the specified Wear Count value occurs (it will not be invoked for every occurrence afterwards when the segment Wear Count exceeds the trigger value).

Parameters

u32WearCountTriggerLevel Wear Count value that triggers a Wear Count event

Returns

None

PDM_eGetSegmentWearCount

```
PDM_teStatus PDM_eGetSegmentWearCount(  
    uint8 u8SegmentIndex,  
    uint32 *pu32WearCount);
```

Description

This function obtains the current Wear Count value of the specified EEPROM segment.

Parameters

<i>u8SegmentIndex</i>	Index of EEPROM segment for which Wear Count needed
<i>pu32WearCount</i>	Pointer to location to receive obtained Wear Count value

Returns

PDM_E_STATUS_OK (success)
PDM_E_STATUS_INVLD_PARAM (an invalid parameter value was supplied)

11. PWRM API

This chapter describes the functions of the JenOS Power Manager (PWRM) API. The API is defined in the header file **pwrn.h**.



Caution: *The Power Manager uses Wake Timer 1 of the JN516x device if scheduled wake events are configured. In this case, do not use this wake timer for any other purpose in your application.*

The PWRM API functions are divided into the following categories:

- 'Core' functions, described in [Section 11.1](#)
- 'Callback Set-up' functions, described in [Section 11.2](#)
- 'Debugging' functions, described in [Section 11.3](#)

11.1 Core Functions

The PWRM core functions are listed below, along with their page references:

Function	Page
PWRM_vInit	136
PWRM_eStartActivity	137
PWRM_eFinishActivity	138
PWRM_u16GetActivityCount	139
PWRM_eScheduleActivity	140
PWRM_vManagePower	141

PWRM_vInit

```
void PWRM_vInit(PWRM_tePowerMode ePowerMode);
```

Description

This function is used to initialise the Power Manager and specify the low-power mode in which the JN516x device should be put when inactive.

There are five possible low-power modes that can be specified:

- Sleep with 32-kHz oscillator running and memory held
- Sleep with 32-kHz oscillator running and memory not held
- Sleep with 32-kHz oscillator not running and memory held
- Sleep with 32-kHz oscillator not running and memory not held
- Deep Sleep (32-kHz oscillator not running and memory not held)

The enumerations for the above power modes are listed below and described in [Section 14.3](#). For further information on these low-power modes and how to wake from them, refer to [Section 5.1](#).

Note that if the Power Manager is unable to put the JN516x device into the specified low-power mode, it will put the device into Doze mode instead - see description of **PWRM_vManagePower()**.

If the 32-kHz oscillator is run, the JN516x device's Wake Timer 1 is calibrated and made available (and then must not be used for any other purpose).

Parameters

<i>ePowerMode</i>	The power mode to be used during sleep, one of: PWRM_E_SLEEP_OSCON_RAMON PWRM_E_SLEEP_OSCON_RAMOFF PWRM_E_SLEEP_OSCOFF_RAMON PWRM_E_SLEEP_OSCOFF_RAMOFF PWRM_E_SLEEP_DEEP
-------------------	--

Returns

None

PWRM_eStartActivity

PWRM_teStatus PWRM_eStartActivity(void);

Description

This function is used to notify the Power Manager that an activity has been started which must not be interrupted by sleep. Thus, while such an activity is running, the JN516x device will not enter sleep mode.

The function **PWRM_eFinishActivity()** must then be called when the activity has completed. However, if **PWRM_eStartActivity()** has also been called for other activities that have not yet finished, the device will not be able to enter sleep mode until **PWRM_eFinishActivity()** has been called for all such activities.

The activity for which **PWRM_eStartActivity()** is called does not need to be identified, since the function simply increments a counter of running activities that must not be interrupted by sleep. There is an upper limit of 64K to the value of this counter. If this limit is exceeded, an overflow error is returned.

Parameters

None

Returns

PWRM_E_OK (success)

PWRM_E_ACTIVITY_OVERFLOW (activity counter limit exceeded)

PWRM_eFinishActivity

PWRM_teStatus PWRM_eFinishActivity(void);

Description

This function is used to notify the Power Manager that an activity has completed which was not to be interrupted by sleep.

The function call must be paired with a previous call to **PWRM_eStartActivity()**. Sleep mode cannot be entered until **PWRM_eFinishActivity()** has been called for all activities for which **PWRM_eStartActivity()** has been previously called.

The activity for which **PWRM_eFinishActivity()** is called does not need to be identified, since the function simply decrements a counter of running activities that must not be interrupted by sleep. Sleep mode must not be entered until this counter reaches zero. If this function is called when the counter is already zero, an underflow error is returned.

Parameters

None

Returns

PWRM_E_OK (success)

PWRM_E_ACTIVITY_UNDERFLOW (activity counter already zero)

PWRM_u16GetActivityCount

```
uint16 PWRM_u16GetActivityCount(void);
```

Description

This function obtains the current value of the activity counter which indicates the number of activities currently running that must not be interrupted by sleep. Sleep mode cannot be entered until the value of this counter is zero.

Parameters

None

Returns

Current value of activity counter

PWRM_eScheduleActivity

```
PWRM_teStatus PWRM_eScheduleActivity(  
    pwrm_tsWakeTimerEvent *psWake,  
    uint32 u32Ticks,  
    void (*prCallbackfn)(void));
```

Description

This function can be used to add a wake point and associated callback function to a list of scheduled wake points and callbacks. The new wake point is linked to an exclusive 32-kHz software wake timer, through the specified structure.

The function takes as input the number of ticks of the wake timer until the scheduled wake point. When the wake timer expires, the JN516x device will be woken from sleep and the specified callback function will be called.

To use this function, the Power Manager must be configured through **PWRM_vlnit()** to implement a low-power mode in which the 32-kHz oscillator is running and memory is held (otherwise, the list of scheduled wake points will be lost when the device enters sleep mode).

The function will return an error (see below) if the 32-kHz oscillator has not been configured to run during sleep or the software wake timer is already running for another wake point.

Parameters

<i>*psWake</i>	Pointer to a structure to be populated with the wake point and callback function (see below)
<i>u32Ticks</i>	The number of ticks of the 32-kHz wake timer until wake point
<i>*prCallbackfn</i>	Pointer to callback function associated with wake point

Returns

PWRM_E_OK (wake timer started successfully)
PWRM_E_TIMER_RUNNING (wake timer already running for another wake point)
PWRM_E_TIMER_INVALID (oscillator not configured to run during sleep)

PWRM_vManagePower

void PWRM_vManagePower(void);

Description

This function instructs the Power Manager to manage the power state of the JN516x device. The device must be idle when this function is called, i.e. the function is typically called from the OS idle task.

Once this function has been called, whenever appropriate, the Power Manager will put the device into the low-power mode specified through the function **PWRM_vInit()**. To allow the device to enter sleep mode:

- No activities that are uninterruptable by sleep must be running - that is, the activity counter must be zero.
- If the 32-kHz oscillator will run during sleep, a wake point must have been scheduled using **PWRM_vScheduleActivity()** (this condition does not apply when the oscillator is not used)

If the above two conditions are not satisfied, the function will put the device into Doze mode instead of sleep mode. Doze mode simply pauses the on-chip CPU, leaving all components powered (e.g. radio), and requires an interrupt to be configured to wake the device.

Before putting the device into sleep mode, this function calls any user-defined callback functions that have been registered using the function **PWRM_vRegisterPreSleepCallback()**.

Parameters

None

Returns

None

11.2 Callback Set-up Functions

The PWRM callback set-up functions are used to introduce user-defined callback functions that must be defined when using the Power Manager.

The functions are listed below, along with their page references:

Function	Page
vAppMain	143
PWRM_vRegisterPreSleepCallback	144
PWRM_vRegisterWakeupCallback	145
vAppRegisterPWRMCallbacks	146
PWRM_vWakeInterruptCallback	147

vAppMain

```
void vAppMain(void);
```

Description

This is a user-defined callback function which is the application entry point when using the Power Manager. This function should never return.

Parameters

None

Returns

None

PWRM_vRegisterPreSleepCallback

```
void PWRM_vRegisterPreSleepCallback(  
    tsCallbackDescriptor *psCBDesc);
```

Description

This function is used to register a user-defined callback function that will be called by the Power Manager before the JN516x device enters sleep mode. You must specify a pointer to a structure containing a descriptor for your callback function.

The callback function must have been declared using the macro **PWRM_CALLBACK(*fn_name*)**, where *fn_name* is the name of the callback function.

The callback descriptor must have been declared using the macro **PWRM_DECLARE_CALLBACK_DESCRIPTOR(*desc_name*, *fn_name*)**, where *desc_name* is the descriptor name and *fn_name* is the callback function name.

For example:

```
PWRM_CALLBACK(vPreSleepCB1);  
PWRM_DECLARE_CALLBACK_DESCRIPTOR(pscbl_desc, vPreSleepCB1);
```

The callback function should perform any housekeeping tasks that are necessary before the device enters sleep mode.

Note that this registration function is normally called within the user-defined function **vAppRegisterPWRMCBcallbacks()**. This ensures that the callback is registered during a cold start.

Parameters

<i>*psCBDesc</i>	Pointer to callback descriptor structure
------------------	--

Returns

None

PWRM_vRegisterWakeupCallback

```
void PWRM_vRegisterWakeupCallback(  
    tsCallbackDescriptor *psCBDesc);
```

Description

This function is used to register a user-defined callback function that will be called by the Power Manager when the JN516x device wakes from sleep (this may be due to a change on a DIO line or comparator input, or the expiry of a wake timer). You must specify a pointer to a structure containing a descriptor for your callback function.

The callback function must have been declared using the macro **PWRM_CALLBACK(*fn_name*)**, where *fn_name* is the name of the callback function.

The callback descriptor must have been declared using the macro **PWRM_DECLARE_CALLBACK_DESCRIPTOR(*desc_name*, *fn_name*)**, where *desc_name* is the descriptor name and *fn_name* is the callback function name.

For example:

```
PWRM_CALLBACK (vWakeUpCB1) ;
```

```
PWRM_DECLARE_CALLBACK_DESCRIPTOR (wucb1_desc, vWakeUpCB1) ;
```

The callback function should perform any housekeeping tasks that are necessary after the device wakes from sleep.

Note that this registration function is normally called within the user-defined function **vAppRegisterPWRMCBcallbacks()**. This ensures that the callback is registered during a cold start.

Parameters

<i>*psCBDesc</i>	Pointer to callback descriptor structure
------------------	--

Returns

None

vAppRegisterPWRMCallbacks

```
void vAppRegisterPWRMCallbacks(void);
```

Description

This is a user-defined function to register pre- and post-sleep callback functions, if required.

The function definition must itself use **PWRM_vRegisterPreSleepCallback()** and **PWRM_vRegisterWakeupCallback()** to register the required callbacks.

Parameters

None

Returns

None

PWRM_vWakeInterruptCallback

```
void PWRM_vWakeInterruptCallback(void);
```

Description

This function is a pre-defined callback function which must be called from the application's interrupt handler to deal with interrupts from Wake Timer 1 on the JN516x device.

The function is needed to maintain the scheduled wake points list, by restarting the wake timer for the next wake-up event (if any) when the previous one has just completed. The function also calls the user-defined callback function specified through **PWRM_vScheduleActivity()**.

Parameters

None

Returns

None

11.3 Debugging Function

The PWRM debugging function can be used to investigate how long the JN516x device spends in Doze mode. The Doze state is output on the JN516x DIO1 pin for external monitoring, allowing you to calculate the proportion of time that the device typically spends in Doze mode for a given application.

The function is listed below, along with its page reference:

Function	Page
PWRM_vSetupDozeMonitor	149

PWRM_vSetupDozeMonitor

```
void PWRM_vSetupDozeMonitor(bool_t bUseIO);
```

Description

This function can be used during debug to start a Doze mode monitoring session on the JN516x device - that is, to investigate the proportion of the time that the device typically spends in Doze mode.

The Doze state of the device is output on the pin DIO1. This allows the times spent in and out of Doze mode to be measured externally.

Parameters

<i>bUseIO</i>	Always set to TRUE
---------------	--------------------

Returns

None

12. PDUM API

This chapter describes the functions of the JenOS Protocol Data Unit Manager (PDUM) API. The API is defined in the header file **pdum.h**.

The PDUM API functions are listed below, along with their page references:

Function	Page
PDUM_vInit	152
PDUM_hAPduAllocateAPduInstance	153
PDUM_eAPduFreeAPduInstance	154
PDUM_u16APduInstanceReadNBO	155
PDUM_u16APduInstanceWriteNBO	156
PDUM_u16APduInstanceWriteStrNBO	157
PDUM_u16SizeNBO	158
PDUM_u16APduGetSize	159
PDUM_pvAPduInstanceGetPayload	160
PDUM_u16APduInstanceGetPayloadSize	161
PDUM_eAPduInstanceSetPayloadSize	162
PDUM_vDBGPrintAPduInstance	163



Note: In ZigBee PRO, the APDUs used by the application must be pre-defined (before building the application) using the ZPS Configuration Editor. This tool is detailed in the *ZigBee PRO Stack User Guide* (JN-UG-3048 or JN-UG-3101).

PDUM_vInit

```
void PDUM_vInit();
```

Description

This function initialises the PDU Manager and must therefore be the first PDUM function called.

Parameters

None

Returns

None

PDUM_hAPduAllocateAPduInstance

```
PDUM_thAPduInstance  
PDUM_hAPduAllocateAPduInstance(  
    PDUM_thAPdu hAPdu);
```

Description

This function allocates an instance of an Application Protocol Data Unit (APDU) - that is, memory space is allocated to the APDU instance.

The available APDUs (types and their handles) are pre-defined using the ZPS Configuration Editor (refer to the *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)*).

The allocated APDU instance can subsequently be populated with data and sent to another node.

Parameters

hAPdu Handle of APDU (type)

Returns

Handle of allocated APDU instance

PDUM_INVALID_HANDLE if no APDU instances free

PDUM_eAPduFreeAPduInstance

```
PDUM_tStatus PDUM_eAPduFreeAPduInstance(  
    PDUM_thAPduInstance hAPduInst);
```

Description

This function de-allocates the specified APDU instance, thus freeing the associated memory space.

Parameters

hAPduInstance Handle of APDU instance

Returns

PDUM_E_INTERNAL_ERROR

PDUM_u16APdulInstanceReadNBO

```
uint16 PDUM_u16APdulInstanceReadNBO(  
    PDUM_thAPdulInstance hAPdulInst,  
    uint16 u16Pos,  
    const char *szFormat,  
    void *pvStruct);
```

Description

This function reads data from the specified APDU instance and inserts the data into a C structure. The byte position of the start (least significant byte) of the data in the APDU instance must be specified, as well as the format of the data.

Data is read from the APDU instance in packed network byte order (little-endian) and translated into unpacked host byte order for the C structure (big-endian for the JN516x device).

Parameters

<i>hAPdulInst</i>	Handle of APDU instance to read the data from
<i>u32Pos</i>	The starting position (least significant byte) of the data within the APDU
<i>*szFormat</i>	Format string of the data: b 8-bit byte h 16-bit half-word (short integer) w 32-bit word l 64-bit long-word (long integer) a\xnn nn (hex) bytes of data (array) p\xnn nn (hex) bytes of packing
<i>*pvStruct</i>	Pointer to C structure to receive the data

Note that the compiler will not correctly interpret the format string “a\xnnb” for a data array followed by a single byte, e.g. “a\x0ab”. In this case, to ensure that the ‘b’ (for byte) is not interpreted as a hex value, use the format “a\xnn” “b”, e.g. “a\x0a” “b”.

Returns

Total number of data bytes read from the APDU instance

PDUM_u16APdulInstanceWriteNBO

```
uint16 PDUM_u16APdulInstanceWriteNBO(  
    PDUM_thAPdulInstance hAPdulInst,  
    uint16 u16Pos,  
    const char *szFormat, ...);
```

Description

This function writes the specified data values into the specified APDU instance. The byte position of the start of the data (least significant byte) in the APDU instance must be specified, as well as the format of the data.

The data values are written into the APDU instance at the specified position in packed network byte order (little-endian). The input data values should be in host byte order (big-endian for the JN516x device).

Parameters

<i>hAPdulInst</i>	Handle of the APDU instance to write the data into
<i>u32Pos</i>	The starting position (least significant byte) of the data within the APDU instance
<i>*szFormat</i>	Format string of the data: b 8-bit byte h 16-bit half-word (short integer) w 32-bit word l 64-bit long-word (long integer) a\xnn nn (hex) bytes of data (array) p\xnnnn (hex) bytes of packing
...	Variable list of data values described by the format string

Note that the compiler will not correctly interpret the format string “a\xnnb” for a data array followed by a single byte, e.g. “a\x0ab”. In this case, to ensure that the ‘b’ (for byte) is not interpreted as a hex value, use the format “a\xnn” “b”, e.g. “a\x0a” “b”.

Returns

Total number of bytes written to the APDU instance

PDUM_u16APdulInstanceWriteStrNBO

```
uint16 PDUM_u16APdulInstanceWriteStrNBO(  
    PDUM_thAPdulInstance hAPdulInst,  
    uint16 u16Pos,  
    const char *szFormat,  
    void *pvStruct);
```

Description

This function writes data from the specified structure into the specified APDU instance. The byte position of the start of the data (least significant byte) in the APDU instance must be specified, as well as the format of the data.

The data values are written into the APDU instance at the specified position in packed network byte order (little-endian). The input data values should be in host byte order (big-endian for the JN516x device).

Parameters

<i>hAPdulInst</i>	Handle of the APDU instance to write the data into
<i>u32Pos</i>	The starting position (least significant byte) of the data within the APDU instance
<i>*szFormat</i>	Format string of the data: b 8-bit byte h 16-bit half-word (short integer) w 32-bit word l 64-bit long-word (long integer) a\xnn nn (hex) bytes of data (array) p\xnnnn (hex) bytes of packing
<i>*pvStruct</i>	Pointer to C structure to containing data

Note that the compiler will not correctly interpret the format string “a\xnnb” for a data array followed by a single byte, e.g. “a\x0ab”. In this case, to ensure that the ‘b’ (for byte) is not interpreted as a hex value, use the format “a\xnn” “b”, e.g. “a\x0a” “b”.

Returns

Total number of bytes written to the APDU instance

PDUM_u16SizeNBO

```
uint16 PDUM_u16SizeNBO(const char *szFormat);
```

Description

This function obtains the size, in bytes, of an APDU data payload, given the format of the data.

Parameters

<i>*szFormat</i>	Format string of the data: <ul style="list-style-type: none">b 8-bit byteh 16-bit half-word (short integer)w 32-bit wordl 64-bit long-word (long integer)a\xnn nn (hex) bytes of data (array)p\xnnnn (hex) bytes of packing
------------------	--

Note that the compiler will not correctly interpret the format string “a\xnnb” for a data array followed by a single byte, e.g. “a\x0ab”. In this case, to ensure that the ‘b’ (for byte) is not interpreted as a hex value, use the format “a\xnn” “b”, e.g. “a\x0a” “b”.

Returns

Number of bytes in data payload

PDUM_u16APduGetSize

```
uint16 PDUM_u16APduGetSize(PDUM_thAPdu hAPdu);
```

Description

This function obtains the maximum size, in bytes, of the specified APDU (type).

Parameters

<i>hAPdu</i>	Handle of APDU
--------------	----------------

Returns

Number of bytes in APDU

PDUM_pvAPdulInstanceGetPayload

```
void * PDUM_pvAPdulInstanceGetPayload(  
    PDUM_thAPdulInstance hAPdulInst);
```

Description

This function obtains a pointer to the payload data of the specified APDU instance.

Parameters

<i>hAPdulInst</i>	Handle of APDU instance to access
-------------------	-----------------------------------

Returns

Pointer to data as an array of bytes

PDUM_u16APdulInstanceGetPayloadSize

```
uint16 PDUM_u16APdulInstanceGetPayloadSize(  
    PDUM_thAPdulInstance hAPdulInst);
```

Description

This function obtains the size, in bytes, of the payload data of the specified APDU instance.

Parameters

hAPdulInst Handle of APDU instance to access

Returns

Size of the payload data, in bytes

PDUM_eAPdulInstanceSetPayloadSize

```
PDUM_teStatus PDUM_eAPdulInstanceSetPayloadSize(  
    PDUM_thAPdulInstance hAPdulInst,  
    uint16 u16Size);
```

Description

This function sets the size, in bytes, of the payload of the specified APDU instance.

Parameters

<i>hAPdulInst</i>	Handle of APDU instance
<i>u16Size</i>	Size of payload to set, in bytes

Returns

PDUM_OK
PDUM_E_APDU_INSTANCE_TOO_BIG

PDUM_vDBGPrintAPdulInstance

```
void PDUM_vDBGPrintAPdulInstance(  
    PDUM_thAPdulInstance hAPdulInst);
```

Description

This function can be used to output the specified APDU instance via the Debug (DBG) module.

For details of the DBG functions, refer to [Chapter 12](#).

Parameters

<i>hAPdu</i>	Handle of APDU instance to output
--------------	-----------------------------------

Returns

None

13. DBG API

The chapter describes the functions of the JenOS Debug (DBG) module API. The API is defined in the header file **dbg.h**.

To use the Debug module, it must be enabled at build-time by defining **DBG_ENABLE** in the build - for example, by adding the **-DDBG_ENABLE** option to the compiler.

By default, the Debug module will just display each line as passed. However, if **DBG_VERBOSE** is defined at build-time then each line displayed will be prefixed with the file name and line number of the debug statement.



Note: Compiling with the DBG option results in a larger application size, requiring a lot more space in RAM.

The DBG API functions are listed below, along with their page references:

Function	Page
DBG_vInit	166
DBG_vUartInit	167
DBG_vPrintf	168
DBG_vAssert	169
DBG_vDumpStack	170
DBG_vFlush	171
DBG_iGetChar	172

DBG_vlnit

```
void DBG_vlnit(tsDBG_FunctionTbl *psFunctionTbl);
```

Description

This function is used to initialise the Debug module.



Note: If a JN516x UART is to be used as the debug output interface, **DBG_vUartInit()** must be called instead. Thus, **DBG_vlnit()** will not be needed by most users, since a UART will normally be used for debug output.

The function can be used during a cold start or a warm start (with memory held). Its parameter accepts a structure containing pointers to four user-defined callback functions concerned with the output interface:

```
typedef struct
{
    void (*prInitHardwareCb)(void);
    void (*prPutchCb)      (char c);
    void (*prFlushCb)      (void);
    void (*prFailedAssertCb)(void);
} tsDBG_FunctionTbl;
```

The callback functions pointed to by this structure are as follows:

- | | |
|--------------------------|---|
| *prInitHardwareCb | Points to function which re-initialises the interface after a warm start, e.g. when JN516x device wakes from sleep |
| *prPutchCb | Points to function used by DBG_vPrintf() to output a single character to the interface |
| *prFlushCb | Points to function used by DBG_vPrintf() to flush the interface buffer to allow buffered output characters to be displayed. If the output is unbuffered, this function should do nothing or wait for the last character output using the putch() function to be made available. Note that the function should not append a newline character, as this should be handled by the formatting string passed to DBG_vPrintf() |
| *prAssertFailedCb | Points to function which is called when DBG_vAssert() fails. The function should stop execution and may reset the device |

Parameters

- *psFunctionTbl** Pointer to structure containing list of callback functions.

Returns

None

DBG_vUartInit

```
void DBG_vUartInit(DBG_teUart eUart,  
                  DBG_teUartBaudRate eBaudRate);
```

Description

This function is used to initialise the Debug module when one of the JN516x on-chip UARTs is to be used as the output interface. In this case, this function should be called instead of **DBG_vInit()**. This will be the case for most users, as a UART will normally be used for debug output.

The function can be used during a cold start or a warm start (with memory held). It is necessary to specify the UART (0 or 1) and the required baud rate.

Note that the callback functions required by **DBG_vInit()** are not needed for **DBG_vUartInit()**, since they are pre-defined by NXP for the on-chip UARTs.

Parameters

<i>eUart</i>	UART to use as output interface, one of: DBG_E_UART_0 (UART0) DBG_E_UART_1 (UART1)
<i>eBaudRate</i>	Baud rate of UART, one of: DBG_E_UART_BAUD_RATE_4800 (4800 bps) DBG_E_UART_BAUD_RATE_9600 (9600 bps) DBG_E_UART_BAUD_RATE_19200 (19200 bps) DBG_E_UART_BAUD_RATE_38400 (38400 bps) DBG_E_UART_BAUD_RATE_76800 (76800 bps) DBG_E_UART_BAUD_RATE_115200 (115200 bps)

Returns

None

DBG_vPrintf

```
void DBG_vPrintf(bool_t bStreamEnabled,  
                const char *pcFormat, ...);
```

Description

This function is an adapted **printf()** function, allowing a formatted string to be output (e.g. via the UART) for display.

The function contains a parameter which allows the output of the string to be enabled or disabled - the value of this Boolean parameter must be a literal. If disabled, the compiler will optimise out this function, but its parameters will still be evaluated.

Parameters

<i>bStreamEnabled</i>	Boolean which determines whether string will be output: TRUE: Output string FALSE: Do not output string (compile out function)
<i>*pcFormat</i>	Pointer to printf-style formatting string
...	As for the standard printf() function

Returns

None

DBG_vAssert

```
void DBG_vAssert(bool_t bStreamEnabled,  
                bool_t bAssertion);
```

Description

This function is an adapted **assert()** function, allowing a Boolean condition to be tested.

The function contains a parameter which allows the test to be enabled or disabled - the value of this Boolean parameter must be a literal. If disabled, the compiler will optimise out this function.

The Boolean condition to be tested is specified as a parameter:

- If the condition is TRUE, program execution continues.
- If the condition is FALSE, an error message is output and execution is passed to a callback function, which stops execution. This callback function is specified when **DGB_vlnit()** is called for a cold start.

Parameters

<i>bStreamEnabled</i>	Boolean which determines whether test will be performed: TRUE: Perform test FALSE: Do not perform test
<i>bAssertion</i>	Boolean expression to be tested

Returns

None

DBG_vDumpStack

```
void DBG_vDumpStack(void);
```

Description

This function outputs the contents of the CPU stack (e.g. via the UART) for display.

Parameters

None

Returns

None

DBG_vFlush

```
void DBG_vFlush(void);
```

Description

This function flushes buffered characters from the JN516x device to the display device. If the JN516x UART is used for debug, this function flushes the UART buffer.

Parameters

None

Returns

None

DBG_iGetChar

```
int DBG_iGetChar(void);
```

Description

This function can be used to obtain a character from an input device (such as a serial terminal connected to the JN516x UART).

Parameters

None

Returns

ASCII value of obtained character, or -1 if no character available

14. JenOS Structures

This chapter describes the structures used by the JenOS APIs.

The structures are listed below along with their page references.

Structure	Page
PDM_tsHwFncTable	173
tSPIflashFncTable	174
PWRM_teSleepMode	176
DBG_tsFunctionTbl	176
tsReg128	177
PDM_tpfvSystemEventCallback	177
PDM_eSystemEventCode	177
OS_teStatus	182

14.1 PDM_tsHwFncTable

This structure is used in the function **PDM_vlnit()** to specify a set of user-defined functions used to interact with a custom NVM device.

```
typedef struct
{
    /* This function is called after a cold or warm start */
    void (*prInitHwCb)(void);

    /* This function is called to erase the given sector */
    void (*prEraseCb) (uint8 u8Sector);

    /*This function is called to write data to an address
    * within a given sector. Address zero is the start of the
    * given sector */
    void (*prWriteCb) (uint8 u8Sector,
                      uint16 ul6Addr,
                      uint16 ul6Len,
                      uint8 *pu8Data);

    /* This function is called to read data from an address
    * within a given sector. Address zero is the start of the
    * given sector */
    void (*prReadCb) (uint8 u8Sector,
                     uint16 ul6Addr,
                     uint16 ul6Len,
                     uint8 *pu8Data);
} PDM_tsHwFncTable;
```

14.2 tSPIflashFncTable

This structure is used in the function **PDM_vSPIFlashConfig()** to specify a set of user-defined functions used to interact with a custom SPI Flash device.

```
typedef struct tagSPIflashFncTable {
    uint32          u32Signature; //always set to 0x1234678
    uint16          u16FlashId;  //(u8ManufactureId<<8)|u8DeviceId
    uint16          u16Reserved; //Reserved
    tpfvZSPIflashInit  vZSPIflashInit; //see below
    tpfvZSPIflashSetSlaveSel vZSPIflashSetSlaveSel; //see below
    tpfvZSPIflashWREN  vZSPIflashWREN; //see below
    tpfvZSPIflashEWRSR vZSPIflashEWRSR; //see below
    tpfu8ZSPIflashRDSR u8ZSPIflashRDSR; //see below
    tpfu16ZSPIflashRDID u16ZSPIflashRDID; //see below
    tpfvZSPIflashWRSR  vZSPIflashWRSR; //see below
    tpfvZSPIflashPP    vZSPIflashPP; //see below
    tpfvZSPIflashRead  vZSPIflashRead; //see below
    tpfvZSPIflashBE    vZSPIflashBE; //see below
    tpfvZSPIflashSE    vZSPIflashSE; //see below
} tSPIflashFncTable;
```

The custom functions specified in this structure are outlined in [Table 1](#) below.

Function Prototype	Description
void vZSPIflashInit(int <i>iDivisor</i>, uint8 <i>u8SlaveSel</i>);	Initialises variables for Flash access. <ul style="list-style-type: none"> • <i>iDivisor</i> Clock divisor • <i>u8SlaveSel</i> Byte used for slave select
void vZSPIflashSetSlaveSel(uint8 <i>u8SlaveSel</i>);	
void vZSPIflashWREN(void);	Enables writes to Flash. Called before erasing or programming Flash.
vZSPIflashEWSR(void);	Enables writes to Flash Status Register. Called before writing to the Flash Status Register.
uint8 u8ZSPIflashRDSR(void);	Reads Flash Status Register and returns Status Register data
uint16 u16ZSPIflashRDID(void);	Reads Flash ID Register and returns ID Register data, 0 on error or 2 bytes [ManufacturerId, DeviceId]
void vZSPIflashWRSR(uint8 <i>u8Data</i>);	Writes data to Flash Status Register <ul style="list-style-type: none"> • <i>u8Data</i> Status Register data
void vZSPIflashPP(uint32 <i>u32Addr</i>, uint16 <i>u16Len</i>, uint8* <i>pu8Data</i>);	Writes data to Flash. <ul style="list-style-type: none"> • <i>u32Addr</i> Address • <i>u16Len</i> Length, in bytes • <i>pu8Data</i> Data to write
void vZSPIflashRead(uint32 <i>u32Addr</i>, uint16 <i>u16Len</i>, uint8* <i>pu8Data</i>);	Reads data from Flash. <ul style="list-style-type: none"> • <i>u32Addr</i> Address • <i>u16Len</i> Length, in bytes • <i>pu8Data</i> Data read
void vZSPIflashBE(void);	Performs a bulk erase of Flash
void vZSPIflashSE(uint8 <i>u8Sector</i>);	Performs a sector erase of Flash <ul style="list-style-type: none"> • <i>u8Sector</i> Sector number

Table 1: Functions of tSPIflashFncTable Structure

14.3 PWRM_teSleepMode

This structure contains the enumerations used to set the power mode of the JN516x device during sleep.

```
typedef enum
{
    PWRM_E_SLEEP_OSCON_RAMON,    /*32-kHz Osc on and RAM on*/
    PWRM_E_SLEEP_OSCON_RAMOFF,   /*32-kHz Osc on and RAM off*/
    PWRM_E_SLEEP_OSCOFF_RAMON,   /*32-kHz Osc off and RAM on*/
    PWRM_E_SLEEP_OSCOFF_RAMOFF,  /*32-kHz Osc off and RAM off*/
    PWRM_E_SLEEP_DEEP,           /*Deep Sleep*/
} PWRM_teSleepMode;
```

14.4 DBG_tsFunctionTbl

This structure contains callback functions used by the Debug (DBG) module to interact with the output interface.

```
typedef struct
{
    void (*prInitHardwareCb)(void);
    void (*prPutchCb)(char c);
    void (*prFlushCb)(void);
    void (*prFailedAssertCb)(void);
} DBG_tsFunctionTbl;
```

For details of the callback functions, refer to the description of [DBG_vInit](#) on page 166.

14.5 tsReg128

This is a constant structure which contains a 128-bit encryption key used by the PDM module - the key is passed into the module via the **PDM_vInit()** function.

```
typedef struct
{
    uint32 u32register0;
    uint32 u32register1;
    uint32 u32register2;
    uint32 u32register3;
} tsReg128;
```

In the above structure, `u32register0` contains the 32 least significant bits and `u32register3` contains the 32 most significant bits of the key.

14.6 PDM_tpfvSystemEventCallback

This type defines the callback function that receives PDM events.

```
typedef void (*PDM_tpfvSystemEventCallback) (
    uint32                                u32eventNumber,
    PDM_eSystemEventCode                  eSystemEventCode);
```

where:

- `u32eventNumber` gives further information about the event depending on the event code, as detailed in [Section 14.7](#)
- `eSystemEventCode` identifies the type of event that triggered the callback.

14.7 PDM_eSystemEventCode

This structure contains enumerations for the events generated by the PDM library. There are two versions of this structure:

- One for the PDM for external Flash memory, used by ZigBee Smart Energy (and described in [Chapter 3](#) and [Chapter 12](#))
- One for the PDM for internal EEPROM, used by ZigBee Home Automation, ZigBee Light Link and JenNet-IP (and described in [Chapter 4](#) and [Chapter 10](#))

PDM for External Flash Memory

```
typedef enum
{
    E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED=0,
    E_PDM_SYSTEM_EVENT_DESCRIPTOR_SAVE_FAILED,
    E_PDM_SYSTEM_EVENT_PDM_NOT_ENOUGH_SPACE,
    E_PDM_SYSTEM_EVENT_EEPROM_SEGMENT_HEADER_REPAIRED,
    E_PDM_SYSTEM_EVENT_SYSTEM_INTERNAL_BUFFER_WEAR_COUNT_SWAP,
    E_PDM_SYSTEM_EVENT_SYSTEM_DUPLICATE_FILE_SEGMENT_DETECTED
} PDM_eSystemEventCode;
```

The events are outlined in [Table 2](#) below.

Event Enumeration	Description
E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.
E_PDM_SYSTEM_EVENT_DESCRIPTOR_SAVE_FAILED	A save has failed. <code>u32eventNumber</code> contains the <code>u16IdValue</code> of the record that failed to save. This is a fatal error as the ZigBee PRO stack records may be inconsistent. Test software should log this error and halt. Production software may need to perform a factory reset.
E_PDM_SYSTEM_EVENT_PDM_NOT_ENOUGH_SPACE	There is not enough space to hold all the PDM records. <code>u32eventNumber</code> contains the <code>u16IdValue</code> of the record that was being processed. This is a fatal error as the ZigBee PRO stack records may be inconsistent. Test software should log this error and halt. Production software may need to perform a factory reset.
E_PDM_SYSTEM_EVENT_EEPROM_SEGMENT_HEADER_REPAIRED	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.
E_PDM_SYSTEM_EVENT_SYSTEM_INTERNAL_BUFFER_WEAR_COUNT_SWAP	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.
E_PDM_SYSTEM_EVENT_SYSTEM_DUPLICATE_FILE_SEGMENT_DETECTED	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.

Table 2: PDM Event Codes (Flash Memory)

PDM for Internal EEPROM

```
typedef enum
{
    E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED=0,
    E_PDM_SYSTEM_EVENT_SAVE_FAILED,
    E_PDM_SYSTEM_EVENT_PDM_NOT_ENOUGH_SPACE,
    E_PDM_SYSTEM_EVENT_LARGEST_RECORD_FULL_SAVE_NO_LONGER_POSSIBLE,
    E_PDM_SYSTEM_EVENT_SEGMENT_DATA_CHECKSUM_FAIL,
    // Debug event codes
    E_PDM_SYSTEM_EVENT_EEPROM_SEGMENT_HEADER_REPAIRED,
    E_PDM_SYSTEM_EVENT_SYSTEM_INTERNAL_BUFFER_WEAR_COUNT_SWAP,
    E_PDM_SYSTEM_EVENT_SYSTEM_DUPLICATE_FILE_SEGMENT_DETECTED,
    E_PDM_SYSTEM_EVENT_SYSTEM_ERROR,
} PDM_eSystemEventCode;
```

The events are outlined in [Table 3](#) below.

Event Enumeration	Description
E_PDM_SYSTEM_EVENT_WEAR_COUNT_TRIGGER_VALUE_REACHED	An EEPROM segment has reached a set Wear Count (set by the user or left at the manufacturer stated maximum value). <code>u32EventNumber</code> carries the EEPROM segment number.
E_PDM_SYSTEM_EVENT_SAVE_FAILED	A save has failed. <code>u32eventNumber</code> contains the <code>u16IdValue</code> of the record that failed to save. This is a fatal error as the stack records may be inconsistent. Test software should log this error and halt. Production software may need to perform a factory reset.
E_PDM_SYSTEM_EVENT_PDM_NOT_ENOUGH_SPACE	There is not enough space to hold all the PDM records. <code>u32eventNumber</code> contains the <code>u16IdValue</code> of the record that was being processed. This is a fatal error as the stack records may be inconsistent. Test software should log this error and halt. Production software may need to perform a factory reset.
E_PDM_SYSTEM_EVENT_LARGEST_RECORD_FULL_SAVE_NO_LONGER_POSSIBLE	The EEPROM occupancy is such that the largest record in the PDM can no longer be fully saved. <code>u32EventNumber</code> carries the <code>u16IdValue</code> of the record that was being processed.
E_PDM_SYSTEM_EVENT_SEGMENT_DATA_CHECKSUM_FAIL	The calculated checksum for the data in an EEPROM segment does not match the stored checksum value. <code>u32EventNumber</code> carries the number of the segment.
E_PDM_SYSTEM_EVENT_EEPROM_SEGMENT_HEADER_REPAIRED	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.

Table 3: PDM Event Codes (EEPROM)

Event Enumeration	Description
E_PDM_SYSTEM_EVENT_SYSTEM_INTERNAL_BUFFER_WEAR_COUNT_SWAP	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.
E_PDM_SYSTEM_EVENT_SYSTEM_DUPLICATE_FILE_SEGMENT_DETECTED	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.
E_PDM_SYSTEM_EVENT_SYSTEM_ERROR	This code can be ignored by the application software and only needs to be logged if requested by NXP Technical Support.

Table 3: PDM Event Codes (EEPROM)

14.8 PDM_teStatus

This structure contains enumerations for the status codes generated by the PDM module.

```
typedef enum
{
    PDM_E_STATUS_OK,
    PDM_E_STATUS_INVLD_PARAM,
    // EEPROM based PDM codes
    PDM_E_STATUS_PDM_FULL,
    PDM_E_STATUS_NOT_SAVED,
    PDM_E_STATUS_RECOVERED,
    PDM_E_STATUS_PDM_RECOVERED_NOT_SAVED,
    PDM_E_STATUS_USER_BUFFER_SIZE,
    PDM_E_STATUS_BITMAP_SATURATED_NO_INCREMENT,
    PDM_E_STATUS_BITMAP_SATURATED_OK,
    PDM_E_STATUS_IMAGE_BITMAP_COMPLETE,
    PDM_E_STATUS_IMAGE_BITMAP_INCOMPLETE,
    PDM_E_STATUS_INTERNAL_ERROR
} PDM_teStatus;
```

The status codes are described in [Table 4](#) below.

Event Enumeration	Description
PDM_E_STATUS_OK	The function completed without error.
PDM_E_STATUS_INVLD_PARAM	An invalid parameter value was supplied.
PDM_E_STATUS_PDM_FULL	There is no available EEPROM space for PDM.
PDM_E_STATUS_NOT_SAVED	A PDM save to EEPROM failed.
PDM_E_STATUS_RECOVERED	The record was recovered from a previous save to NVM.
PDM_E_STATUS_PDM_RECOVERED_NOT_SAVED	The record was not recovered from a previous save to NVM.
PDM_E_STATUS_USER_BUFFER_SIZE	Not used.
PDM_E_STATUS_BITMAP_SATURATED_NO_INCREMENT	Counter increment not made because the EEPROM segment is saturated.
PDM_E_STATUS_BITMAP_SATURATED_OK	Counter increment made but the EEPROM segment is now saturated.
PDM_E_STATUS_IMAGE_BITMAP_COMPLETE	For internal use.
PDM_E_STATUS_IMAGE_BITMAP_INCOMPLETE	For internal use.
PDM_E_STATUS_INTERNAL_ERROR	An unspecified internal PDM error has occurred.

Table 4: PDM Status Codes

14.9 OS_teStatus

This structure contains enumerations for the JenOS status codes generated by the core OS library.

```
typedef enum {  
    OS_E_OK = 0,  
    OS_E_BADTASK = 1,  
    OS_E_BADMUTEX = 2,  
    OS_E_BADMESSAGE = 3,  
    OS_E_BADVALUE = 4,  
    OS_E_OVERACTIVATION = 5,  
    OS_E_QUEUE_EMPTY = 6,  
    OS_E_QUEUE_FULL = 7,  
    OS_E_UNQUEUED = 8,  
    OS_E_OSINTOVERFLOW = 9,  
    OS_E_OSINTUNDERFLOW = 10,  
    OS_E_BADSWTIMER = 11,  
    OS_E_BADHWCOUNTER = 12,  
    OS_E_SWTIMER_STOPPED = 13,  
    OS_E_SWTIMER_EXPIRED = 14,  
    OS_E_SWTIMER_RUNNING = 15,  
    OS_E_HWCOUNTERIDLE = 16,  
    OS_E_NOTHINGTOEXPIRE = 17,  
    OS_E_PRIORITY_ERROR = 18,  
    OS_E_BAD_NESTING = 19,  
    OS_E_TICKS_TOO_BIG = 20,  
    OS_E_CURRENT_TASK_NOT_A_MUTEX_MEMBER = 21,  
    OS_E_TASK_NOT_A_MESSAGE_POSTER = 22,  
    OS_E_TASK_NOT_A_MESSAGE_COLLECTOR = 23  
} OS_teStatus;
```

The status codes are described in [Table 5](#) below. Fatal errors must be handled using the mechanism described in [Section 2.5](#).

Event Enumeration	Description
OS_E_OK	The function completed without error.
OS_E_BADTASK	A bad task handle has been passed to a function. This is a fatal error.
OS_E_BADMUTEX	A bad mutex handle has been passed to a function. This is a fatal error.
OS_E_BADMESSAGE	A bad message handle has been passed to a function. This is a fatal error.
OS_E_BADVALUE	An out-of-range value has been passed to a function. This is a fatal error.
OS_E_OVERACTIVATION	A task has been activated too many times. This is a fatal error.
OS_E_QUEUE_EMPTY	An attempt has been made to read from an empty queue.
OS_E_QUEUE_FULL	An attempt has been made to post to a queue that is full. Whilst the OS can recover from this situation, this error should normally be treated as fatal. If a ZigBee PRO stack queue overflows, the stack can be left in an inconsistent state.
OS_E_UNQUEUED	Returned by OS_eGetMessageStatus when the queue is not empty or full.
OS_E_OSINTOVERFLOW	There have been too many nested interrupts. This is a fatal error.
OS_E_OSINTUNDERFLOW	A resume from interrupts has failed due to there being no matching interrupt. This is a fatal error.
OS_E_BADSWTIMER	A bad timer handle has been passed to a function. This is a fatal error.
OS_E_BADHWCOUNTER	A bad hardware counter handle has been passed to a function. This is a fatal error.
OS_E_SWTIMER_STOPPED	An attempt has been made to stop a timer that is already stopped. This is not necessarily a fatal error.
OS_E_SWTIMER_EXPIRED	The software timer has expired. This is not a fatal error.
OS_E_SWTIMER_RUNNING	The software timer is running. This is not normally a fatal error. However, when calling OS_eContinueSWTimer() on a timer that is already running, the expiry time of the timer will remain at the time previously set and will not be changed by this call.
OS_E_HWCOUNTERIDLE	A code used internally that is not presented to application software.
OS_E_NOTHINGTOEXPIRE	The hardware timer has expired but no software timers are due to expire. This is a fatal error.

Table 5: OS Status Codes

Event Enumeration	Description
OS_E_PRIORITY_ERROR	The priority levels internal to the OS are not consistent. This is a fatal error.
OS_E_BAD_NESTING	There have been two calls to enter a critical section without a leave call. This is a fatal error.
OS_E_TICKS_TOO_BIG	An attempt has been made to start a timer too far into the future. This is a fatal error.
OS_E_CURRENT_TASK_NOT_A_MUTEX_MEMBER	A task has attempted to take a mutex when the task is not in the mutex group in the OS configuration diagram. This is a fatal error.
OS_E_TASK_NOT_A_MESSAGE_POSTER	A task has attempted to post a message to a queue when the task is not connected to the queue in the OS configuration diagram. This is a fatal error.
OS_E_TASK_NOT_A_MESSAGE_COLLECTOR	A task has attempted to collect a message from a queue when the task is not connected to the queue in the OS configuration diagram. This is a fatal error.

Table 5: OS Status Codes

Part III:

Configuration Information

15. JenOS Configuration

The use of certain JenOS resources must be statically configured before building an application, particularly resources relating to the RTOS and PDU Manager, such as timers, mutexes and ISRs. This chapter introduces the JenOS Configuration Editor, the graphical tool used to perform this configuration.

The JenOS Configuration Editor is an NXP-devised plug-in for the Eclipse IDE and is therefore also compatible with 'BeyondStudio for NXP'. In developing a ZigBee PRO application, ZigBee network parameters must be pre-configured using the ZPS Configuration Editor, which is also an NXP-devised plug-in for Eclipse/BeyondStudio.

The Eclipse or BeyondStudio platform is provided as part of the JN516x toolchain, JN-SW-4041 or JN-SW-4141 respectively. Both of the above plug-ins are provided in the JN516x SDKs for the ZigBee application profiles (SE, HA, ZLL).

The principles of the build-time configuration for a ZigBee PRO application are described in [Section 15.2](#). The ZPS Configuration Editor is described in the *ZigBee PRO Stack User Guide (JN-UG-3048 or JN-UG-3101)*.

15.1 CPU Stack and Heap Sizes

For ZigBee PRO, the CPU stack and heap sizes are respectively set to 5000 bytes and 2000 bytes, by default. If you wish to change these sizes, you can over-ride the default values in your application makefile using `__stack_size` and `__minimum_heap_size`.

For example:

```
__stack_size = 6000;  
__minimum_heap_size = 2500;
```

You should increase the stack size when including the ZigBee OTA Upgrade feature.

15.2 Configuration Principles

The build process for a ZigBee PRO application takes a number of configuration files, in addition to the application source file and header file. The following files are generated from Eclipse to feed into the build process:

- ZigBee PRO Stack files:
 - **zps_gen.c**
 - **zps_gen.h**
- PDU Manager files:
 - **pdum_gen.c**
 - **pdum_gen.h**

- RTOS files:
 - **os_gen.c**
 - **os_gen.h**
 - **os_irq.s**

All of the above files are produced according to the same basic principles. The NXP plug-ins in Eclipse are used to edit the configuration data and output this data as XML files (the XML files can be coded manually, outside of Eclipse, but this is not recommended). As part of the build process, the application's makefile invokes command line utilities that use the XML files to generate the files listed above.

The full build process is illustrated in [Figure 1](#).

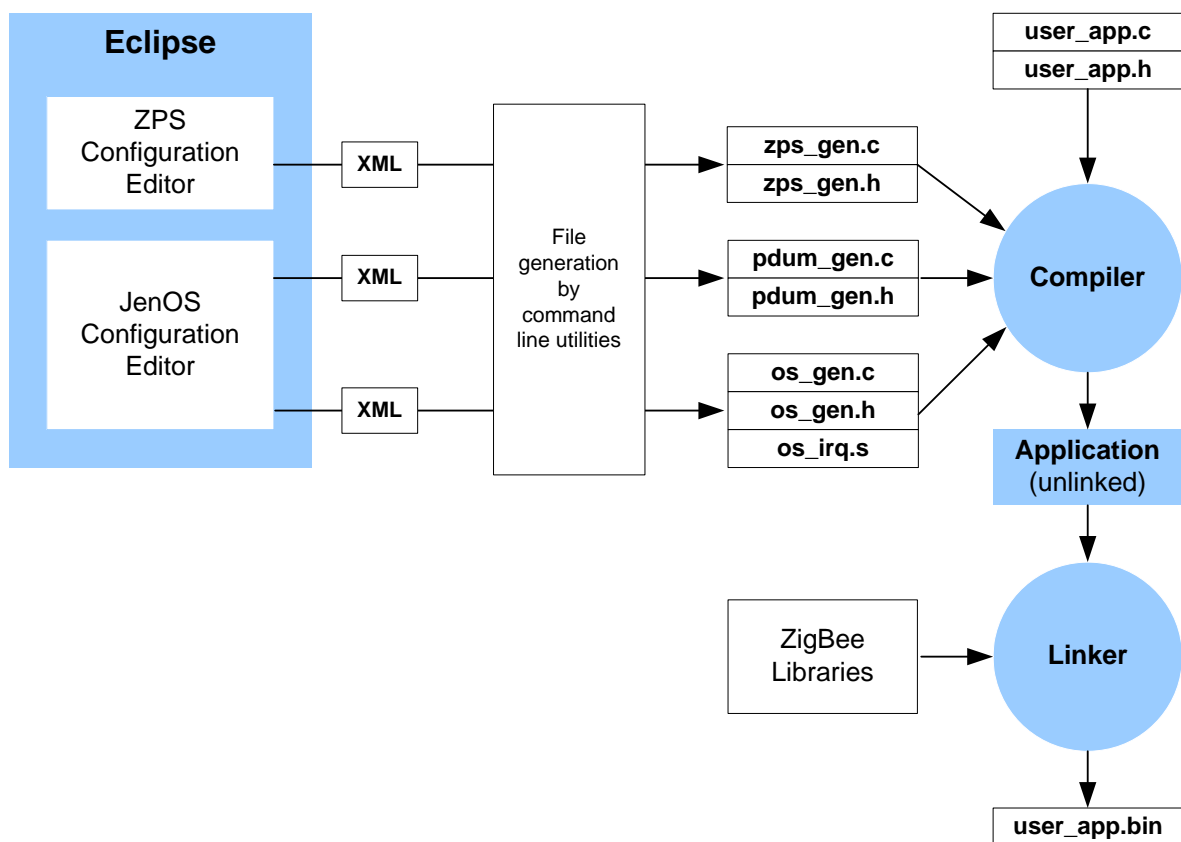


Figure 1: Application Build Process

15.3 Configuring JenOS Resources

The JenOS configuration is closely linked to application coding and can be performed before, during or after coding. This configuration is largely concerned with the management of task/ISR execution and scheduling. Typical settings include:

- The execution priorities of individual user tasks and ISRs
- A user task's membership of a particular mutex group
- The interrupt which 'stimulates' execution of a particular ISR
- The software timers derived from a particular hardware counter
- The callback functions associated with a particular hardware counter
- The user task that is activated when a particular software timer expires
- The message types that can be sent and received by a particular user task
- The length of a user task's message queue for a particular message type

This information is represented in the editor in graphical form. The editor's graphical window is divided into separate regions for the application, ZigBee PRO stack and CPU exceptions. An example of this window is shown below.

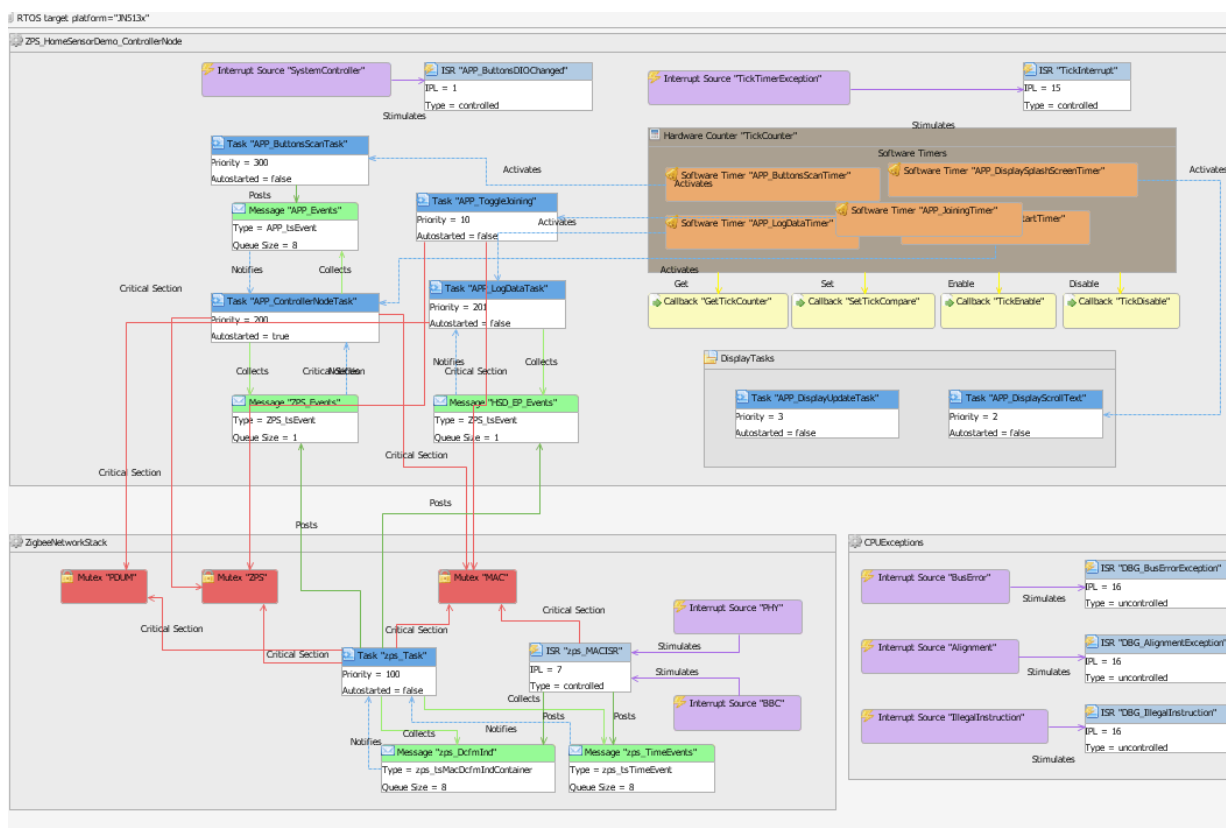



Figure 2: Graphical Configuration Window

Generally, the stack region of the diagram is the same for all applications and can be taken from the examples provided in the Application Notes listed in [Section 1.2](#). Thus, most of the new configuration is required in the application region of the diagram.

In the diagram, the following objects are represented by colour-coded boxes:

- Task (blue)
- ISR (light blue)
- Callback (yellow)
- Message queue (green)
- Mutex group (red)
- Interrupt source (purple)
- Hardware counter (brown)
- Software timer (orange)

For example, the following box represents a task:




The blue banner indicates a user task and the name of the task is included in quotes ("APP_UpdateDisplayTask" here). The white area of the box contains configurable attributes of the task (here, the execution priority and autostart status).

Relationships between the objects are indicated using colour-coded lines that run between them. For example:



- A red line between a task/ISR and a mutex group indicates that the task/ISR belongs to the mutex group.
- A dark-green arrowed line between a task and a message queue indicates that the task is allowed to post messages in the queue.
- A light-green arrowed line between a task and a message queue indicates that the task is allowed to collect messages from the queue.
- A blue arrowed line between a software timer and a task indicates that expiry of the timer will activate the task.
- A purple arrowed line shows the interrupt source that stimulates a certain ISR.

The graphic below defines "TickTimerException" as the interrupt source that triggers execution of the ISR "TickInterrupt".



Part IV: Appendices

A. Hardware Counter Details

A hardware counter is an abstract device used as a source of timing events to drive a number of software timers which are associated with it (see [Section 2.4.6](#)). It requires four user-defined callback functions to implement the following set of operations:

Enable and **Disable** the timer, **Get** the current count value, **Set** the compare value.

These callback functions are used by JenOS to manage the software timers and are defined using JenOS macros (detailed in [Section 8.1](#)).

A.1 Hardware Counter Operation

JenOS maintains the expiry times of the software timers associated with a hardware counter as a list of delta values - in other words, the number of ticks between now and the point at which the software timer is to expire. When a software timer is started, an entry in this list is created for its expiry time, the delta time being calculated relative to its closest earlier neighbour. The nearest timer later than the one inserted will have its delta value adjusted so that its expiry time is now relative to the newly inserted entry. Timers expiring later than the adjusted timer entry will not need their delta values changing since they are still measured relative to the adjusted timer event.

This can be more easily understood with the following example. Timer A is set to expire 25 ms from now (time 0), Timer B is set to expire 50 ms from now and Timer C at 75 ms from now. The list delta values are thus A=25, B=25, C=25, since B will expire 25 ms after A and C will expire 25 ms after B. 10 ms later, Timer D is set to expire after 20 ms. Timer D entry will expire in-between Timer A and Timer B (i.e. 30 ms from time 0), so it is inserted between A and B in the list. Its delta value from timer A is 5 ms (i.e. Timer D expires 5 ms after Timer A expires). However, since Timer D is now in-between Timers A and B in the list, Timer B's delta value must be adjusted so that it is relative to Timer D's expiry time. Therefore, Timer B's adjusted delta value is now 20. However, since Timer C's delta is calculated relative to Timer B, it does not need to change. So after the insertion, the list appears as A=25, D=5, B=20, C=25.

To run the counter, the first element from the list is removed and JenOS uses the Get and Set functions to calculate and set the compare value for the next expiry point (current count + delta from the list). The counter runs to the compare value and then generates an interrupt. JenOS then expires all software timers that match this point, and any others that may have expired during the time it was processing those which caused the interrupt. After performing all the processing to expire the timers (e.g. activating tasks associated with the timers), the next value from the list is added to the current hardware counter value and becomes the new compare value. This is managed by the function **OS_eExpireSWTimers()**.

Any free running counter with the ability to generate an interrupt when the counter reaches a compare value can be used.

A.2 Use of Tick Timer as Hardware Counter

The Tick Timer of the JN516x device is normally used as the hardware counter for the JenOS software timers:

- The Tick Timer is a free-running counter running at 16 MHz (62.5 ns). An interrupt occurs when the Tick Timer counter value matches the value in the Tick Timer compare register.
- The macro **APP_TIME_MS(*t*)** in **app_timer_driver.h** (in **Components/Utilities/Include**) gives the number of tick timer ticks that will occur in *t* ms (1 ms = 62.5 ns x 16000).
- The **OS_eStartSWTimer()** and **OS_eContinueSWTimer()** functions must be called with the number of ticks less than 2147483647. This equates to approximately two minutes with the 16-MHz hardware timer. Application software can time a longer period by maintaining a count of timer expiries.
- Routines which implement the callbacks required by the hardware counter can be found in **components/utilities/source/app_timer_driver.c** and consist of the following functions:
 - **APP_cbEnableTickTimer()**: Clears pending tick interrupts, sets the timer for continuous running and then enables Tick Timer interrupts.
 - **APP_cbDisableTickTimer()**: Disables tick interrupts and stops the Tick Timer.
 - **APP_cbGetTickTimer()**: Returns the current value of the Tick Timer.
 - **APP_cbSetTickTimerCompare()**: Sets the compare value of the Tick Timer with the value calculated - when this matches the Tick Timer counter value, an interrupt is generated.
 - **APP_isrTickTimer**: ISR called when Tick Timer value matches compare value - calls the OS function that manages the OS software timers, **OS_eExpireSWTimers()**.

B. Clearing Interrupts

When using JenOS, it is the programmer's responsibility to clear down an interrupt source within the relevant Interrupt Service Routine (ISR). JenOS only manages the Programmable Interrupt Controller (PIC) to set interrupt priorities and does not clear interrupts.

Unless an interrupt source is cleared down, the corresponding interrupt status bits will remain set, resulting in continuous interrupts. For interrupt sources among the JN516x peripherals, the JN516x Integrated Peripherals API contains a number of functions for clearing down some but not all peripheral interrupt sources. Clearing down JN516x peripheral interrupts is described below for the various peripheral blocks.

Where no API function exists for clearing a peripheral interrupt, a workaround is detailed. Callback functions registered through the Integrated Peripherals API are not permitted when using JenOS. Instead, you should link an interrupt source (purple box) and connecting it to an ISR (using a purple arrow), as illustrated in the diagram below.



Figure 1: Linking an Interrupt Source to an ISR



Note: Where a dedicated function is referenced below for clearing down a peripheral interrupt source, the function is fully detailed in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)*.

System Controller

The following interrupt sources associated with the System Controller can be cleared using the function **vAHB_ClearSystemEventStatus()**:

- System clock
- Comparator
- Pulse counter
- Random number generator
- Brownout detector

For further information, refer to the function description in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)*.

Interrupts from the following System Controller sources must be cleared using the functions indicated:

- DIO interrupts may be cleared using **u32AHI_DioInterruptStatus()** or **vAHI_DioWakeStatus()**
- Wake timer interrupts are cleared using **u8AHI_WakeTimerFiredStatus()**

Analogue Peripherals

There are two analogue peripheral interrupts:

- ADC/DAC conversion complete (CAPT) - this is set when an ADC conversion and a DAC conversion have taken place (ADC and DAC running concurrently)
- ADC conversion complete in accumulation mode (ADCACC) - this is set when a number of digital samples have been accumulated (2, 4, 8 or 16)

There are no API functions to clear down these interrupts and therefore a workaround is needed:

- The relevant interrupt status bits can be read (bit 0 for CAPT, bit 1 for ADCACC) using the following function call:

```
u32Data=u32REG_AnaRead(REG_ANPER_IS);
```

- The bits can be written back using the following function call:

```
vREG_AnaWrite(REG_ANPER_IS, u32Data);
```

These register access functions are not described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)* but are available in the header file **PeripheralRegs.h**, which is included in the SDK.

You should check both bits in the ISR and process accordingly, as there are situations in which both bits are set.

UARTs

The UART interrupts are set in response to a number of external conditions. The interrupt status of a UART (0 or 1) can be read using the function **u8AHI_UartReadInterruptStatus()**. For details of the returned value, refer to the function description in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)*.

The accessed register is read only. The only way to clear a UART interrupt is to remove the cause-condition.

Timers

A timer can generate interrupts on the rising and/or falling edges of the timer output and the function **u8AHI_TimerFired()** can be used to clear the interrupt source.

Tick Timer

Any pending Tick Timer interrupt can be cleared using the function **vAHI_TickTimerIntPendClr()**.

Serial Interface (2-wire)

A Serial Interface interrupt is asserted to indicate the status of a data transfer, such as a byte transfer having completed or loss of arbitration. The interrupt can be cleared using the function **bAHI_SiMasterSetCmdReg()** in the following call:

```
bAHI_SiMasterSetCmdReg( FALSE, FALSE, FALSE, FALSE, FALSE, TRUE );
```

SPI Master

If enabled, a SPI interrupt is generated when each transfer has completed. There is no API function to clear down this interrupt and therefore a workaround is needed:

- The relevant interrupt status bit can be read using the following function call (bit 0 is set if the interrupt is pending):

```
u32Data=u32REG_SpiRead(REG_SPIM_IS);
```

- The bit can be written back using the following function call:

```
u32REG_SpiWrite(REG_SPIM_IS,u32Data);
```

These register access functions are not described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)* but are available in the header file **PeripheralRegs.h**, which is included in the SDK.

Intelligent Peripheral

A Serial Interface interrupt is asserted to indicate the status of a data transfer, such as transaction completed. There is no API function to clear down this interrupt and therefore a workaround is needed:

- The relevant interrupt status bit can be read using the following function call (bit 6 is set if the interrupt is pending):

```
u32Data= u32REG_SpiIpRead(REG_INTPER_CTRL);
```

- The bit can be written back using the following function call:

```
u32REG_SpiIpWrite(REG_INTPER_CTRL,u32Data);
```

These register access functions are not described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)* but are available in the header file **PeripheralRegs.h**, which is included in the SDK.

Digital Audio Interface (DAI)

If enabled, a DAI interrupt is generated to indicate the completion of a data transfer. There is no API function to clear down this interrupt and therefore a workaround is needed:

- The relevant interrupt status bit can be read using the following function call (bit 0 is set if the interrupt is pending):

```
u32Data=u32REG_DaiRead(REG_DAI_INT);
```

- The bit can be written back using the following function call:

```
vREG_DaiWrite(REG_DAI_INT, u32Data);
```

These register access functions are not described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)* but are available in the header file **PeripheralRegs.h**, which is included in the SDK.

Sample FIFO

The Sample FIFO interrupts indicate the status of the Transmit and Receiver buffers. There is no API function to clear down this interrupt and therefore a workaround is needed:

- The relevant interrupt status bit can be read using the following function call (bit 0 is set if the interrupt is pending):

```
u32Data=u32REG_SampleFifoRead(REG_SFF_INT);
```

- The bit can be written back using the following function call:

```
u32REG_SampleFifoWrite(REG_SFF_INT, u32Data);
```

These register access functions are not described in the *JN516x Integrated Peripherals API User Guide (JN-UG-3087)* but are available in the header file **PeripheralRegs.h**, which is included in the SDK.

However, the application must also remove the condition(s) that caused the interrupt, otherwise the just-cleared interrupt will be immediately set again.

Revision History

Version	Date	Comments
1.0	25-Nov-2010	First release containing JenOS information taken from <i>ZigBee PRO Stack User Guide (JN-UG-3048)</i> , former <i>ZigBee PRO APIs Reference Manual (JN-RM-2041)</i> and former <i>ZigBee PRO Configuration Guide (JN-UG-3065)</i>
1.1	3-Mar-2011	Overlays feature added, co-operative tasks described and other minor updates made. Accompanying software files also updated
1.2	20-Sept-2011	Appendix on clearing interrupts added. Advice on de-activating software timers before sleeping added. Software timer 'Start' and 'Expire' functions modified. Other minor updates/corrections also made
1.3	24-Aug-2012	Minor updates/corrections made
1.4	19-Dec-2012	Updated for JN516x
1.5	4-July-2014	Internal edition with debug module updated and support for JN514x removed
1.6	9-Feb-2015	Removed tutorials and minor updates/corrections made
1.7	26-Aug-2015	Added Persistent Data Manager (PDM) for JN516x EEPROM
1.8	25-Aug-2016	Added section on CPU stack and heap sizes. Web addresses for NXP Wireless Connectivity pages updated

Important Notice

Limited warranty and liability - Information in this document is believed to be accurate and reliable. However, NXP Semiconductors does not give any representations or warranties, expressed or implied, as to the accuracy or completeness of such information and shall have no liability for the consequences of use of such information. NXP Semiconductors takes no responsibility for the content in this document if provided by an information source outside of NXP Semiconductors.

In no event shall NXP Semiconductors be liable for any indirect, incidental, punitive, special or consequential damages (including - without limitation - lost profits, lost savings, business interruption, costs related to the removal or replacement of any products or rework charges) whether or not such damages are based on tort (including negligence), warranty, breach of contract or any other legal theory.

Notwithstanding any damages that customer might incur for any reason whatsoever, NXP Semiconductors' aggregate and cumulative liability towards customer for the products described herein shall be limited in accordance with the *Terms and conditions of commercial sale* of NXP Semiconductors.

Right to make changes - NXP Semiconductors reserves the right to make changes to information published in this document, including without limitation specifications and product descriptions, at any time and without notice. This document supersedes and replaces all information supplied prior to the publication hereof.

Suitability for use - NXP Semiconductors products are not designed, authorized or warranted to be suitable for use in life support, life-critical or safety-critical systems or equipment, nor in applications where failure or malfunction of an NXP Semiconductors product can reasonably be expected to result in personal injury, death or severe property or environmental damage. NXP Semiconductors and its suppliers accept no liability for inclusion and/or use of NXP Semiconductors products in such equipment or applications and therefore such inclusion and/or use is at the customer's own risk.

Applications - Applications that are described herein for any of these products are for illustrative purposes only. NXP Semiconductors makes no representation or warranty that such applications will be suitable for the specified use without further testing or modification.

Customers are responsible for the design and operation of their applications and products using NXP Semiconductors products, and NXP Semiconductors accepts no liability for any assistance with applications or customer product design. It is customer's sole responsibility to determine whether the NXP Semiconductors product is suitable and fit for the customer's applications and products planned, as well as for the planned application and use of customer's third party customer(s). Customers should provide appropriate design and operating safeguards to minimize the risks associated with their applications and products.

NXP Semiconductors does not accept any liability related to any default, damage, costs or problem which is based on any weakness or default in the customer's applications or products, or the application or use by customer's third party customer(s). Customer is responsible for doing all necessary testing for the customer's applications and products using NXP Semiconductors products in order to avoid a default of the applications and the products or of the application or use by customer's third party customer(s). NXP does not accept any liability in this respect.

Export control - This document as well as the item(s) described herein may be subject to export control regulations. Export might require a prior authorization from competent authorities.

NXP Semiconductors

For online support resources and contact details of your local NXP office or distributor, refer to:

www.nxp.com